



Linux
Plumbers
Conference

| Richmond, VA | Nov. 13-15, 2023

Developing a Continuous eBPF Profiler

Looking Beneath the Kernel to Beyond the Clouds

[by Sumera | @sylfrena | @SumoOfShinovar](#)

 Polar Signals

Agenda

- High Resolution Continuous Profiling
- Userspace and Kernelspace
- Compilers and Runtimes
- Clouds and Kernels
- Low effort Debugging
- Future work



Linux
Plumbers
Conference | Richmond, VA | Nov. 13-15, 2023

STORY TIME



Linux
Plumbers
Conference | Richmond, VA | Nov. 13-15, 2023

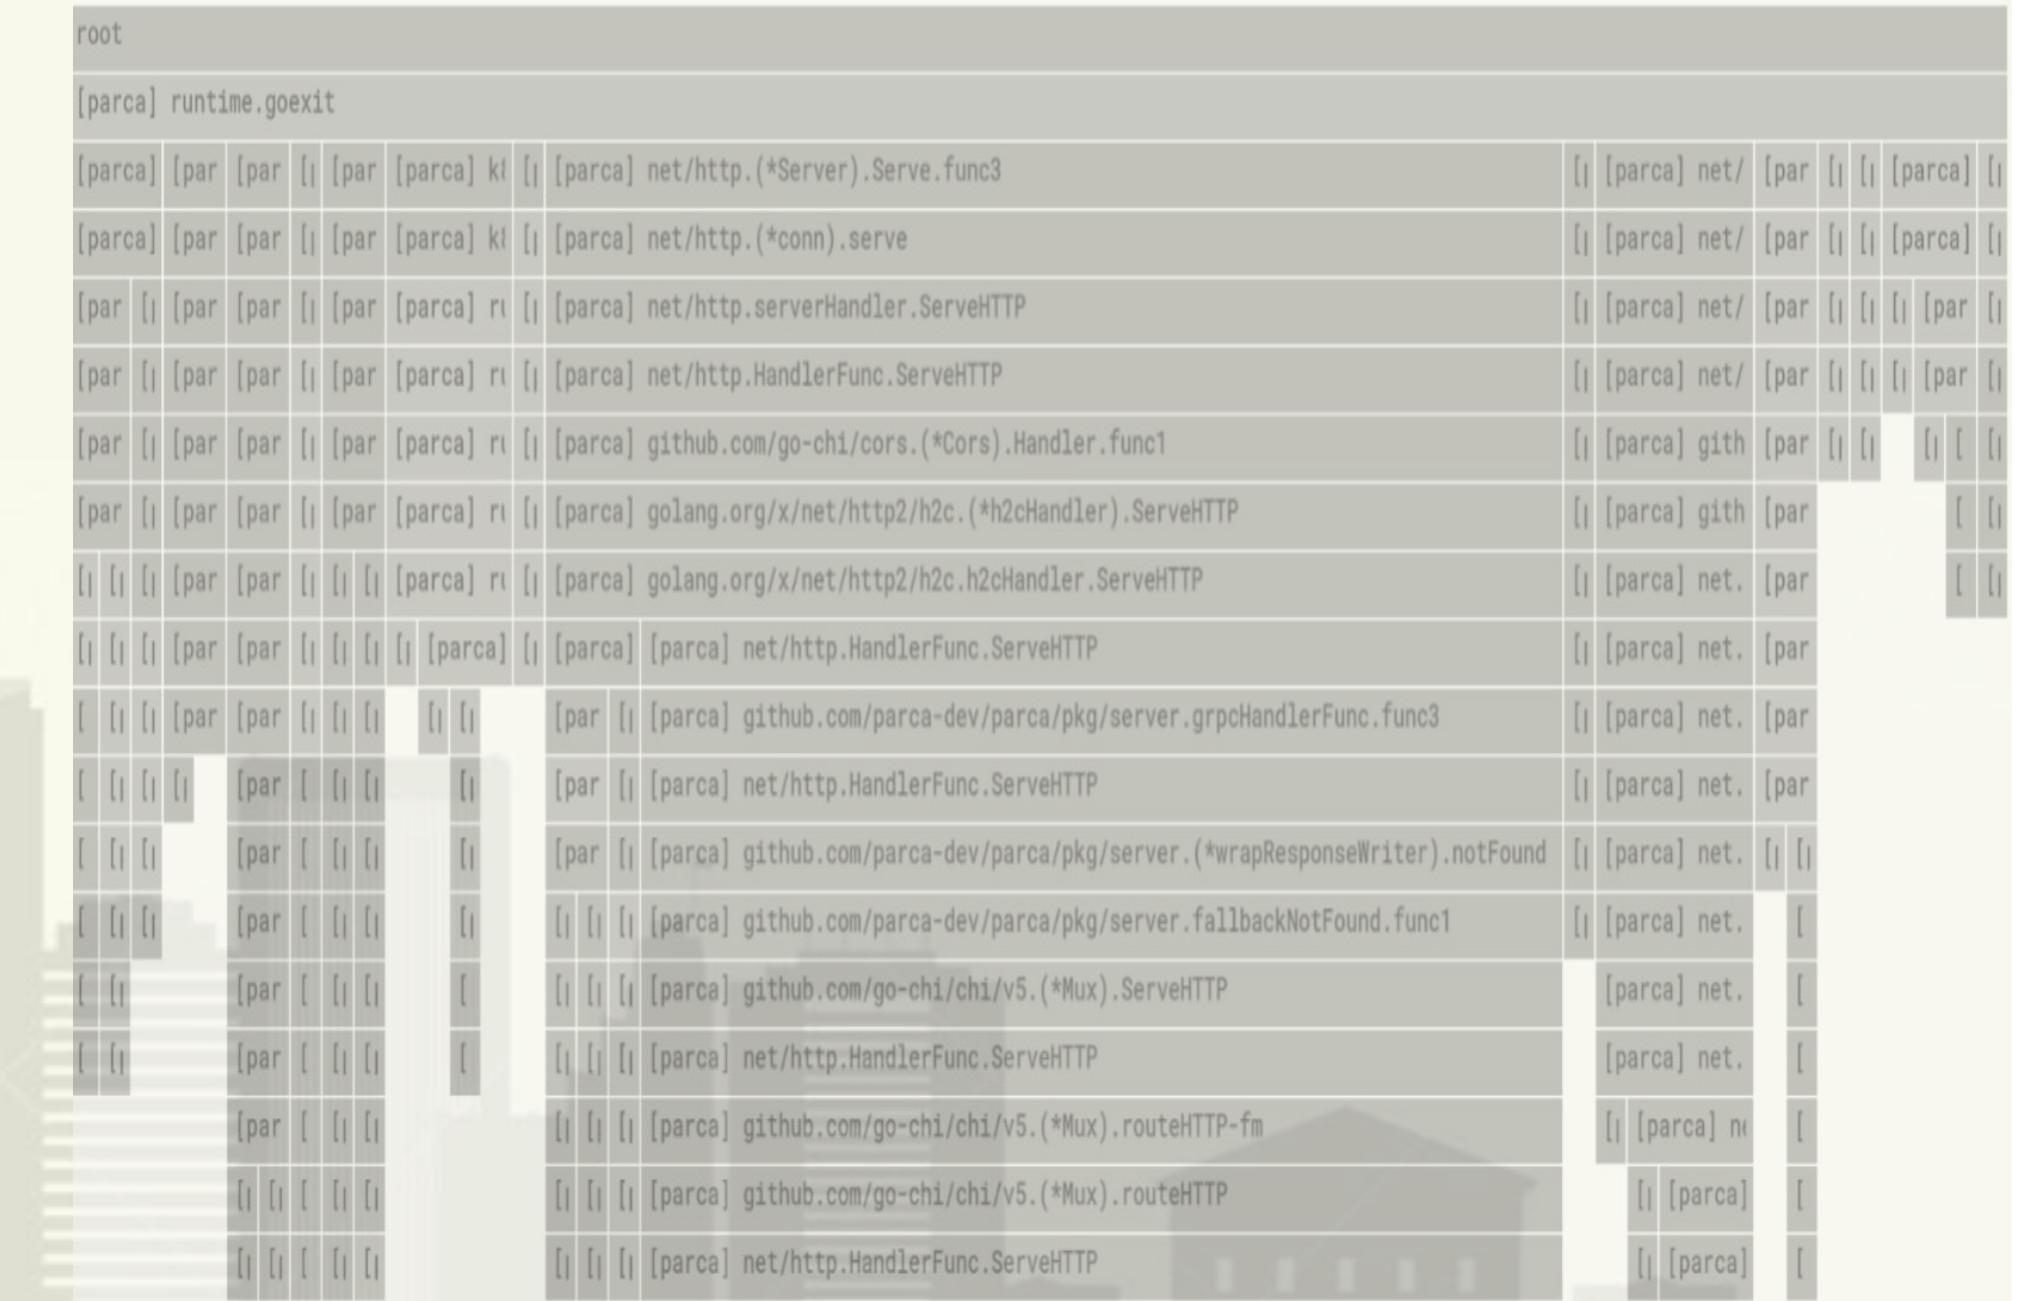
newbie kernel contributor

TO

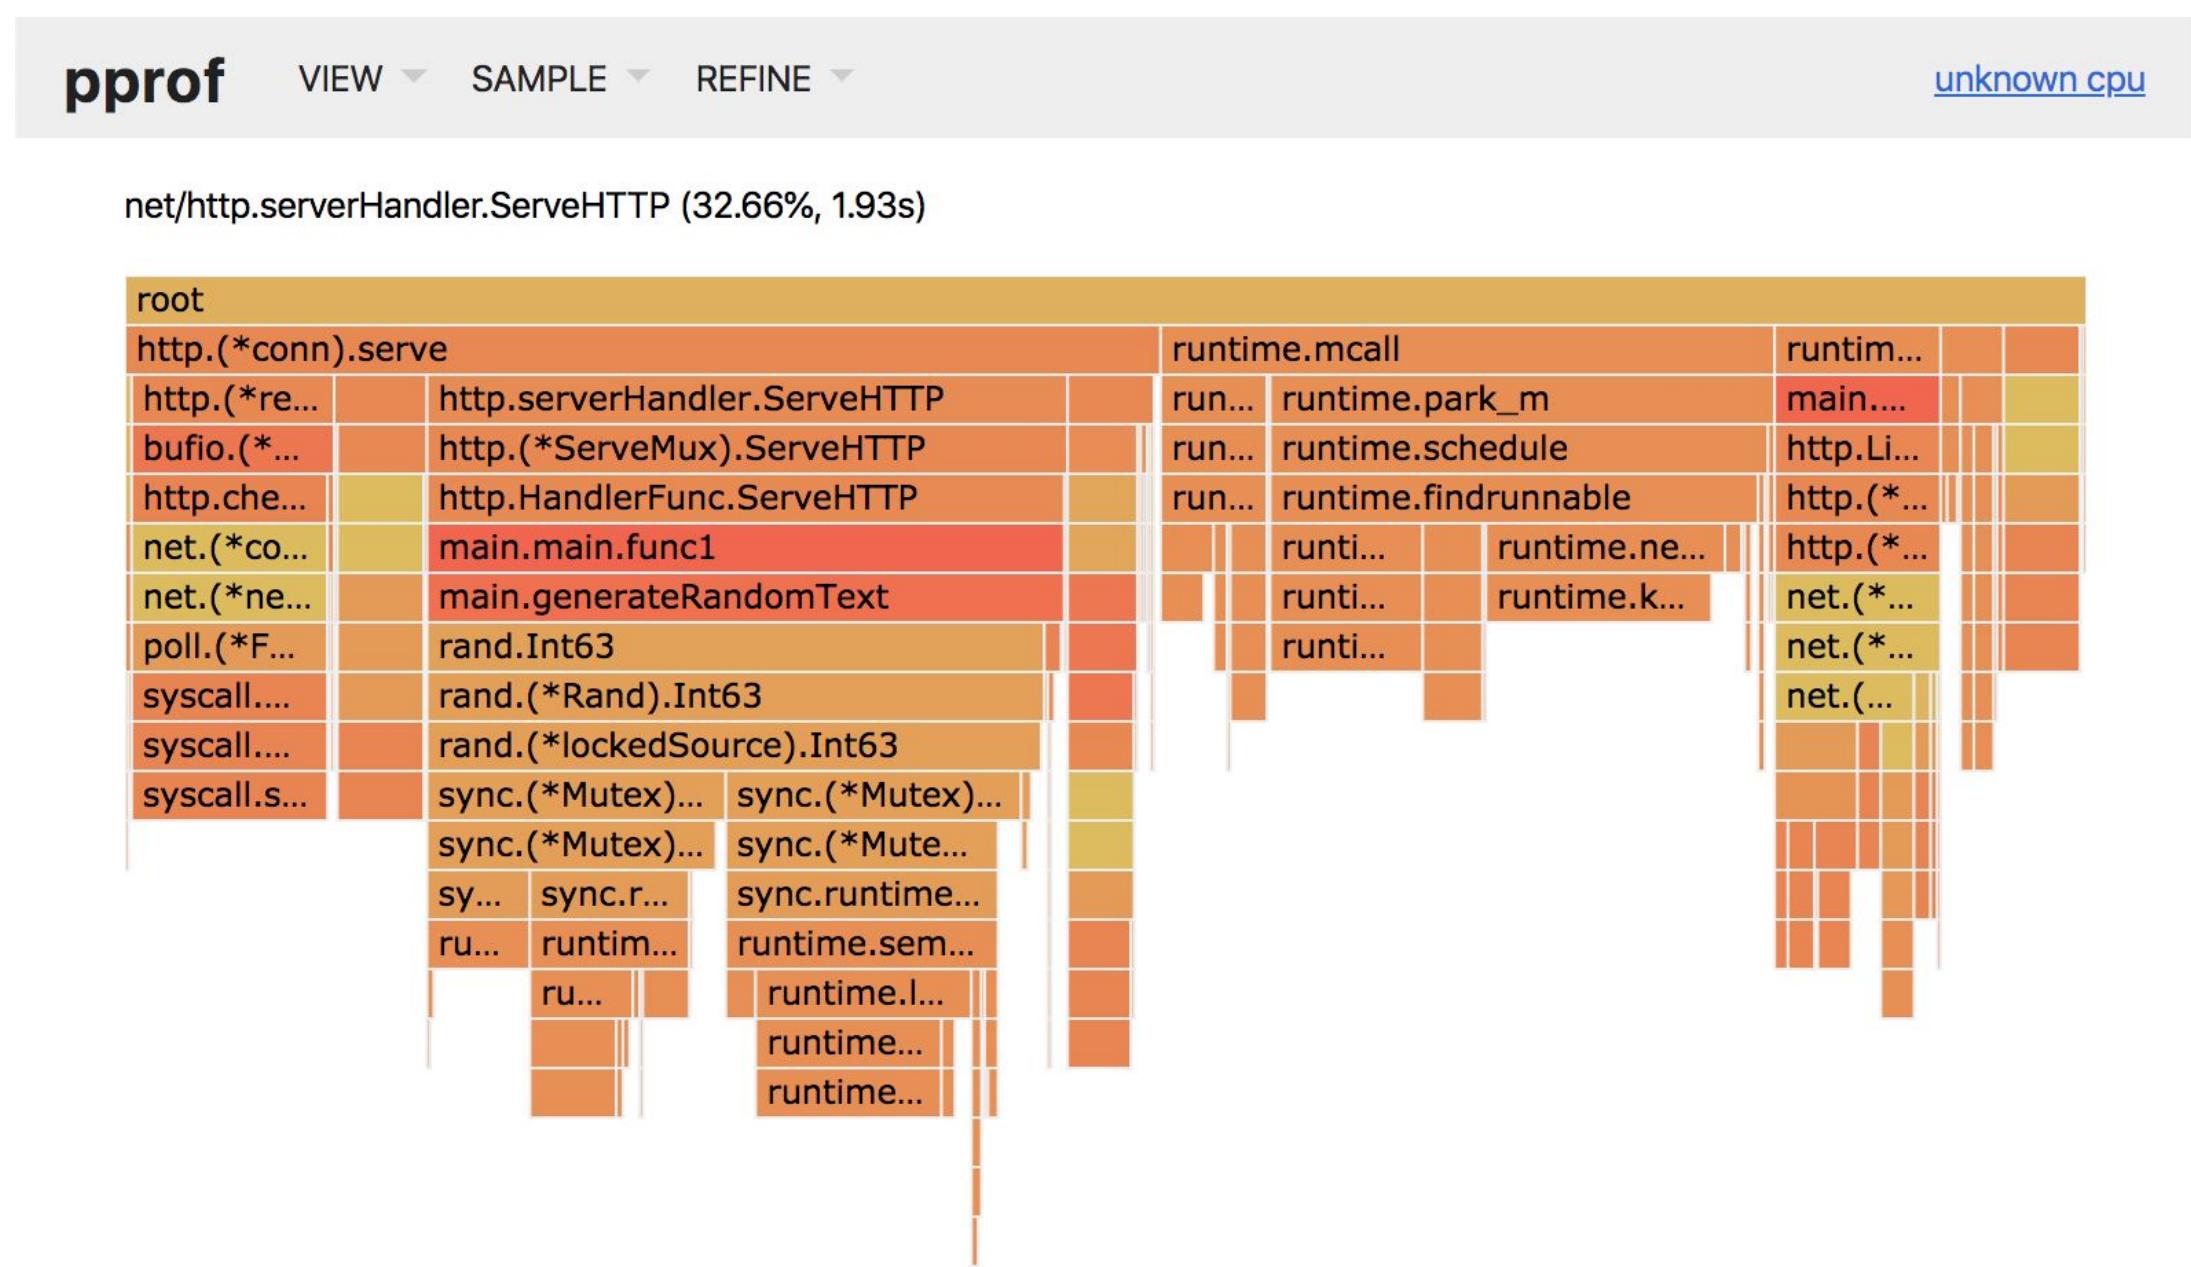
newbie profiler maintainer

parca

[kernel.kallsyms] parca runtime Everything else



- **ftrace**: dynamic tracing
 - **dmesg**: dynamic logging
 - **flamegraphs**: static, infrawide profile snapshots



Debugging tools

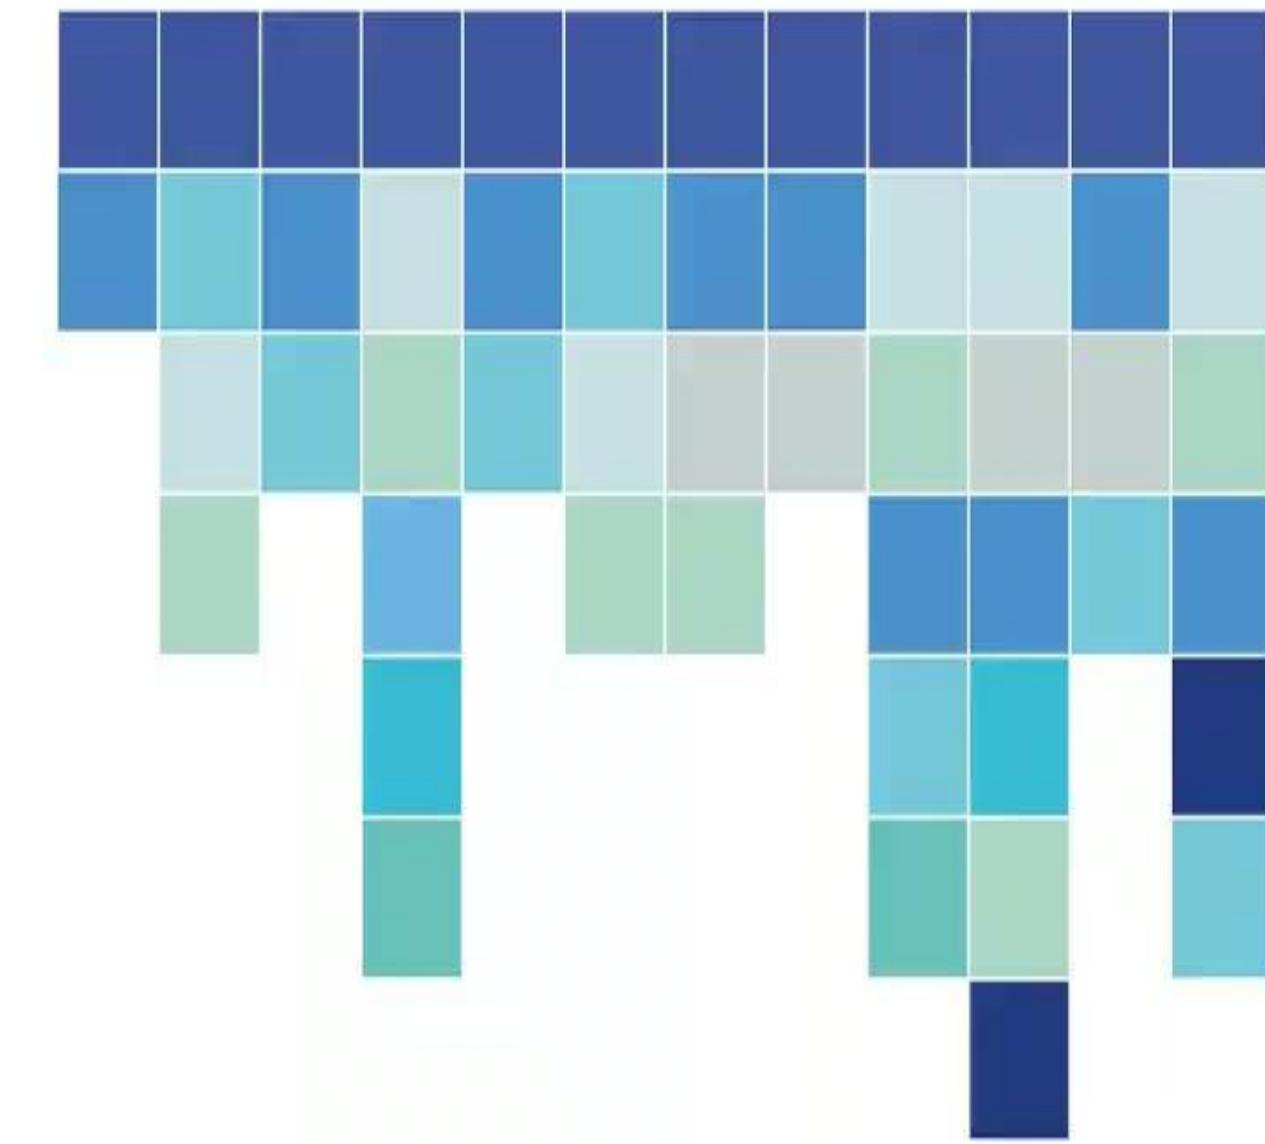
About Me

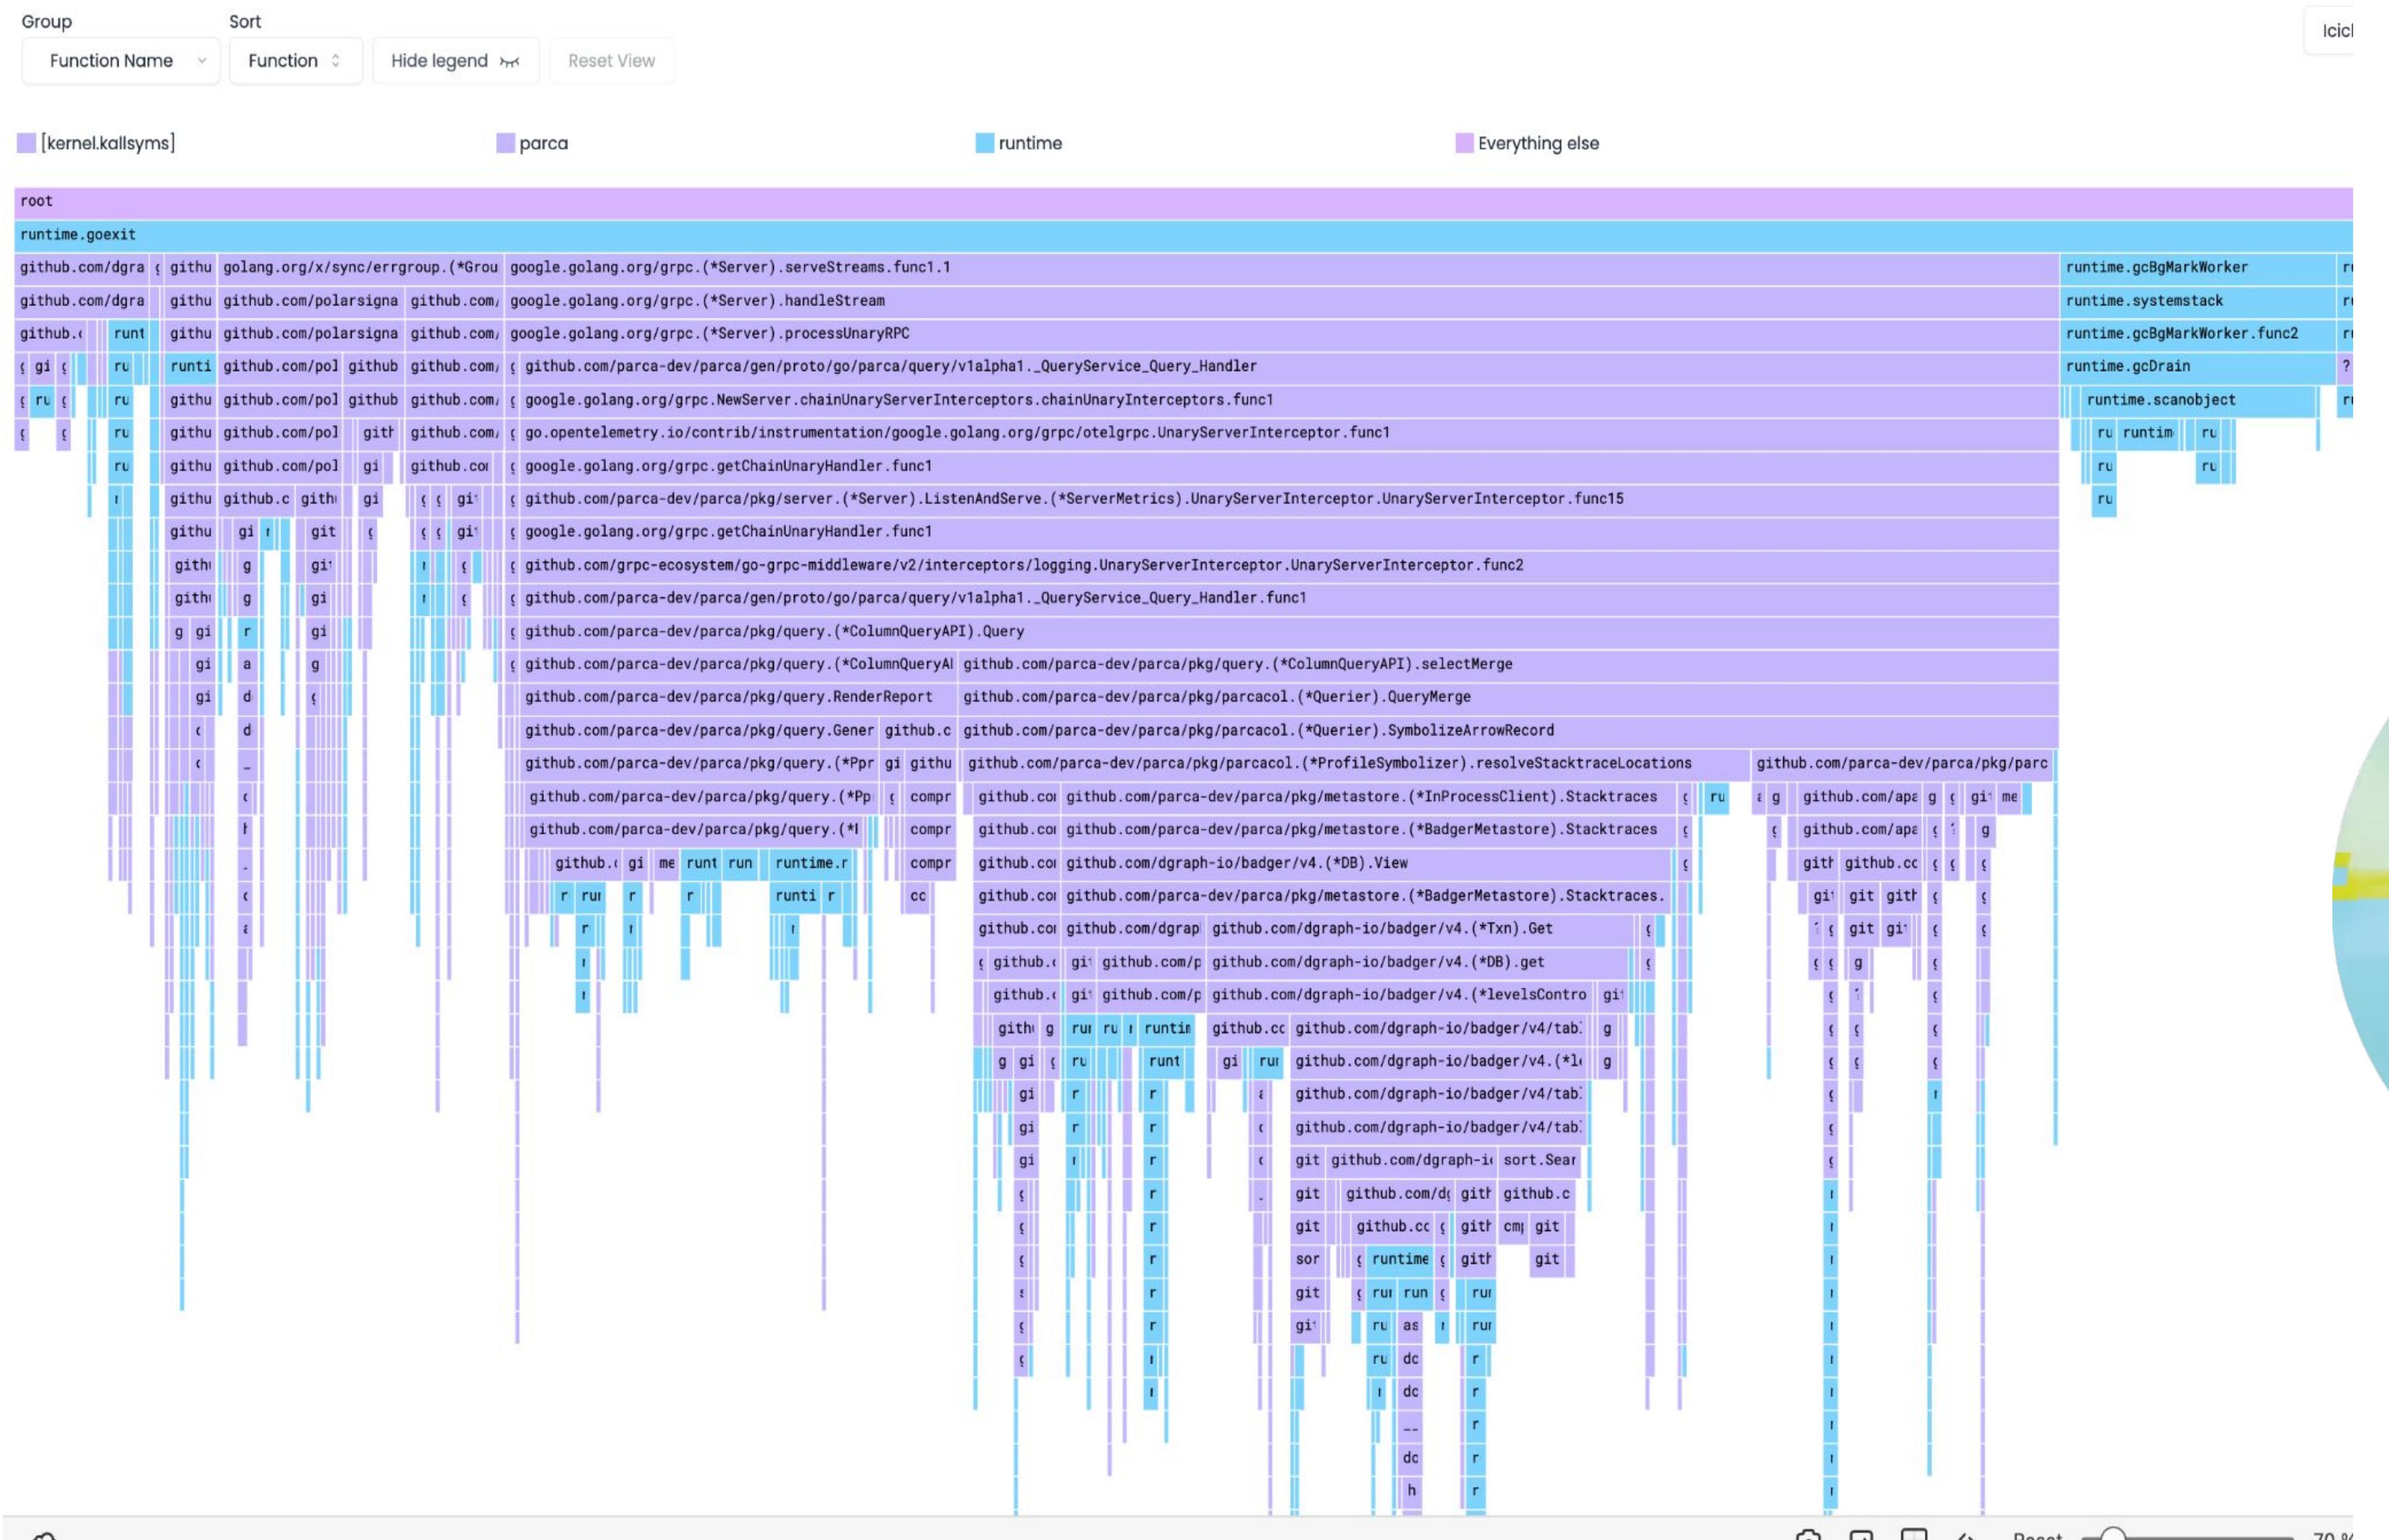
- Junior Software Engineer at Polar Signals
- Open Source Maintainer
 - Parca and Parca-Agent



Sumera Priyadarsini

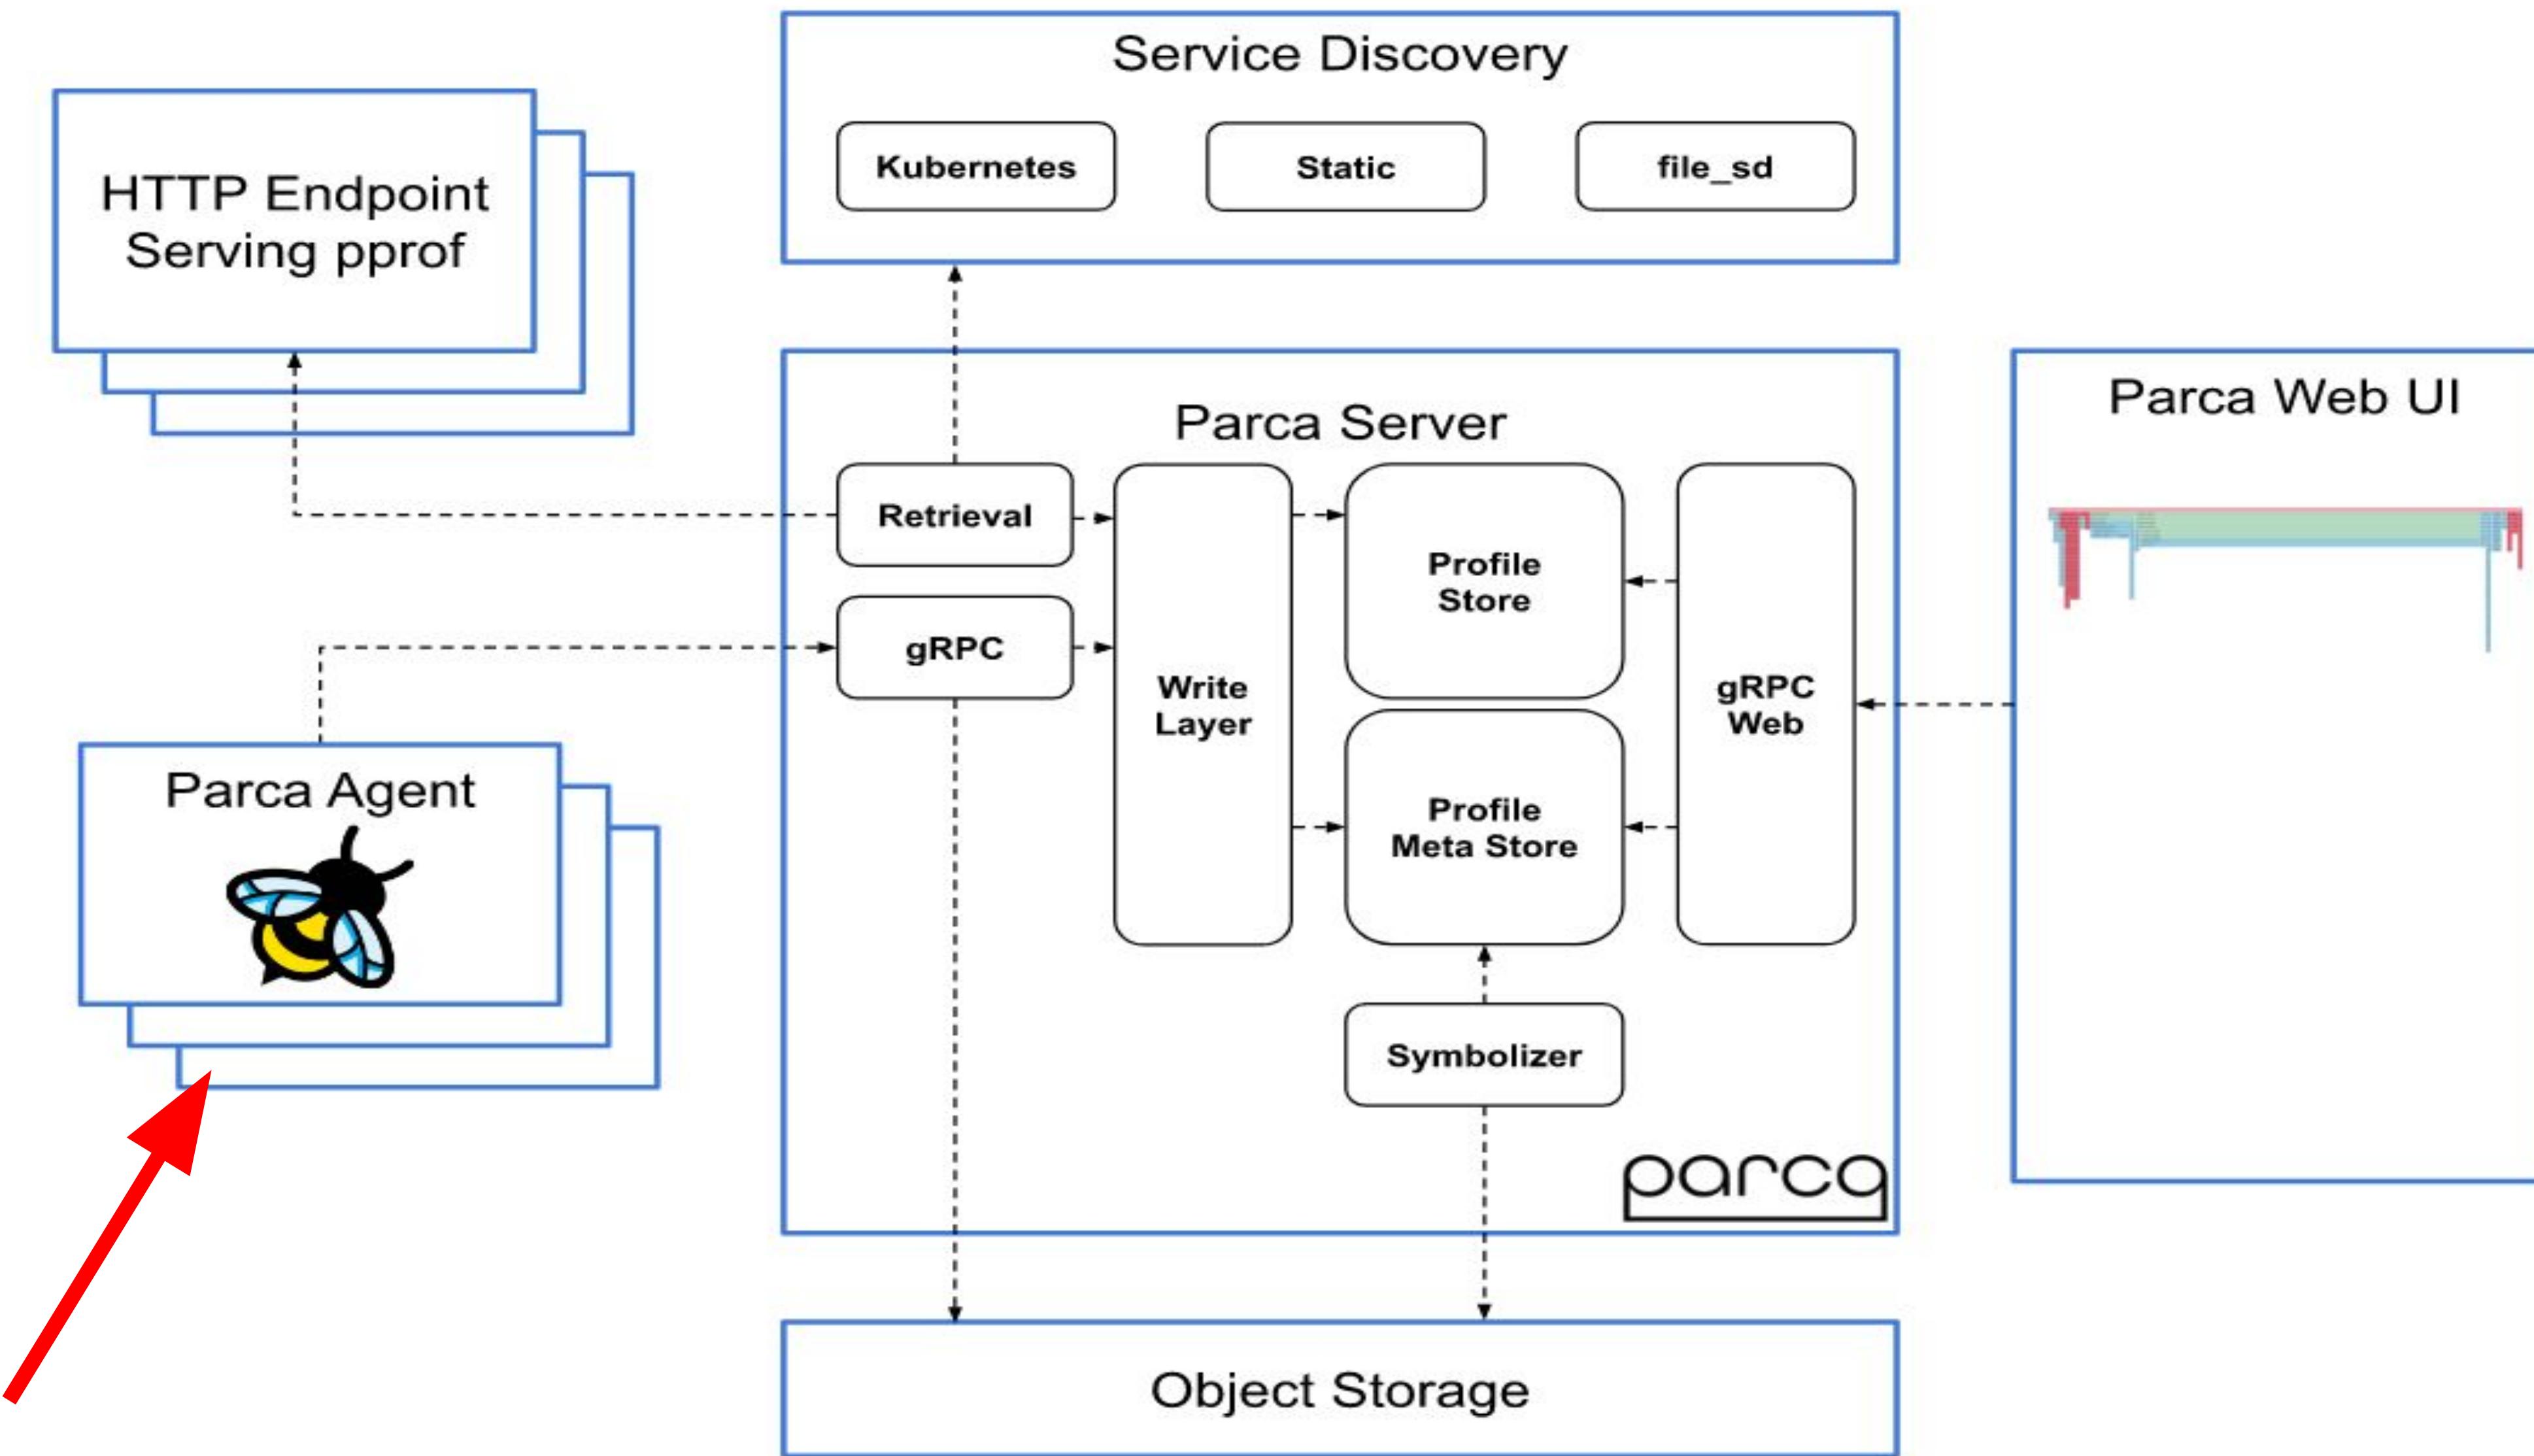
[@sylfrena](https://twitter.com/sylfrena)
[@SumoOfShinovar](https://github.com/SumoOfShinovar)





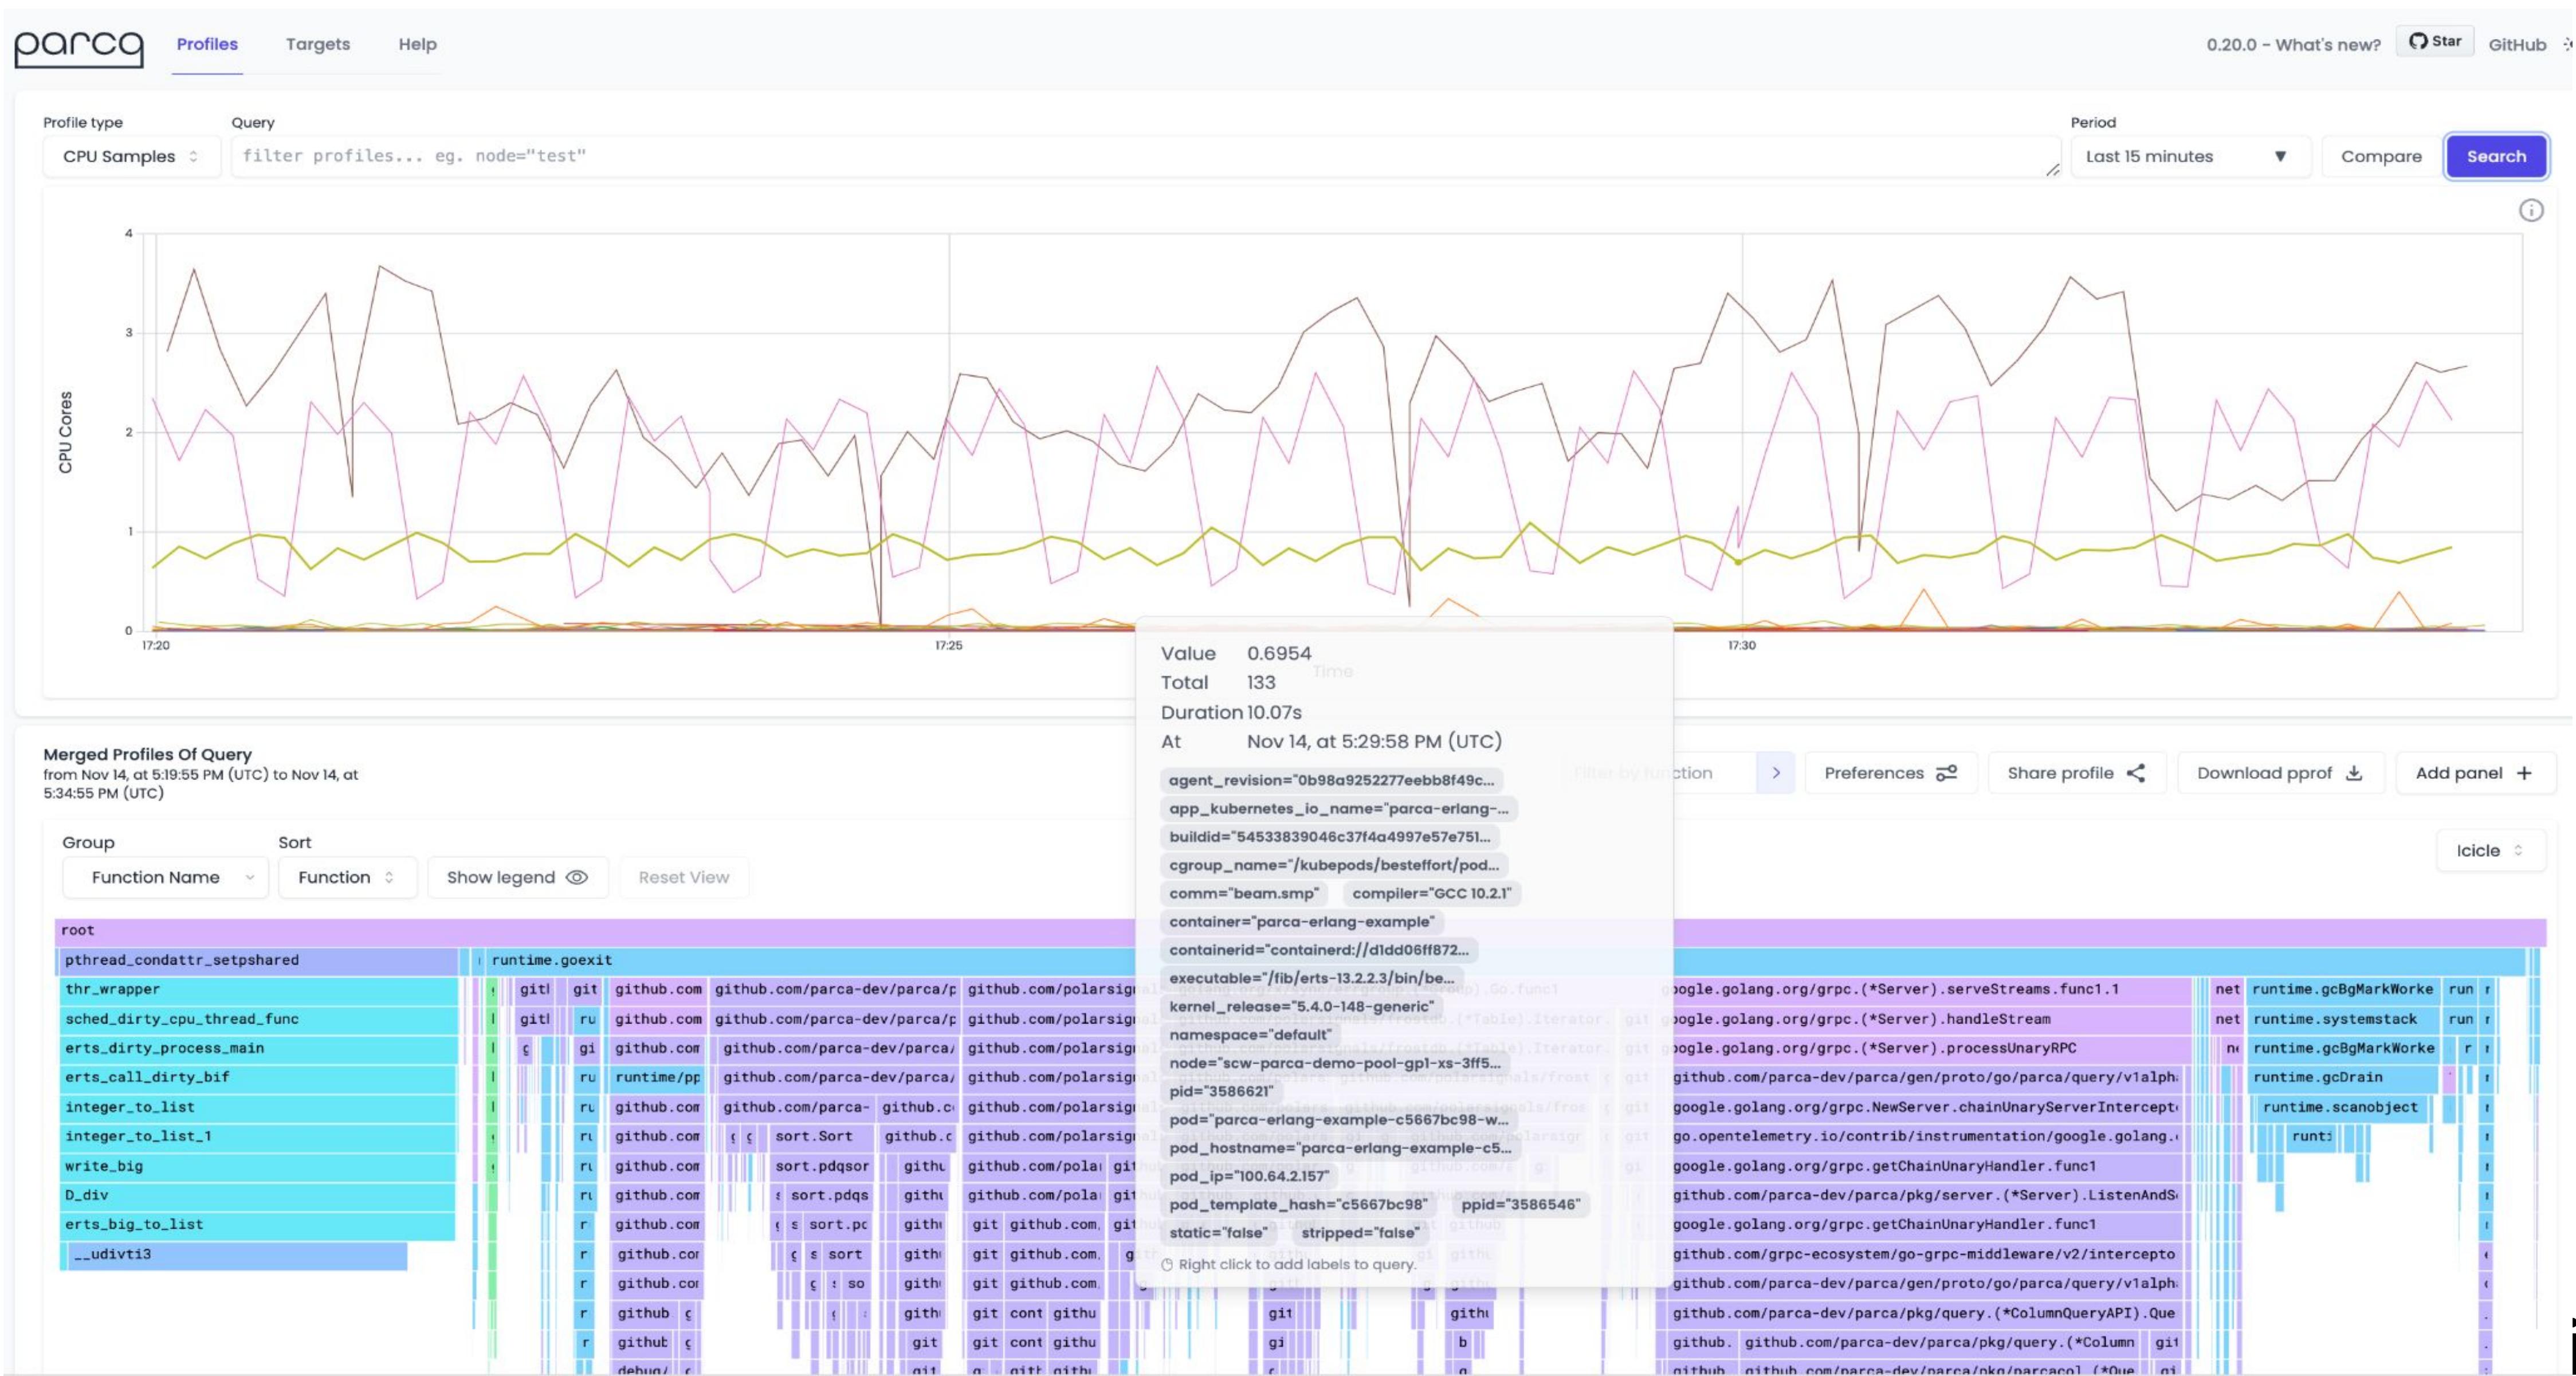


Polar Signals

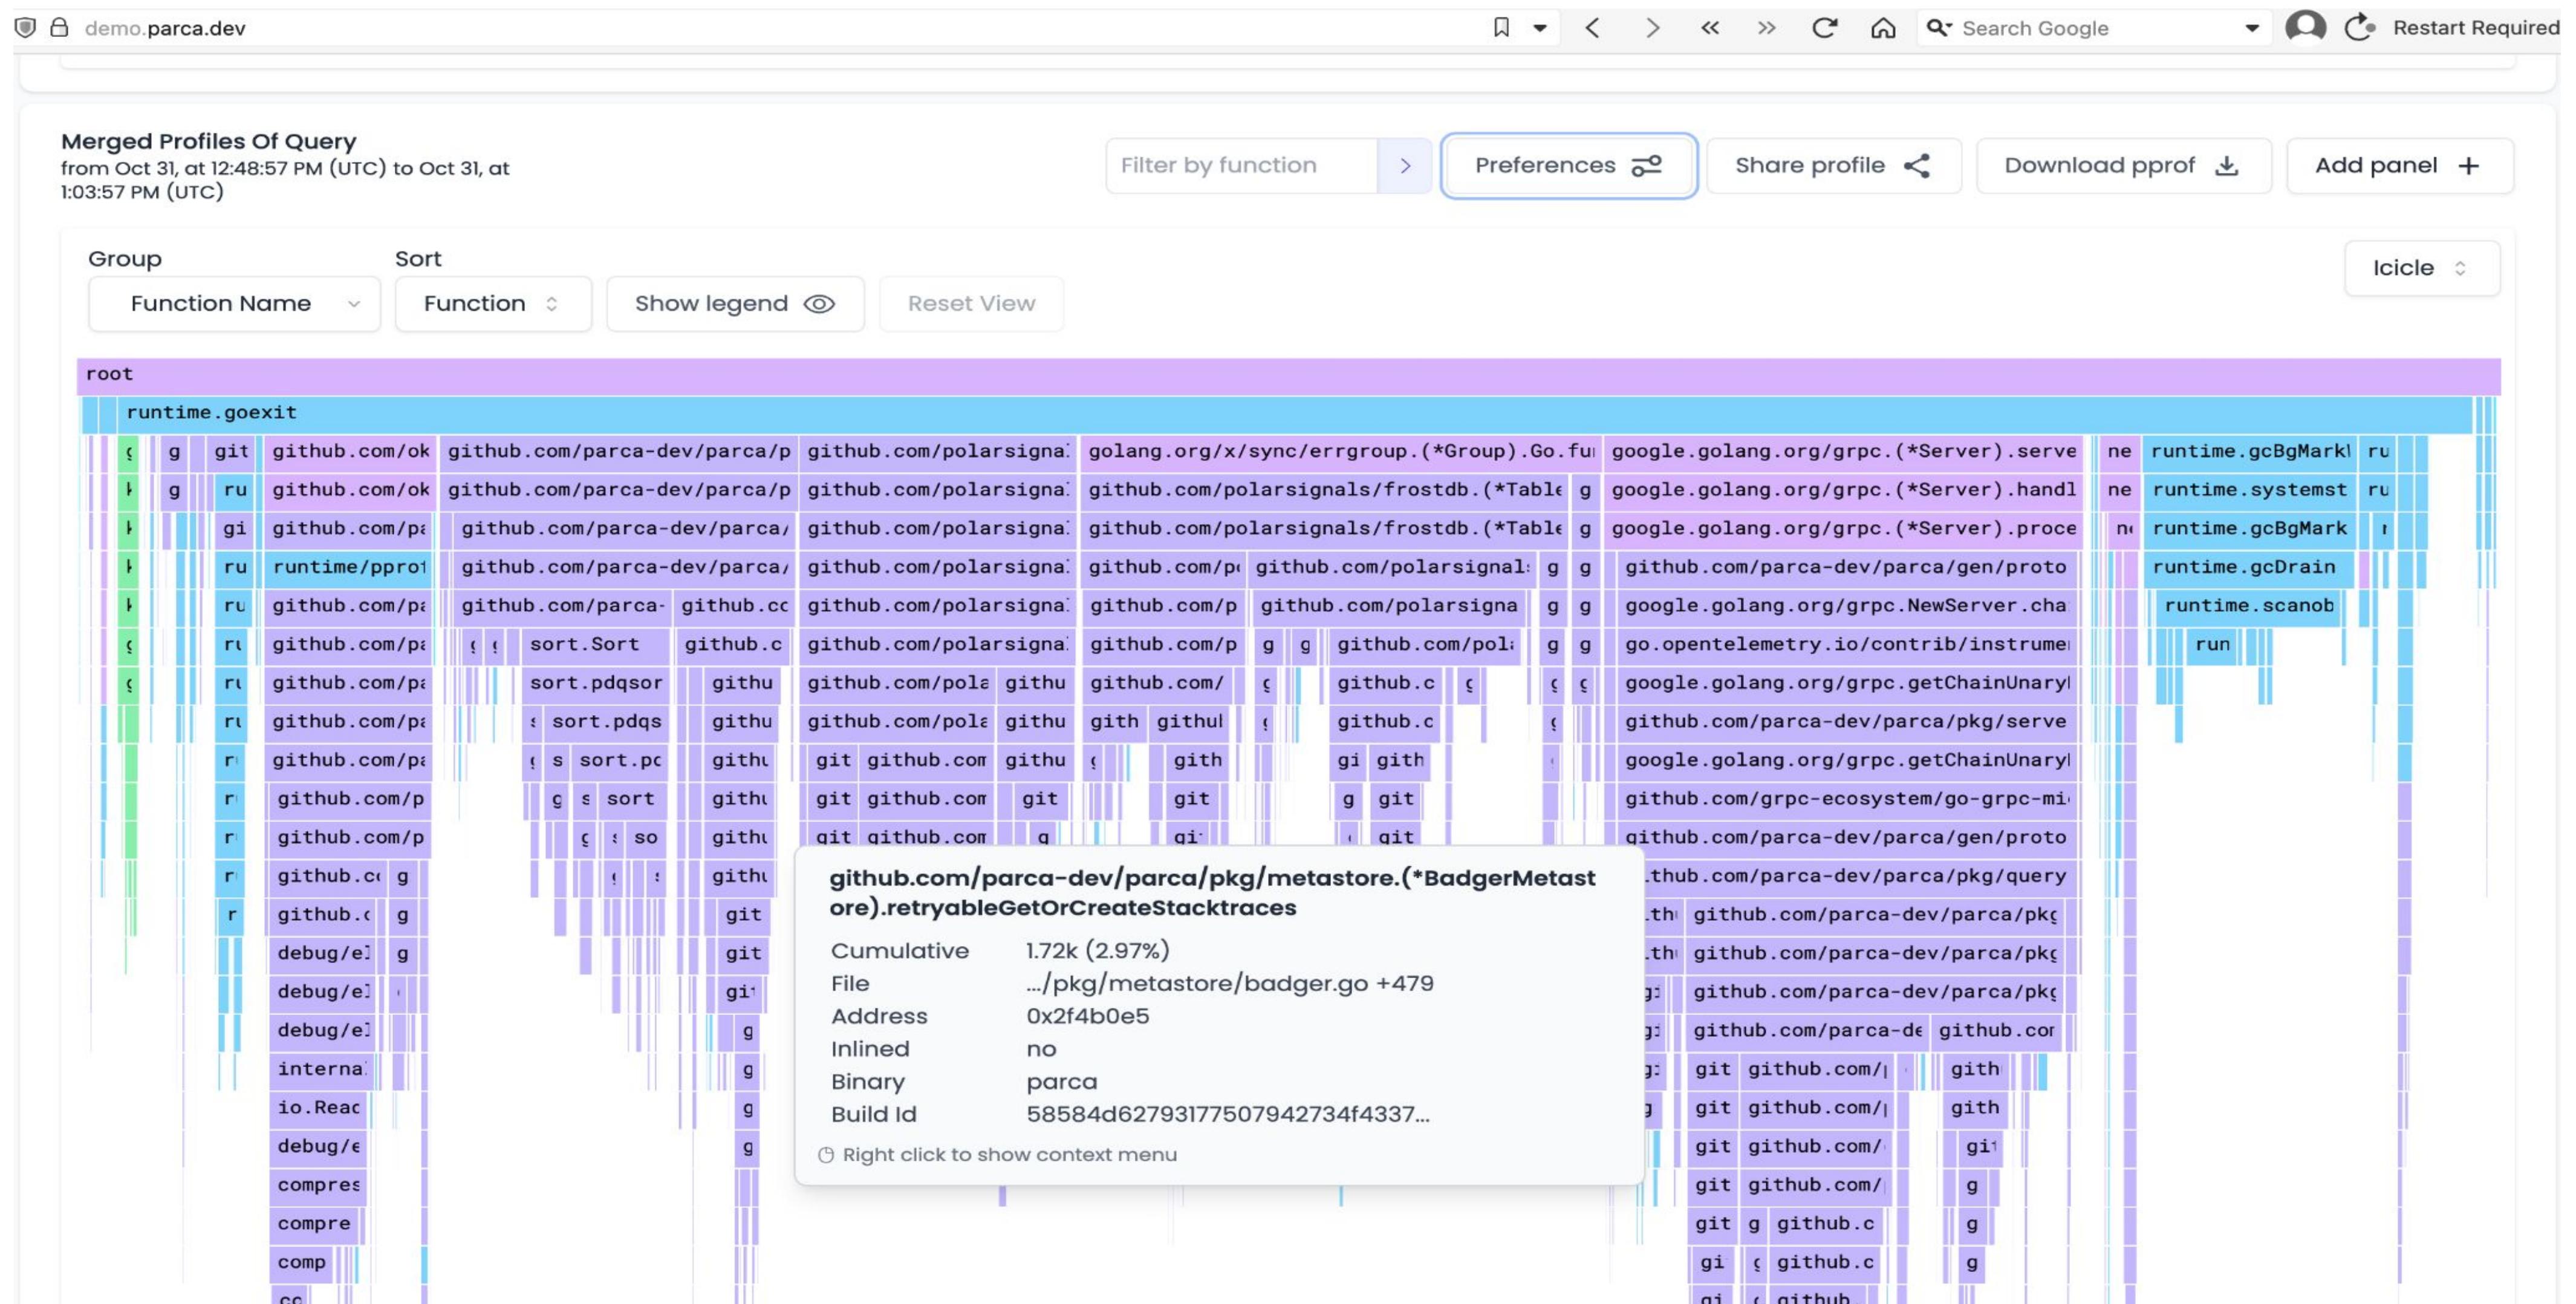


High Resolution Dynamic Continuous Profiling

High Resolution Profiling



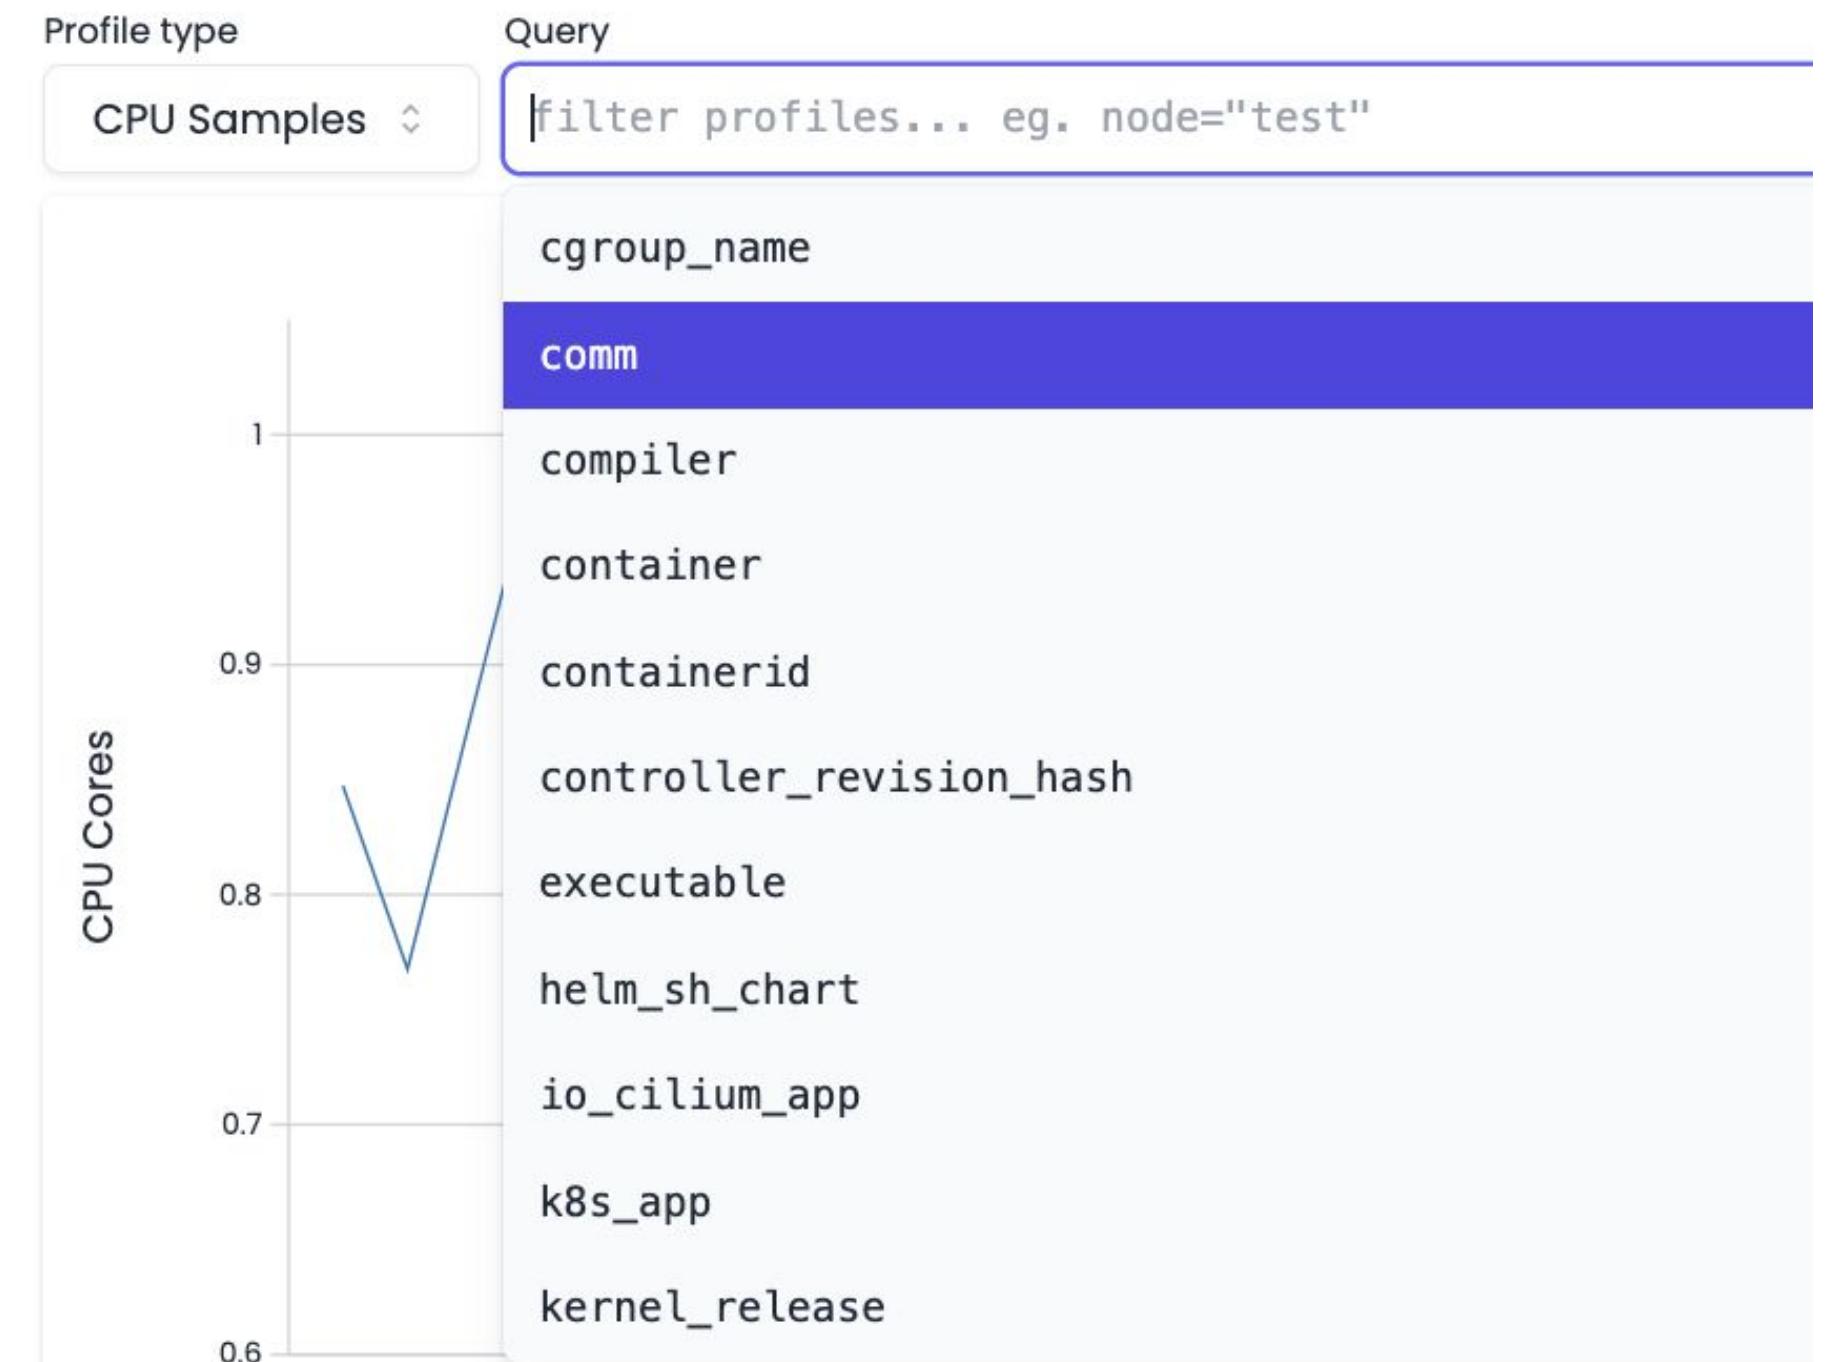
High Resolution Profiling



High Resolution Profiling

Flat	Cumulative	Cumulative (%)	Name
0	13.52k	106.27%	google.golang.org/grpc.getChainUnaryHandler.func1
0	12.7k	99.83%	runtime.goexit
0	6.77k	53.18%	google.golang.org/grpc.(*Server).processUnaryRPC
0	6.77k	53.18%	google.golang.org/grpc.(*Server).handleStream
0	6.77k	53.18%	google.golang.org/grpc.(*Server).serveStreams.func1.1
0	6.76k	53.16%	go.opentelemetry.io/contrib/instrumentation/google.golang.org/grpc/otelgrpc.UnaryServerInterceptor.func1
0	6.76k	53.16%	google.golang.org/grpc.NewServer.chainUnaryServerInterceptors.chainUnaryInterceptors.func1
0	6.76k	53.14%	github.com/parca-dev/parca/pkg/server.(*Server).ListenAndServe.(*ServerMetrics).UnaryServerInterceptor.UnaryServerInterceptor.func15
0	6.76k	53.13%	github.com/grpc-ecosystem/go-grpc-middleware/v2/interceptors/logging.UnaryServerInterceptor.UnaryServerInterceptor.func2
0	6.52k	51.23%	github.com/parca-dev/parca/pkg/query.(*ColumnQueryAPI).Query
0	6.52k	51.23%	github.com/parca-dev/parca/gen/proto/go/parca/query/v1alpha1._QueryService_Query_Handler.func1
0	6.52k	51.23%	github.com/parca-dev/parca/gen/proto/go/parca/query/v1alpha1._QueryService_Query_Handler
0	5.42k	42.59%	github.com/parca-dev/parca/pkg/parcacol.(*Querier).QueryMerge
0	5.42k	42.59%	github.com/parca-dev/parca/pkg/query.(*ColumnQueryAPI).selectMerge
1	5.41k	42.52%	github.com/parca-dev/parca/pkg/parcacol.(*Querier).SymbolizeArrowRecord
44	3.28k	25.79%	github.com/parca-dev/parca/pkg/parcacol.(*ProfileSymbolizer).resolveStacktraceLocations
5	2.99k	23.51%	runtime.systemstack
0	2.95k	23.19%	github.com/dgraph-io/badger/v4.(*DB).View
0	2.52k	19.82%	github.com/parca-dev/parca/pkg/metastore.(*BadgerMetastore).Stacktraces
0	2.52k	19.82%	github.com/parca-dev/parca/pkg/metastore.(*InProcessClient).Stacktraces
4	2.49k	19.59%	github.com/parca-dev/parca/pkg/metastore.(*BadgerMetastore).Stacktraces.func1
10	2.33k	18.33%	github.com/dgraph-io/badger/v4.(*Txn).Get
0	2.22k	17.44%	runtime.gcBgMarkWorker
0	2.21k	17.36%	runtime.gcBgMarkWorker.func2
94	2.21k	17.34%	runtime.gcDrain
13	2.15k	16.86%	github.com/dgraph-io/badger/v4.(*DB).get
808	2.13k	16.74%	runtime.scanobject
148	2.11k	16.58%	github.com/parca-dev/parca/pkg/parcacol.BuildArrowLocations
0	1.98k	15.55%	golang.org/x/sync/errgroup.(*Group).Go.func1

parca



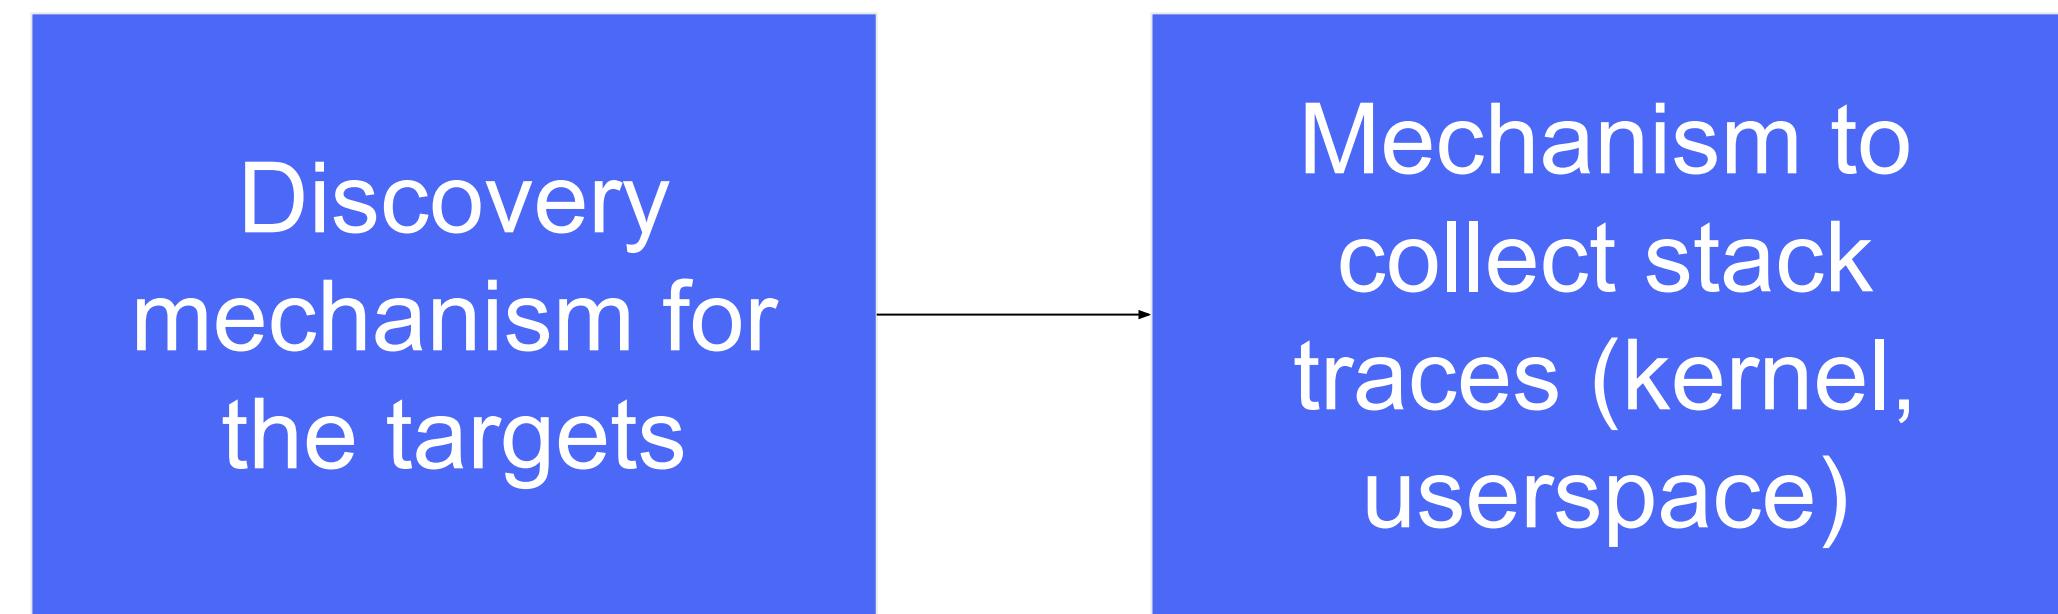
 Polar Signals

eBPF Profilers: Userspace and Kernelspace

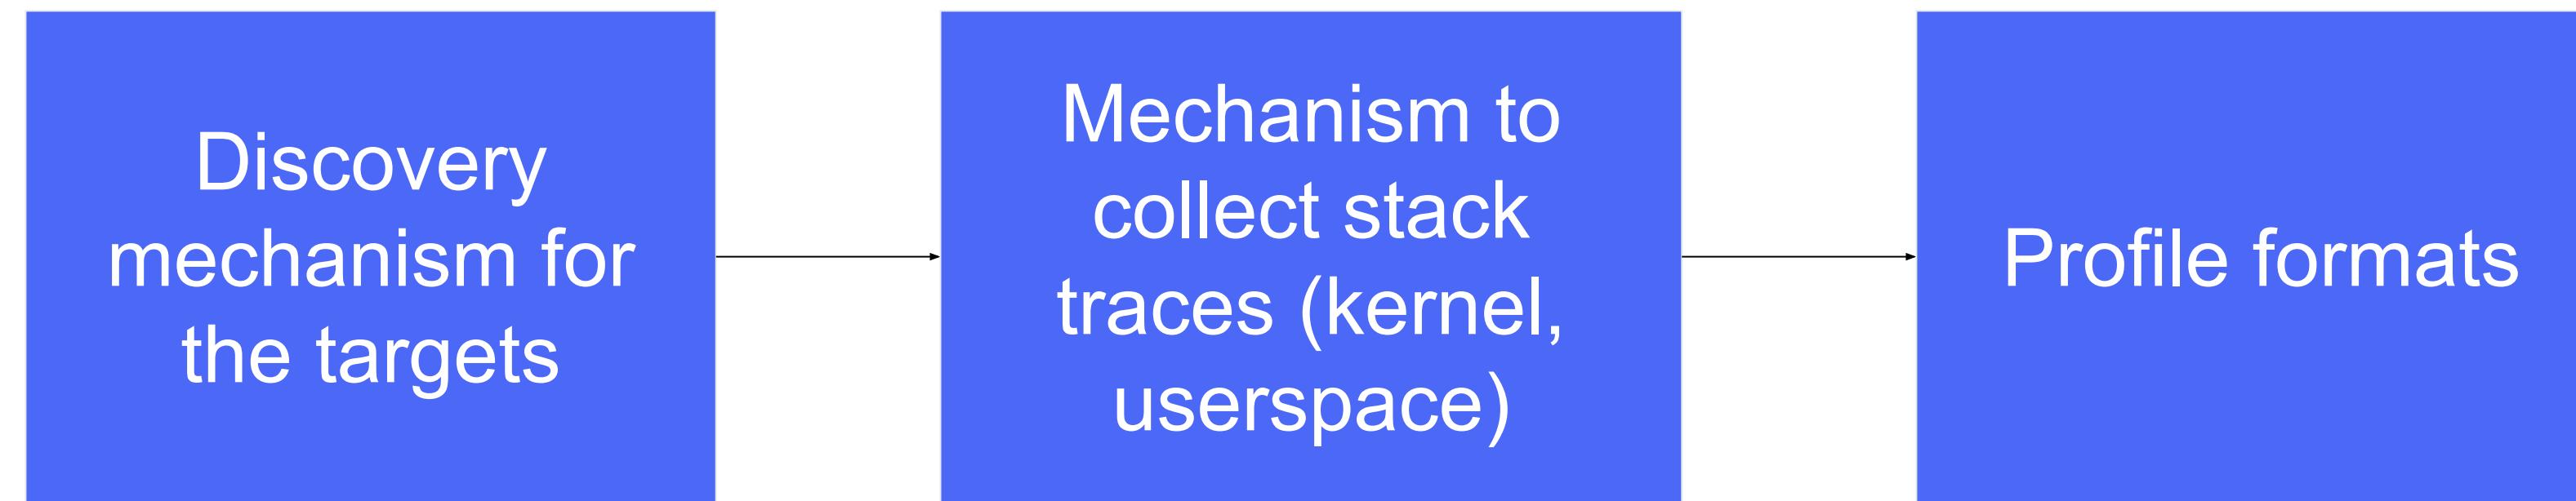
Profilers for the cloud native environment

Discovery
mechanism for
the targets

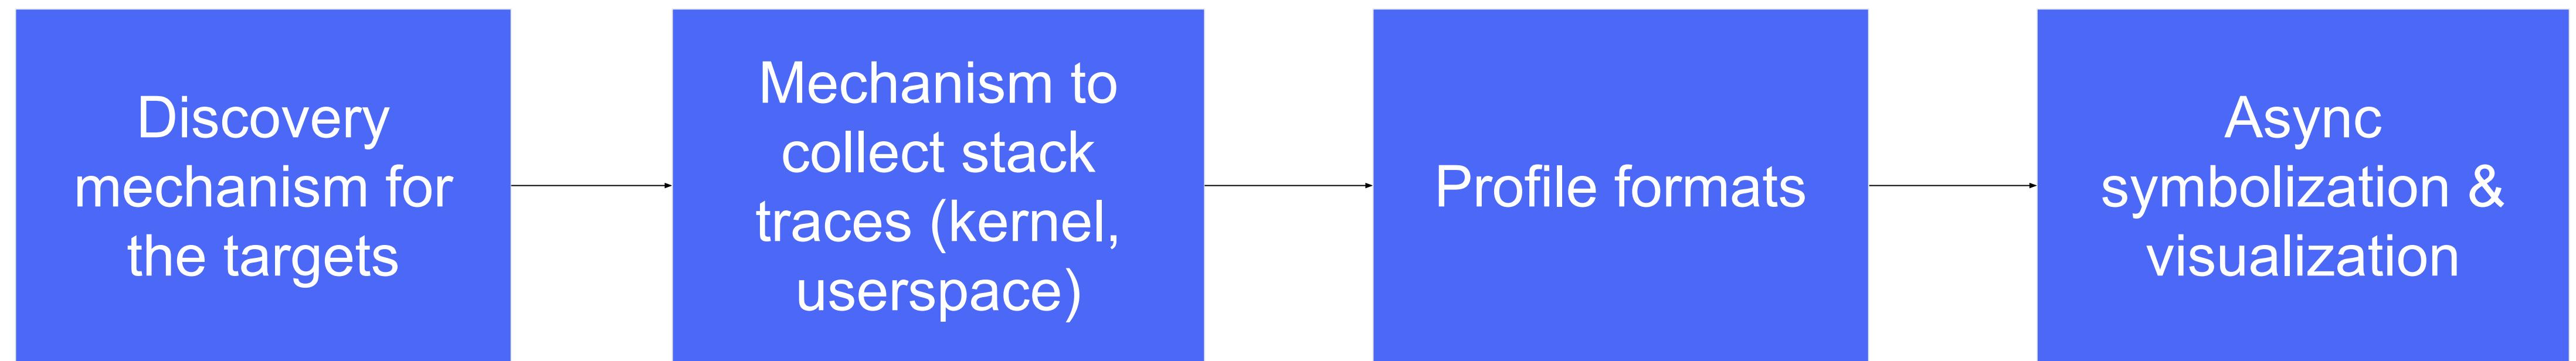
Profilers for the cloud native environment



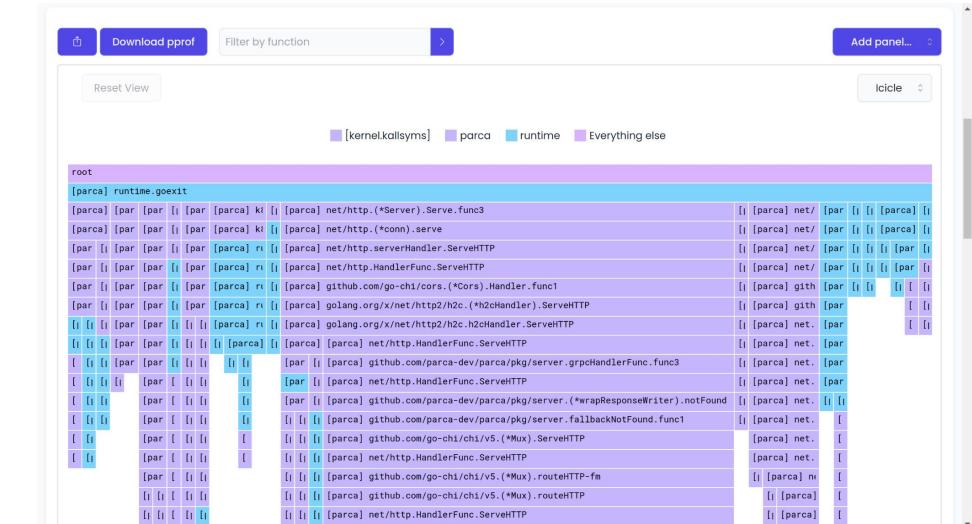
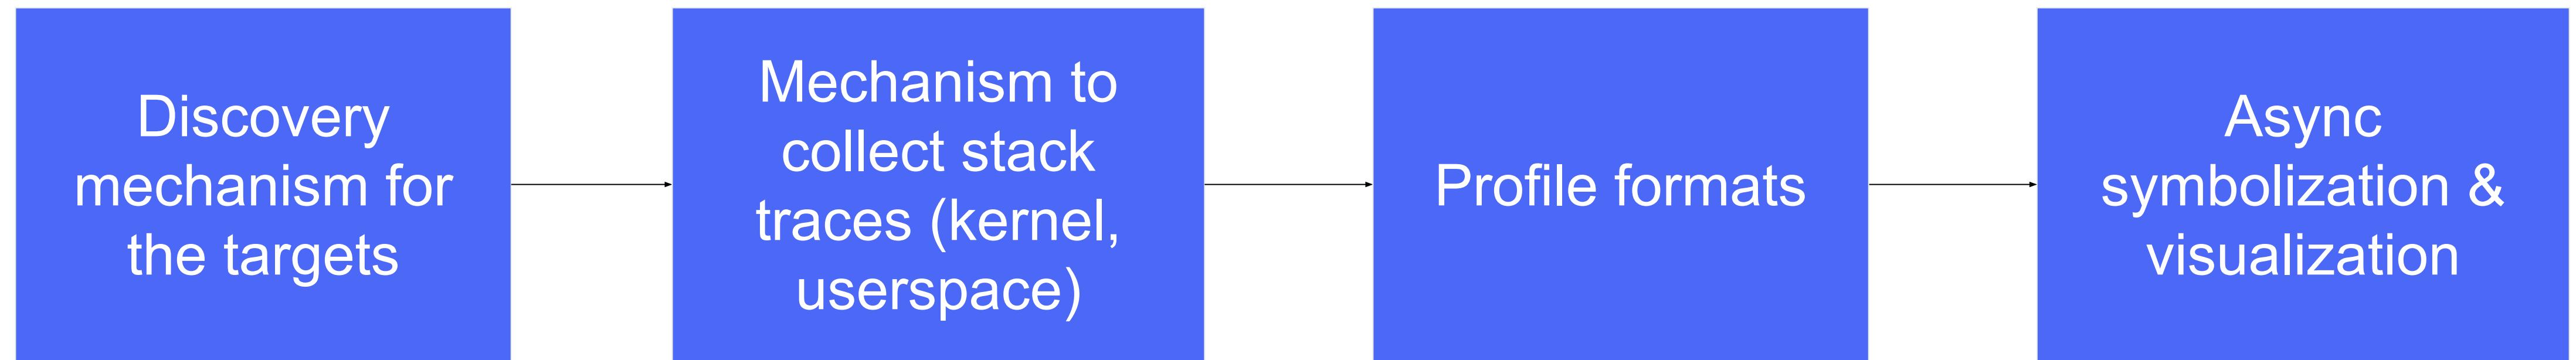
Profilers for the cloud native environment



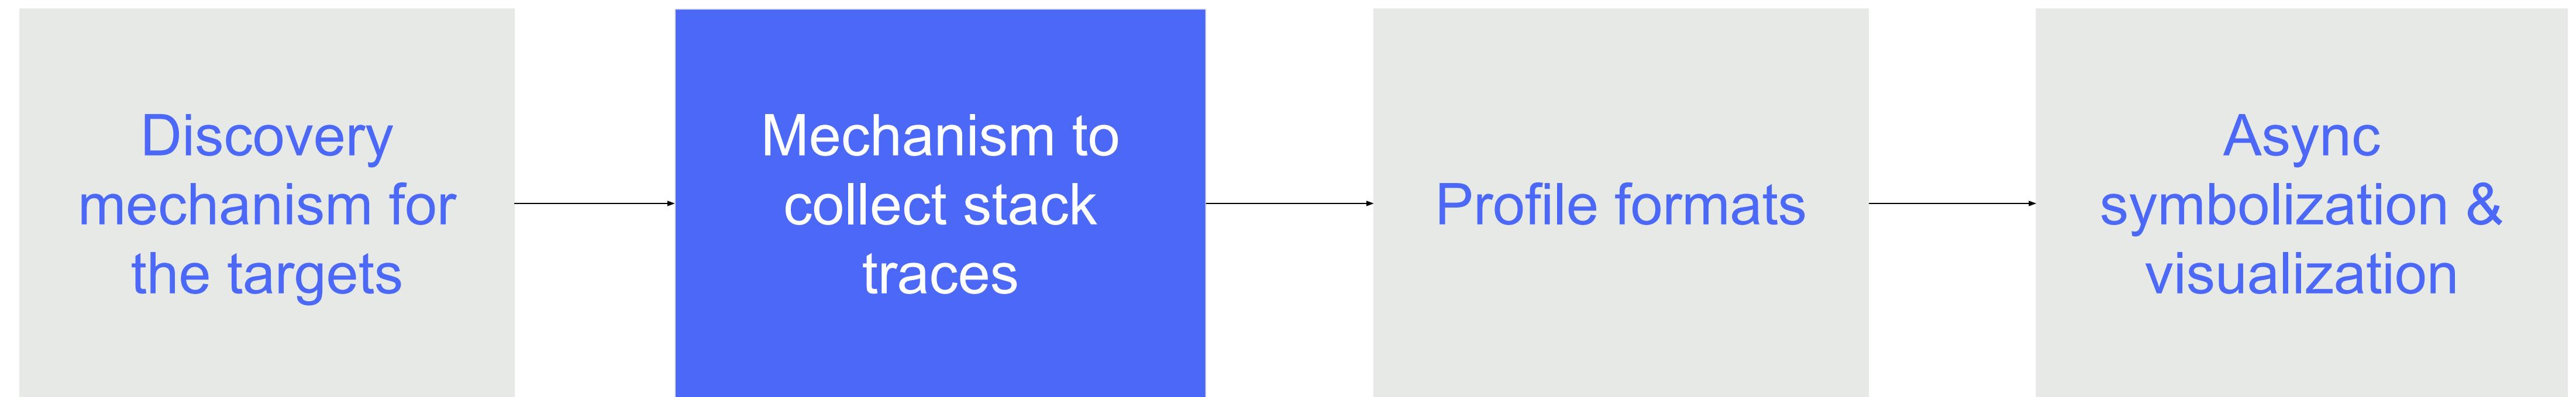
Profilers for the cloud native environment



Profilers for the cloud native environment

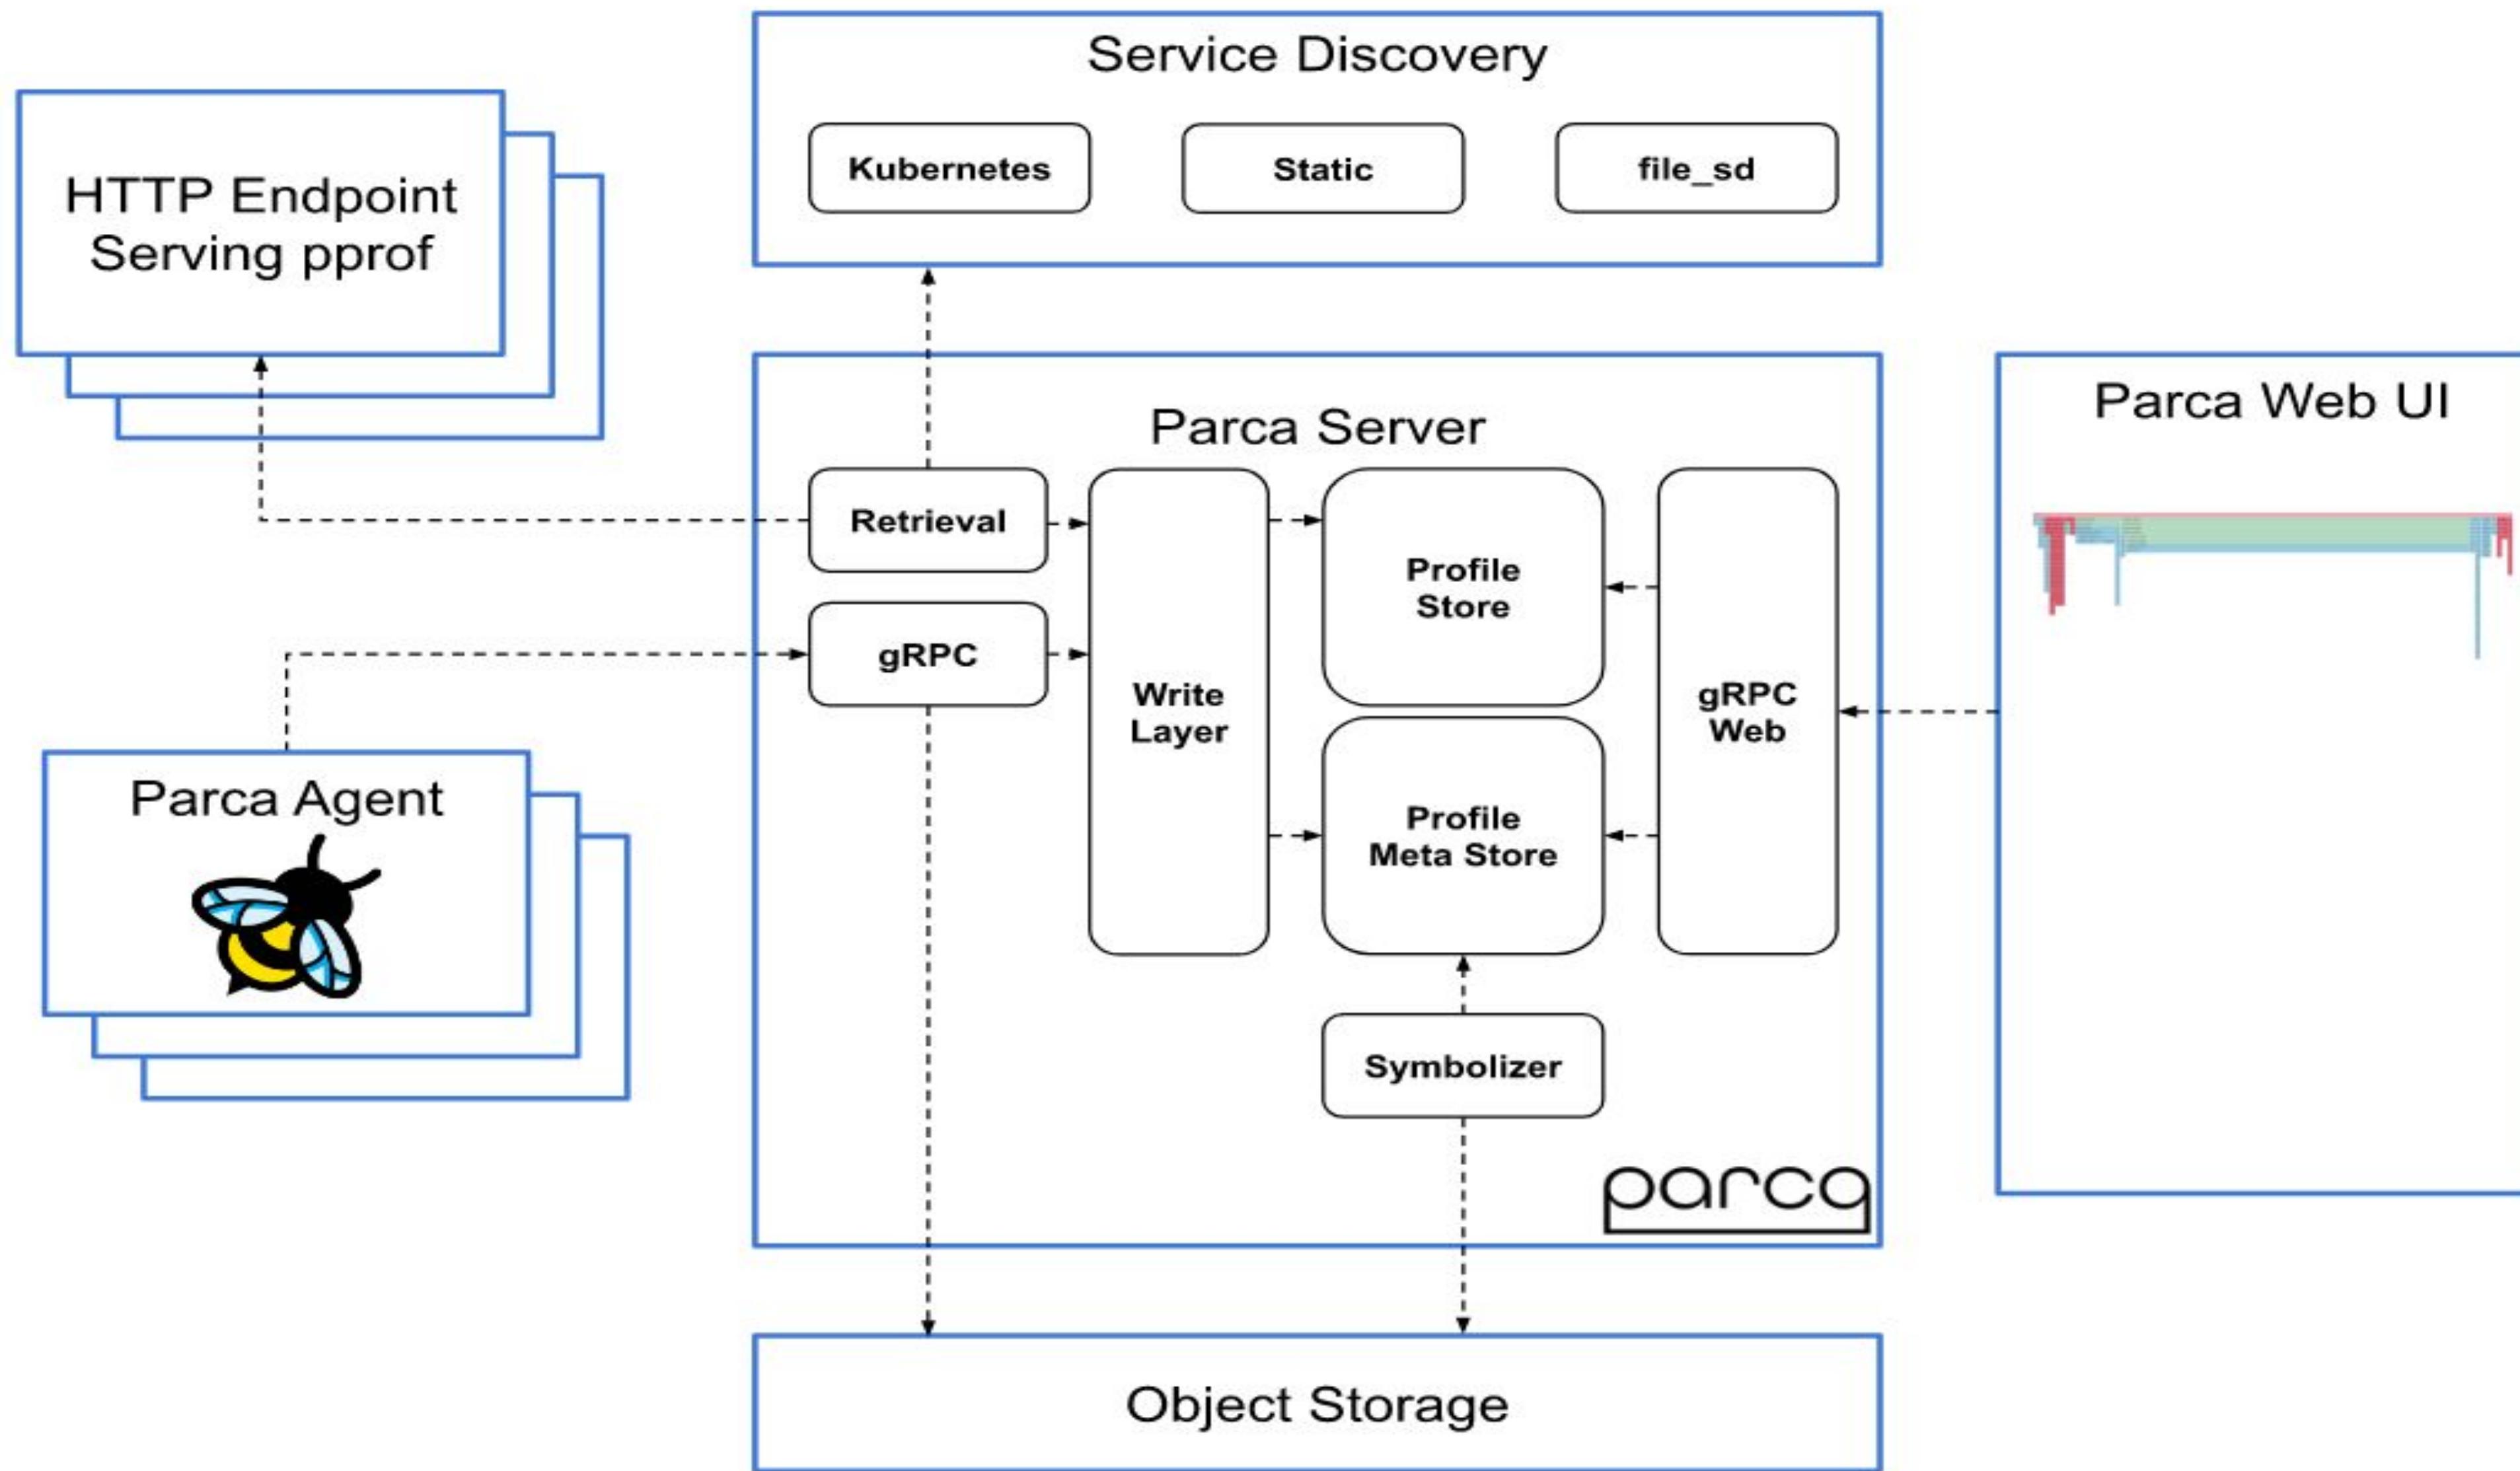


Profilers for the cloud native environment



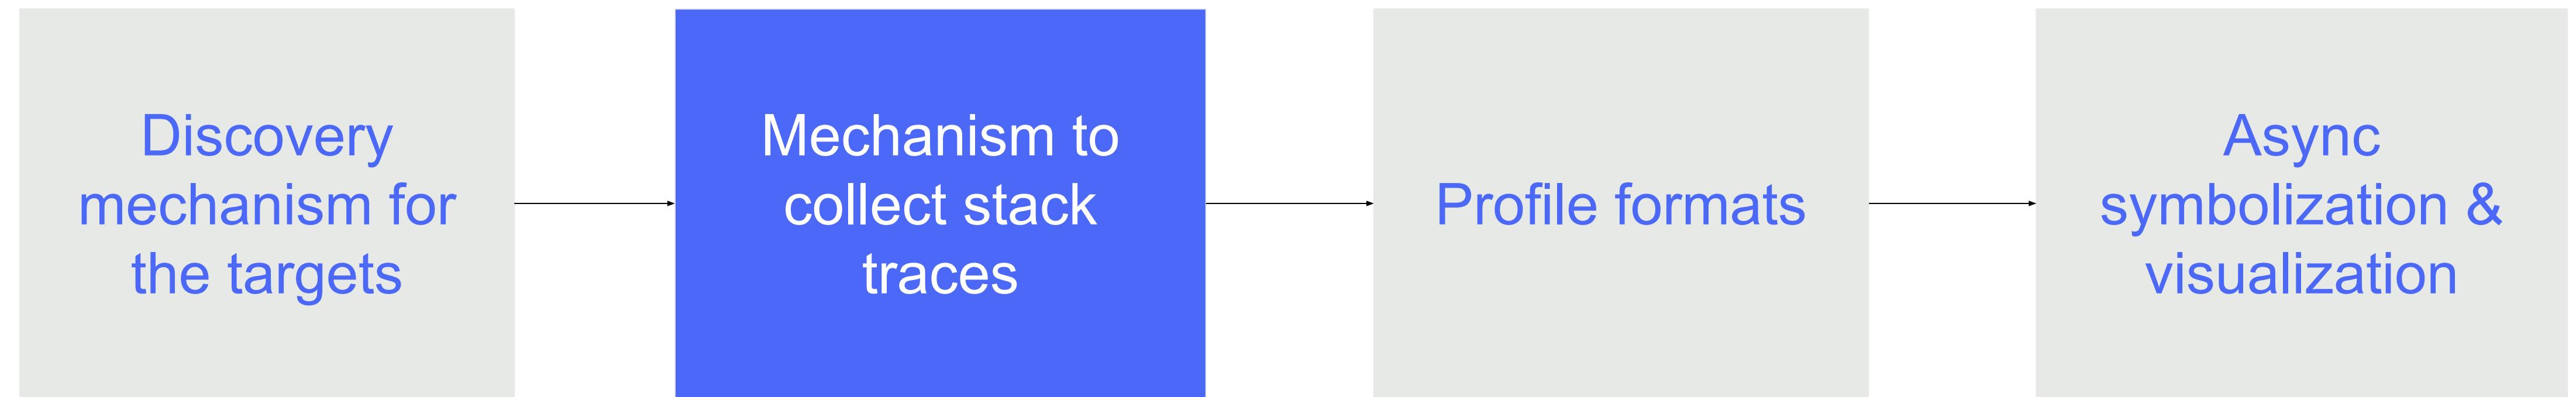
userpace

kernelspace



Userspace

Profilers for the cloud native environment



userpace

kernelspace

Userspace

- Discovery of targets
 - pods
 - binaries
- Debuginfo extraction
 - can be uploaded separately

Userspace

- Discovery of targets
- Debuginfo extraction
- Extract Unwind information from DWARF
 - **Executable Linkable format - ELF:** For obj file, executable program, shared object etc
 - **DWARF - widely used debugging format:** CIE - Common Information Entry
 - **Tools to read ELF and/or DWARF information:** readelf, objdump, elfutils, llvm-dwarfdump

```
// UnwindTableRow represents a single row in the unwind table.  
// x86_64: rip (instruction pointer register), rsp (stack pointer register), rbp (base pointer/frame pointer register)  
// aarch64: lr(link register), sp(stack pointer register), fp(frame pointer register)  
Francisco Javier Honduvilla Coto, 13 months ago | 2 authors (Francisco Javier Honduvilla Coto and others)  
type UnwindTableRow struct {  
    // The address of the machine instruction.  
    // Each row covers a range of machine instruction, from its address (Loc) to that of the row below.  
    Loc uint64  
    // CFA, the value of the stack pointer in the previous frame.  
    CFA frame.DWRule  
    // The value of the RBP register.  
    RBP frame.DWRule  
    // The value of the saved return address. This is not needed in x86_64 as it's part of the ABI but is necessary  
    // in arm64.  
    RA frame.DWRule  
}
```

Userspace: DWARF Unwind Tables

```
// UnwindTableRow represents a single row in the unwind table.  
// x86_64: rip (instruction pointer register), rsp (stack pointer register), rbp (base pointer/frame pointer register)  
// aarch64: lr(link register), sp(stack pointer register), fp(frame pointer register)  
Francisco Javier Honduvilla Coto, 13 months ago | 2 authors (Francisco Javier Honduvilla Coto and others)  
type UnwindTableRow struct {  
    // The address of the machine instruction.  
    // Each row covers a range of machine instruction, from its address (Loc) to that of the row below.  
    Loc uint64  
    // CFA, the value of the stack pointer in the previous frame.  
    CFA frame.DWRule  
    // The value of the RBP register.  
    RBP frame.DWRule  
    // The value of the saved return address. This is not needed in x86_64 as it's part of the ABI but is necessary  
    // in arm64.  
    RA frame.DWRule  
}
```

✓ // CompactUnwindTableRows encodes unwind information using 2x 64 bit words.
// `lrOffset` is the link register for arm64; it is initialized to 0 for x86.

You, 3 months ago | 2 authors (You and others)

✓ type CompactUnwindTableRow struct {
 pc uint64
 lrOffset int16
 cfaType uint8
 rbpType uint8
 cfaOffset int16
 rbpOffset int16
}

compact
unwind table

Userspace

- DWARF information is included in (ELF) binaries
- Extract Unwind tables from DWARF info

```
=> Function start: 293c0, Function end: 29427
pc: 293c0 cfa_type: 2 rbp_type: 0 cfa_offset: 8 rbp_offset: 0
pc: 293c5 cfa_type: 2 rbp_type: 0 cfa_offset: 16 rbp_offset: 0
pc: 293cb cfa_type: 2 rbp_type: 0 cfa_offset: 32 rbp_offset: 0
pc: 29408 cfa_type: 2 rbp_type: 0 cfa_offset: 16 rbp_offset: 0
pc: 29409 cfa_type: 2 rbp_type: 0 cfa_offset: 8 rbp_offset: 0
pc: 2940e cfa_type: 2 rbp_type: 0 cfa_offset: 32 rbp_offset: 0

compact unwind table snippet for `libc` on x86
```

```
=> Function start: 293c0, Function end: 29427
```

```
pc: 293c0 cfa_type: 2 rbp_type: 0 cfa_offset: 8 rbp_offset: 0  
pc: 293c5 cfa_type: 2 rbp_type: 0 cfa_offset: 16 rbp_offset: 0  
pc: 293cb cfa_type: 2 rbp_type: 0 cfa_offset: 32 rbp_offset: 0  
pc: 29408 cfa_type: 2 rbp_type: 0 cfa_offset: 16 rbp_offset: 0  
pc: 29409 cfa_type: 2 rbp_type: 0 cfa_offset: 8 rbp_offset: 0  
pc: 2940e cfa_type: 2 rbp_type: 0 cfa_offset: 32 rbp_offset: 0
```

x86

```
compact unwind table snippet for `libc` on x86
```

Aarch64

```
=> Function start: 26c00, Function end: 26c80
```

```
pc: 26c00 cfa_type: 2 rbp_type: 0 cfa_offset: 0 rbp_offset: 0 lr_offset: 0  
pc: 26c04 cfa_type: 2 rbp_type: 1 cfa_offset: 48 rbp_offset: -48 lr_offset: -  
pc: 26c14 cfa_type: 2 rbp_type: 1 cfa_offset: 48 rbp_offset: -48 lr_offset: -  
pc: 26c6c cfa_type: 2 rbp_type: 0 cfa_offset: 0 rbp_offset: 0 lr_offset: 0  
pc: 26c70 cfa_type: 2 rbp_type: 1 cfa_offset: 48 rbp_offset: -48 lr_offset: -
```

```
=> Function start: 26a00, Function end: 26a0c
```

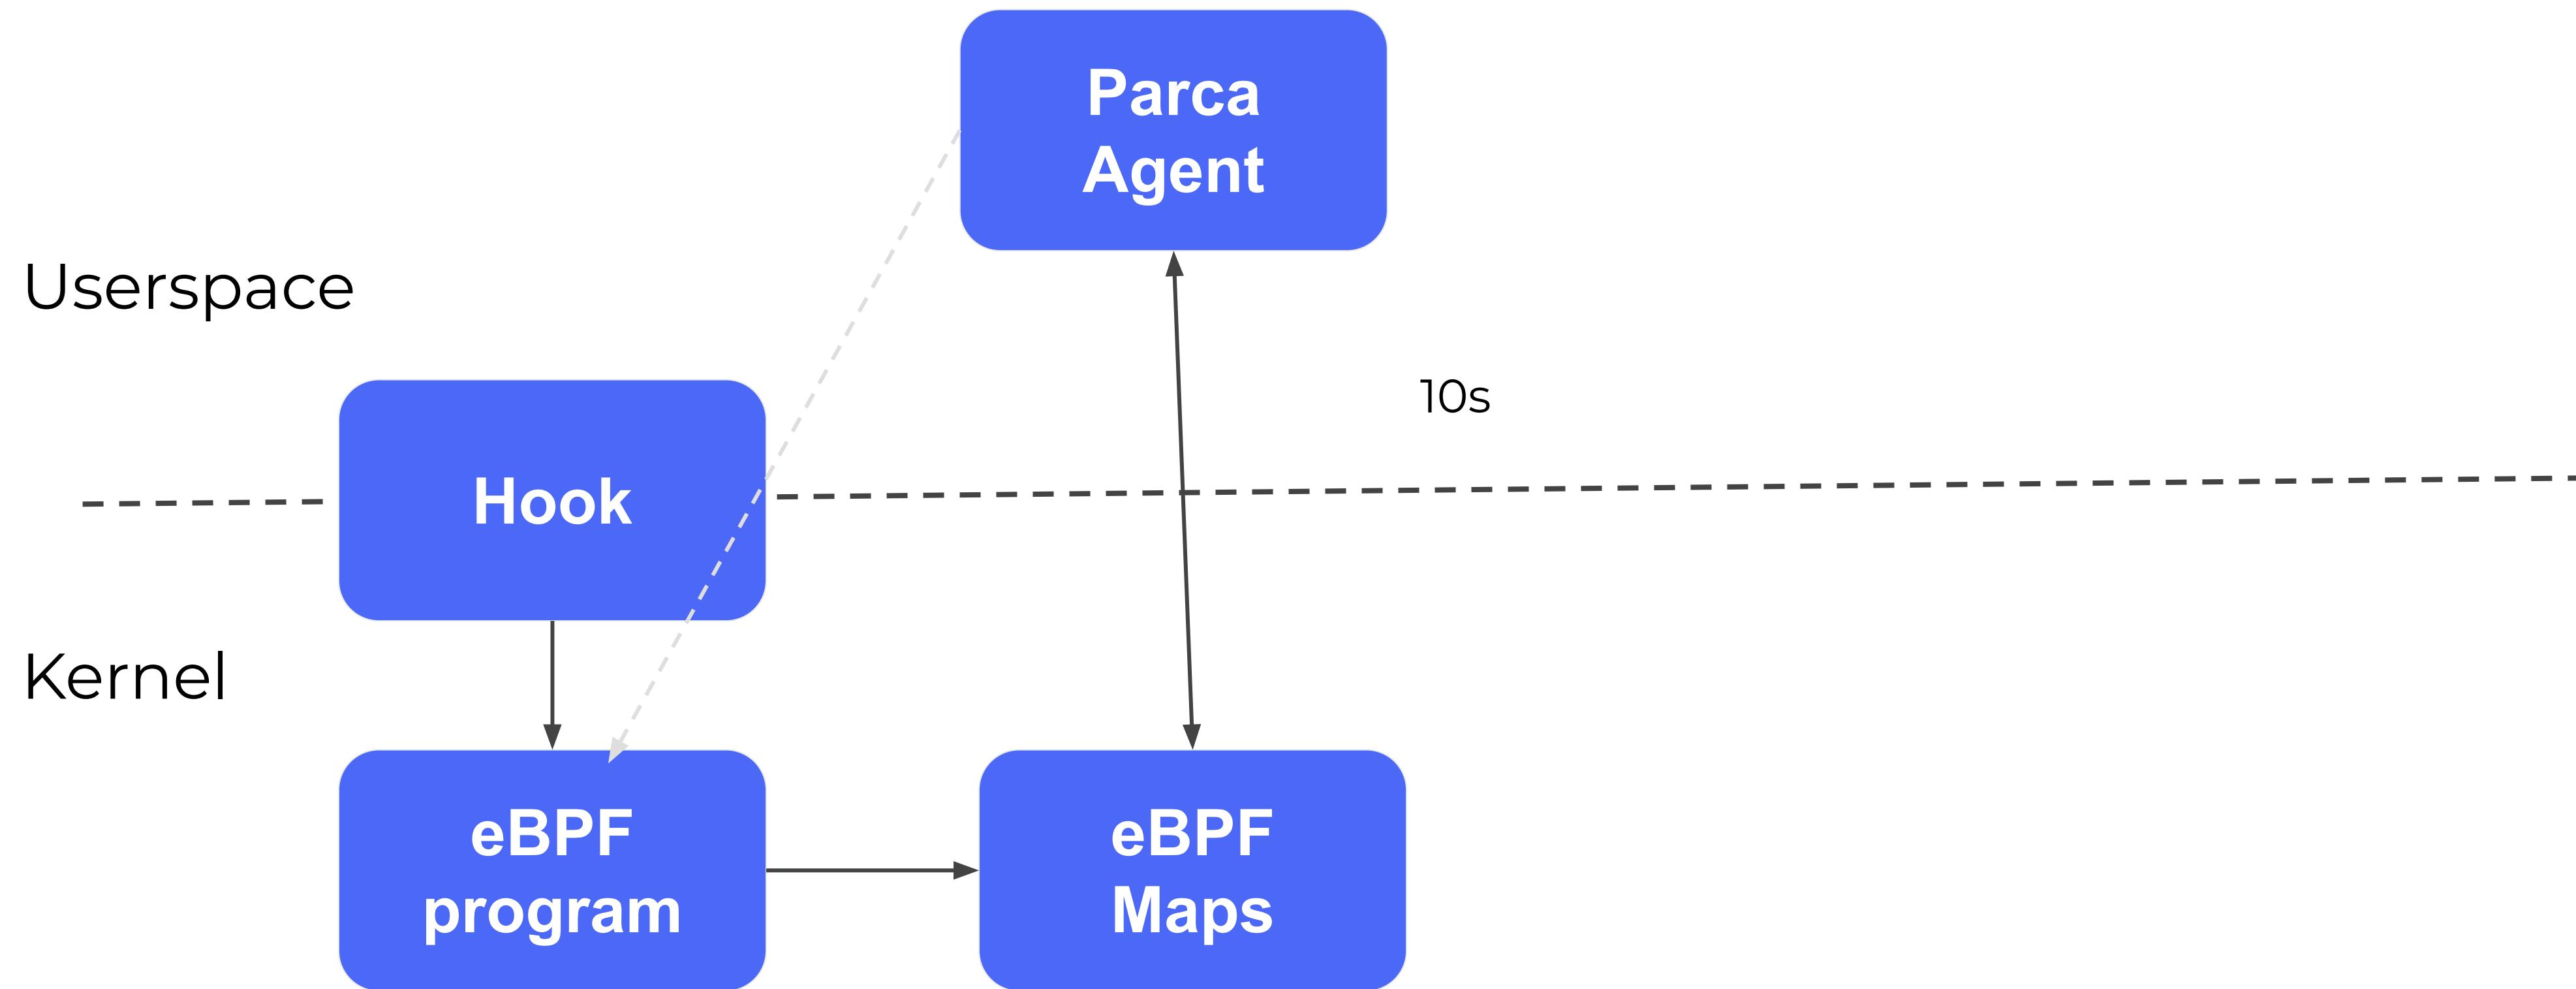
```
pc: 26a00 cfa_type: 2 rbp_type: 0 cfa_offset: 0 rbp_offset: 0 lr_offset: 0  
pc: 26a04 cfa_type: 2 rbp_type: 1 cfa_offset: 16 rbp_offset: -16 lr_offset: -
```

```
compact unwind table snippet for `libc` on Arm64
```

Polar Signals

Kernelspace

Communicating with Userspace

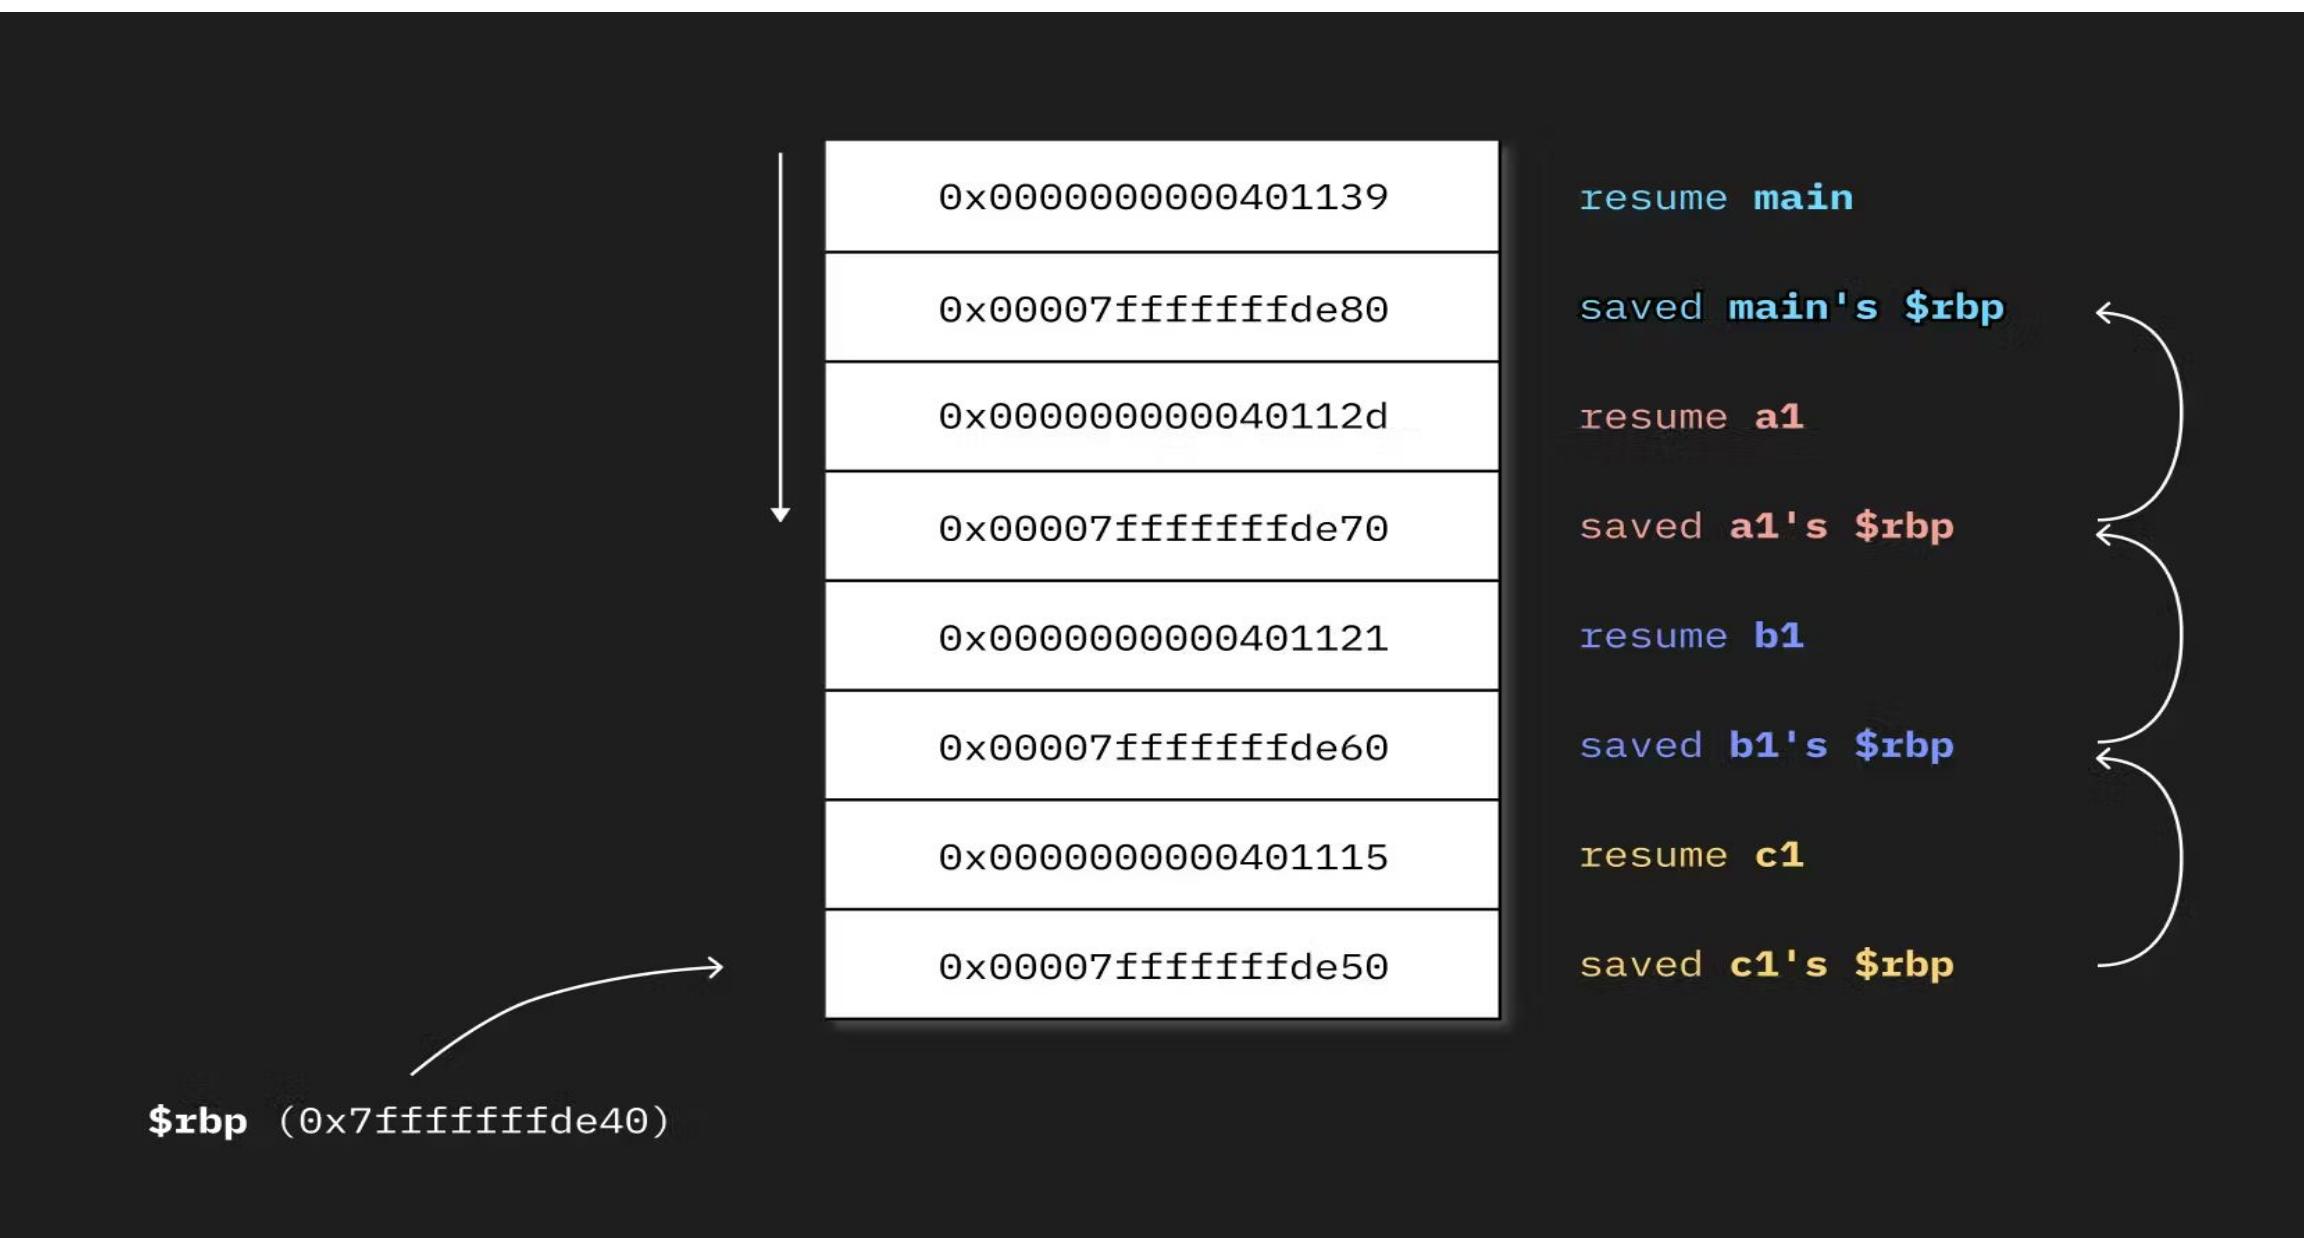
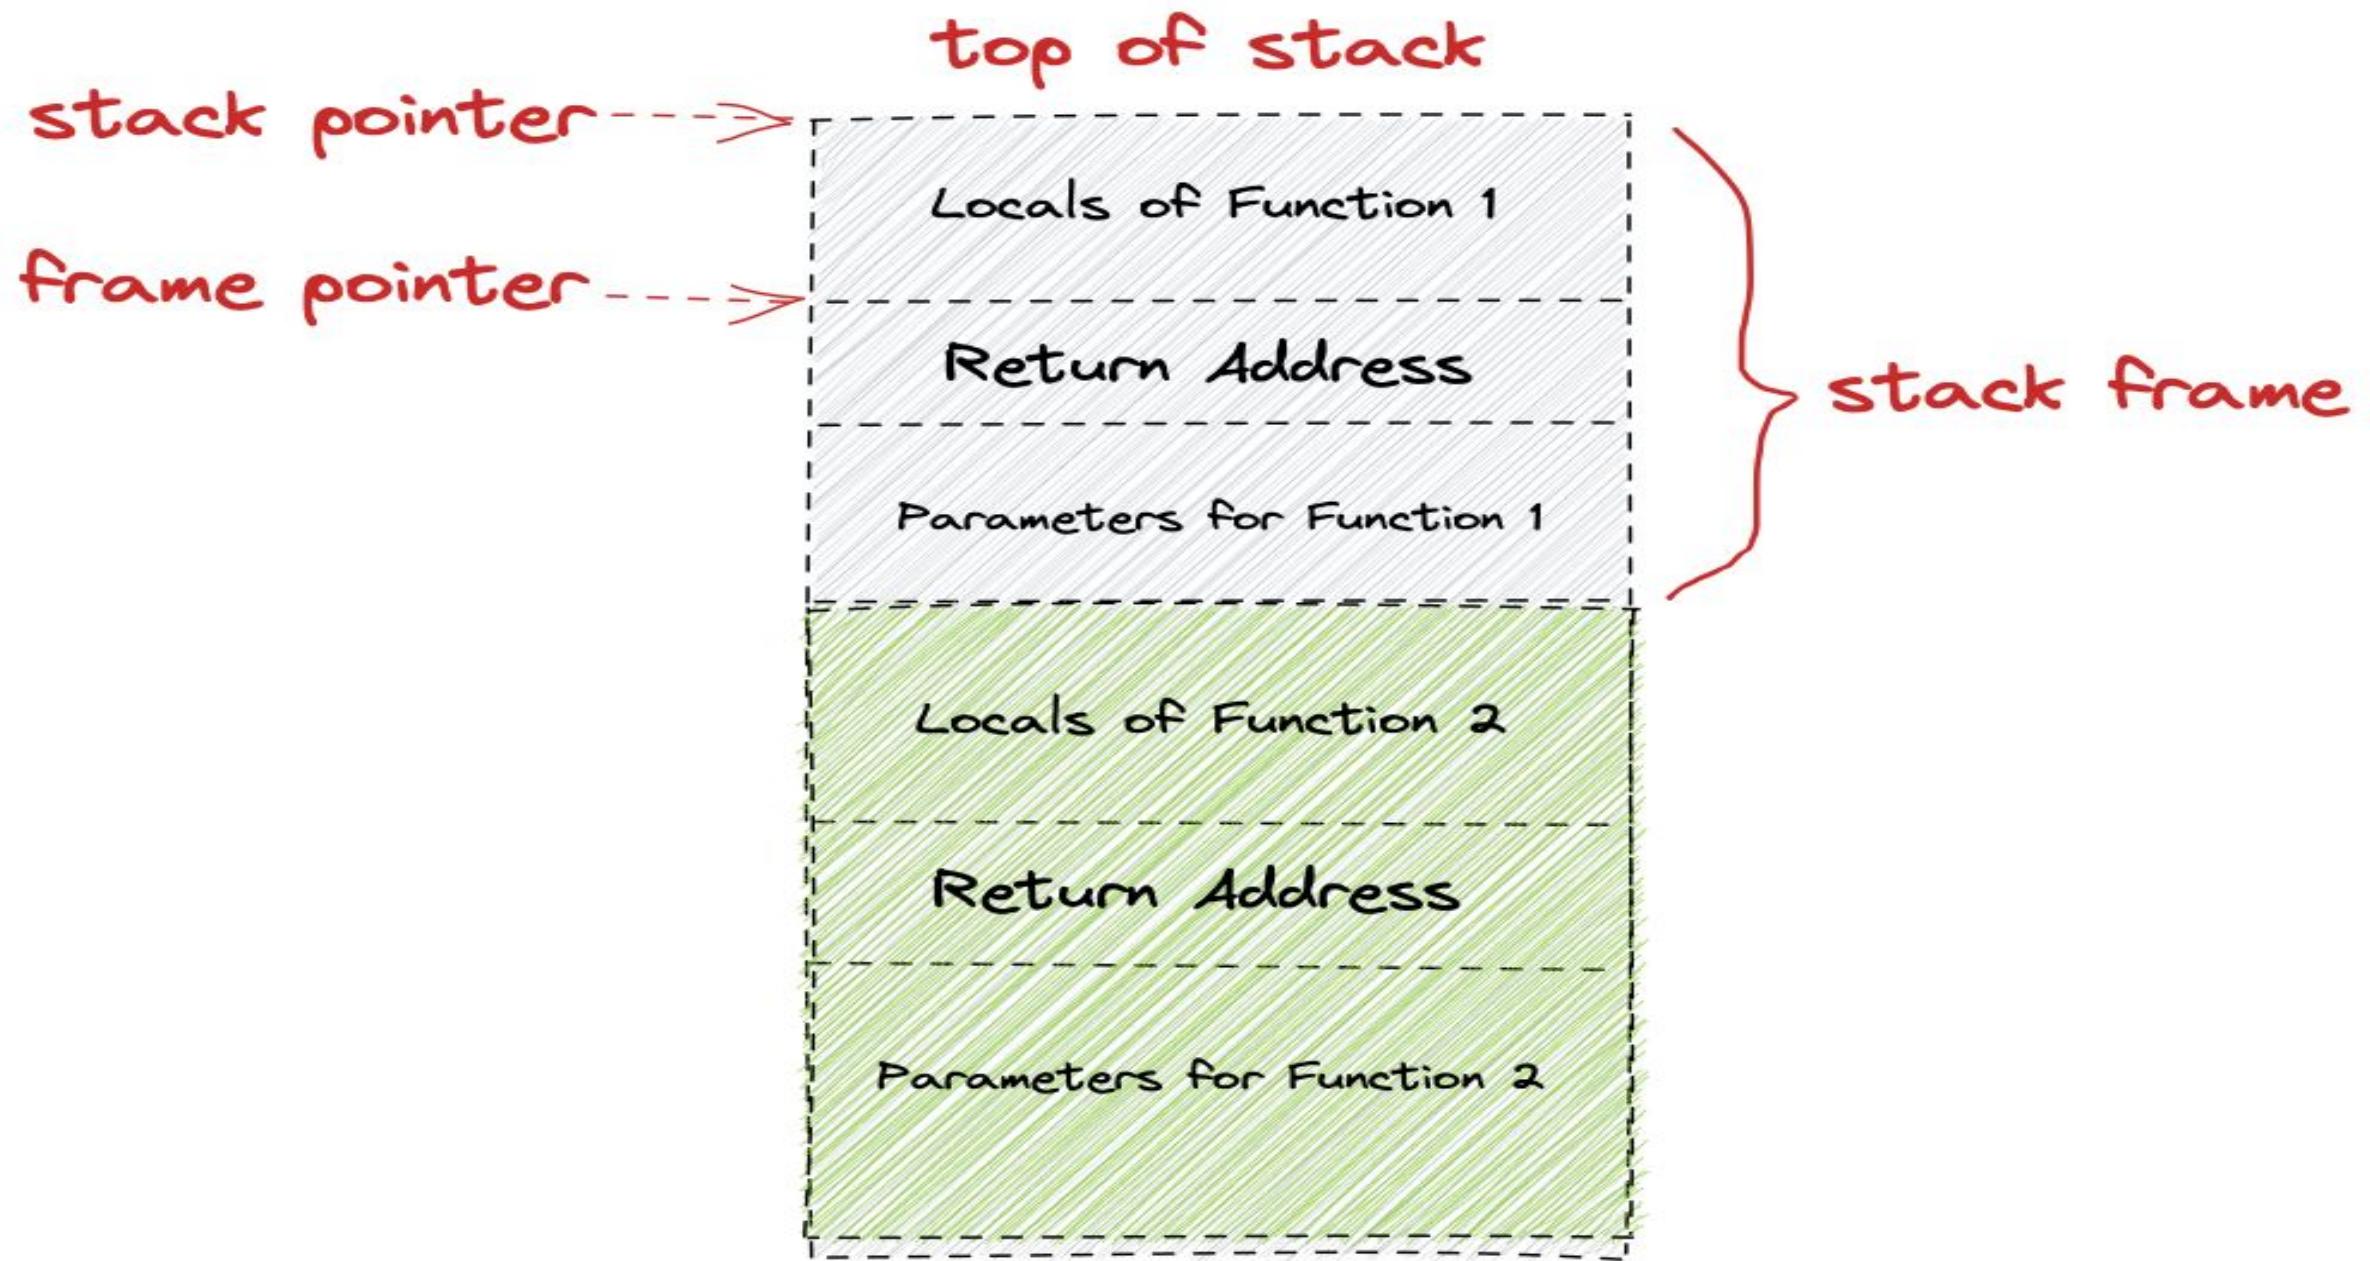


Communicating with Userspace

```
    }
    r
    pythonPIDToProcessInfo *libbpf.BPFMap
    pythonVersionSpecificOffsets *libbpf.BPFMap
    pythonVersionToOffsetIndex map[string]uint32

mapsH
if err != nil {
    r
    unwindShards *libbpf.BPFMap
}
programs *libbpf.BPFMap
m.proc.processInfo *libbpf.BPFMap
return
}

// Unwind stuff
processCache *processCache
mappingInfoMemory profiler.EfficientBuffer
// writeU
You
func (m *Maps) writeUnwindTableRow(rowSlice *profiler.EfficientBuffer, row unwind.CompactUnwindTableRow, arch elf.Machine) {
    // .pc
    rowSlice.PutUint64(row.Pc())
    if arch == elf.EM_AARCH64 {
        // .lr_offset
        rowSlice.PutInt16(row.LrOffset())
    }
    // .cfa_type
    rowSlice.PutUint8(row.CfaType())
    // .rbp_type
    rowSlice.PutUint8(row.RbpType())
    // .cfa_offset
    rowSlice.PutInt16(row.CfaOffset())
    // .rbp_offset
    rowSlice.PutInt16(row.RbpOffset())
}
```



Kernelspace:

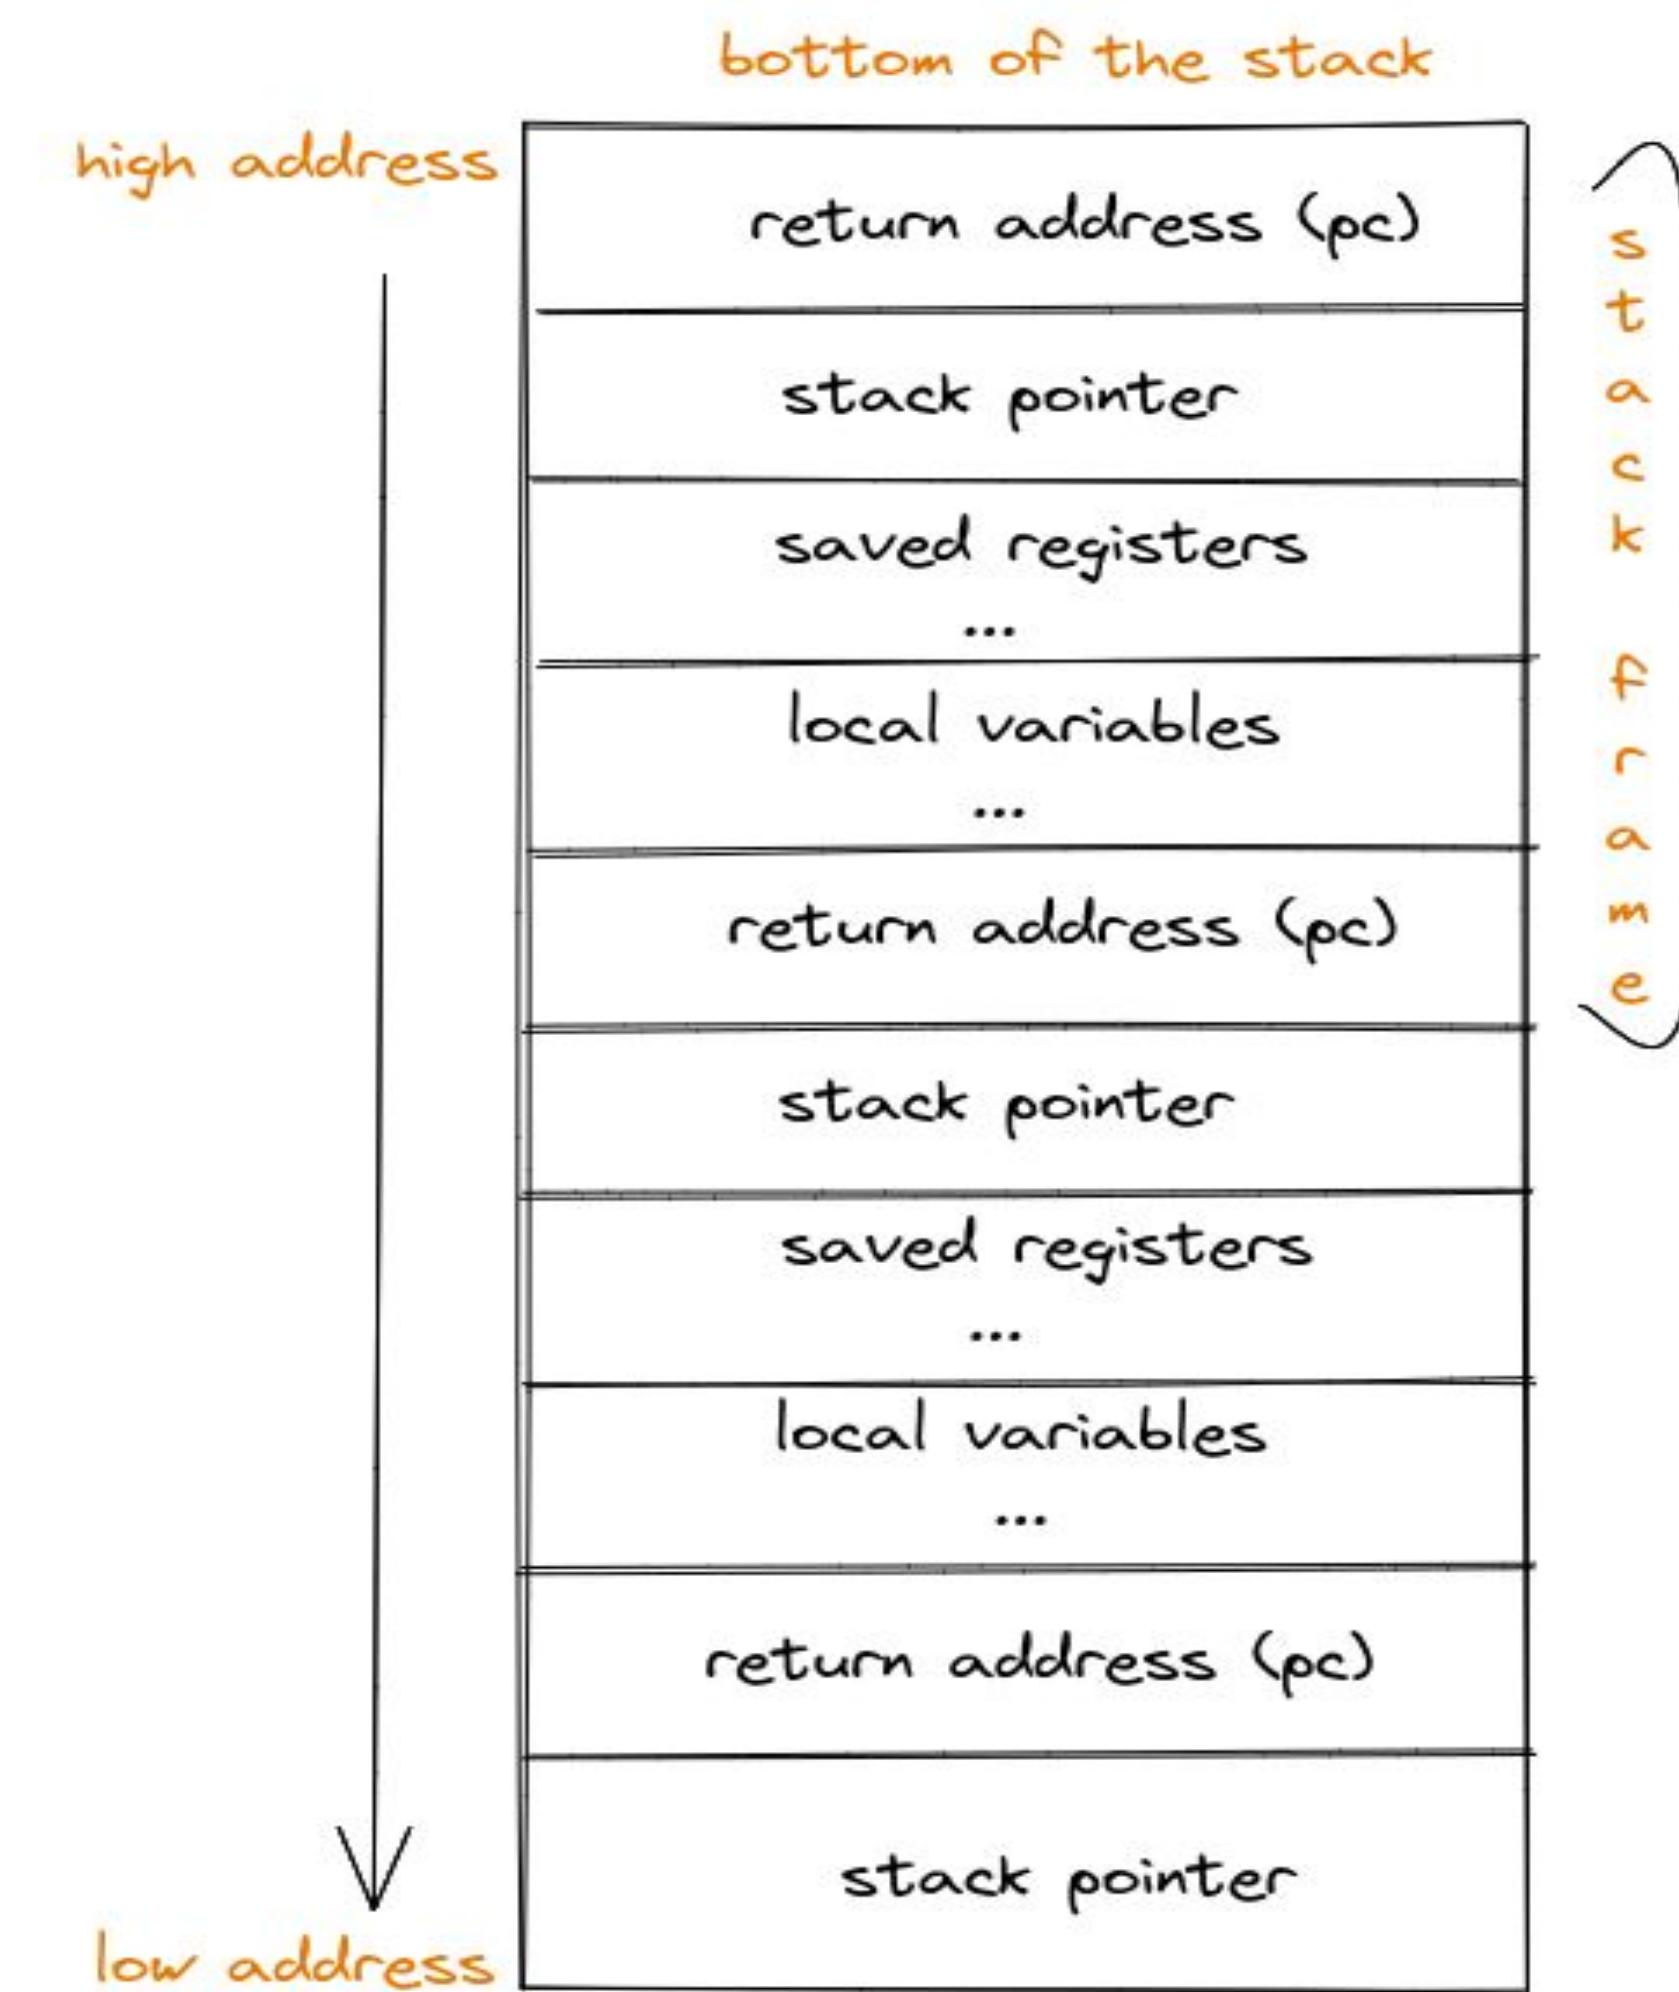
Walking the Stacks

Stacktraces

- What collecting stack traces involve
 - Kernel stacks
 - Application stacks
- Direction of stack growth
- So what are stack pointers, where do they come from

Figure 3.3: Stack Frame with Base Pointer

Position	Contents	Frame
$8n+16(%rbp)$	argument eightbyte n ...	Previous
$16(%rbp)$	argument eightbyte 0	
$8(%rbp)$	return address	Current
$0(%rbp)$	previous %rbp value	
$-8(%rbp)$	unspecified ...	
$0(%rsp)$	variable size	
$-128(%rsp)$	red zone	



From: x86_64 ABI specification

Userspace

=> Function start: 293c0, Function end: 29427

```
pc: 293c0 cfa_type: 2 rbp_type: 0 cfa_offset: 8 rbp_offset: 0
pc: 293c5 cfa_type: 2 rbp_type: 0 cfa_offset: 16 rbp_offset: 0
pc: 293cb cfa_type: 2 rbp_type: 0 cfa_offset: 32 rbp_offset: 0
pc: 29408 cfa_type: 2 rbp_type: 0 cfa_offset: 16 rbp_offset: 0
pc: 29409 cfa_type: 2 rbp_type: 0 cfa_offset: 8 rbp_offset: 0
pc: 2940e cfa_type: 2 rbp_type: 0 cfa_offset: 32 rbp_offset: 0
```

x86

compact unwind table snippet for `libc` on x86

=> Function start: 26c00, Function end: 26c80

```
pc: 26c00 cfa_type: 2 rbp_type: 0 cfa_offset: 0 rbp_offset: 0 lr_offset: 0
pc: 26c04 cfa_type: 2 rbp_type: 1 cfa_offset: 48 rbp_offset: -48 lr_offset: -
pc: 26c14 cfa_type: 2 rbp_type: 1 cfa_offset: 48 rbp_offset: -48 lr_offset: -
pc: 26c6c cfa_type: 2 rbp_type: 0 cfa_offset: 0 rbp_offset: 0 lr_offset: 0
pc: 26c70 cfa_type: 2 rbp_type: 1 cfa_offset: 48 rbp_offset: -48 lr_offset: -
```

Aarch64

=> Function start: 26a00, Function end: 26a0c

```
pc: 26a00 cfa_type: 2 rbp_type: 0 cfa_offset: 0 rbp_offset: 0 lr_offset: 0
pc: 26a04 cfa_type: 2 rbp_type: 1 cfa_offset: 16 rbp_offset: -16 lr_offset: -
```

compact unwind table snippet for `libc` on Arm64

DWARF Unwinding in x86

```
pc: 293c0 cfa_type: 2 rbp_type: 0 cfa_offset: 8    rbp_offset: 0
pc: 293c5 cfa_type: 2 rbp_type: 0 cfa_offset: 16   rbp_offset: 0
pc: 293cb cfa_type: 2 rbp_type: 0 cfa_offset: 32   rbp_offset: 0
pc: 29408 cfa_type: 2 rbp_type: 0 cfa_offset: 16   rbp_offset: 0
pc: 29409 cfa_type: 2 rbp_type: 0 cfa_offset: 8    rbp_offset: 0
pc: 2940e cfa_type: 2 rbp_type: 0 cfa_offset: 32   rbp_offset: 0
```

compact unwind table snippet for `libc` on x86

```
// HACK(javierhonduco): This is an architectural shortcut we can take. As we
// only support x86_64 at the minute, we can assume that the return address
// is *always* 8 bytes ahead of the previous stack pointer.
#ifndef __TARGET_ARCH_x86
    u64 previous_rip_addr = previous_rsp - 8;
    int err = bpf_probe_read_user(&previous_rip, 8, (void *)(previous_rip_addr));
    if (err < 0) {
        LOG("\n[error] Failed to read previous rip with error: %d", err);
    }
    LOG("\tprevious ip: %llx (@ %llx)", previous_rip, previous_rip_addr);
#endif
```

Unwind table in x86



DWARF Unwinding in Aarch64

```
=> Function start: 26c00, Function end: 26c80
  pc: 26c00 cfa_type: 2 rbp_type: 0 cfa_offset: 0 rbp_offset: 0 lr_offset: 0
  pc: 26c04 cfa_type: 2 rbp_type: 1 cfa_offset: 48 rbp_offset: -48 lr_offset: -
  pc: 26c14 cfa_type: 2 rbp_type: 1 cfa_offset: 48 rbp_offset: -48 lr_offset: -
  pc: 26c6c cfa_type: 2 rbp_type: 0 cfa_offset: 0 rbp_offset: 0 lr_offset: 0
  pc: 26c70 cfa_type: 2 rbp_type: 1 cfa_offset: 48 rbp_offset: -48 lr_offset: -
=> Function start: 26a00, Function end: 26a0c
  pc: 26a00 cfa_type: 2 rbp_type: 0 cfa_offset: 0 rbp_offset: 0 lr_offset: 0
  pc: 26a04 cfa_type: 2 rbp_type: 1 cfa_offset: 16 rbp_offset: -16 lr_offset: -
  
compact unwind table snippet for `libc` on Arm64
```

Aarch64

```
#if __TARGET_ARCH_arm64
    // For the leaf frame, the saved pc/ip is always be stored in the link register i
    if (found_lr_offset == 0) {
        previous_rip = PT_REGS_RET(&ctx->regs);
    } else {
        u64 previous_rip_addr = previous_rsp + found_lr_offset;
        int err = bpf_probe_read_user(&previous_rip, 8, (void *) (previous_rip_addr));
        if (err < 0) {
            LOG("\n[error] Failed to read previous rip with error: %d", err);
        }
        LOG("\tprevious ip: %llx (@ %llx)", previous_rip, previous_rip_addr);
    }
#endif
```



Compilers and Runtimes

Parca Agent Language Support

Compiled Languages

- C
- C++
- Rust
- Go
- and more!

*With or without frame
pointers!*

JIT (Just in Time Compiled) Languages

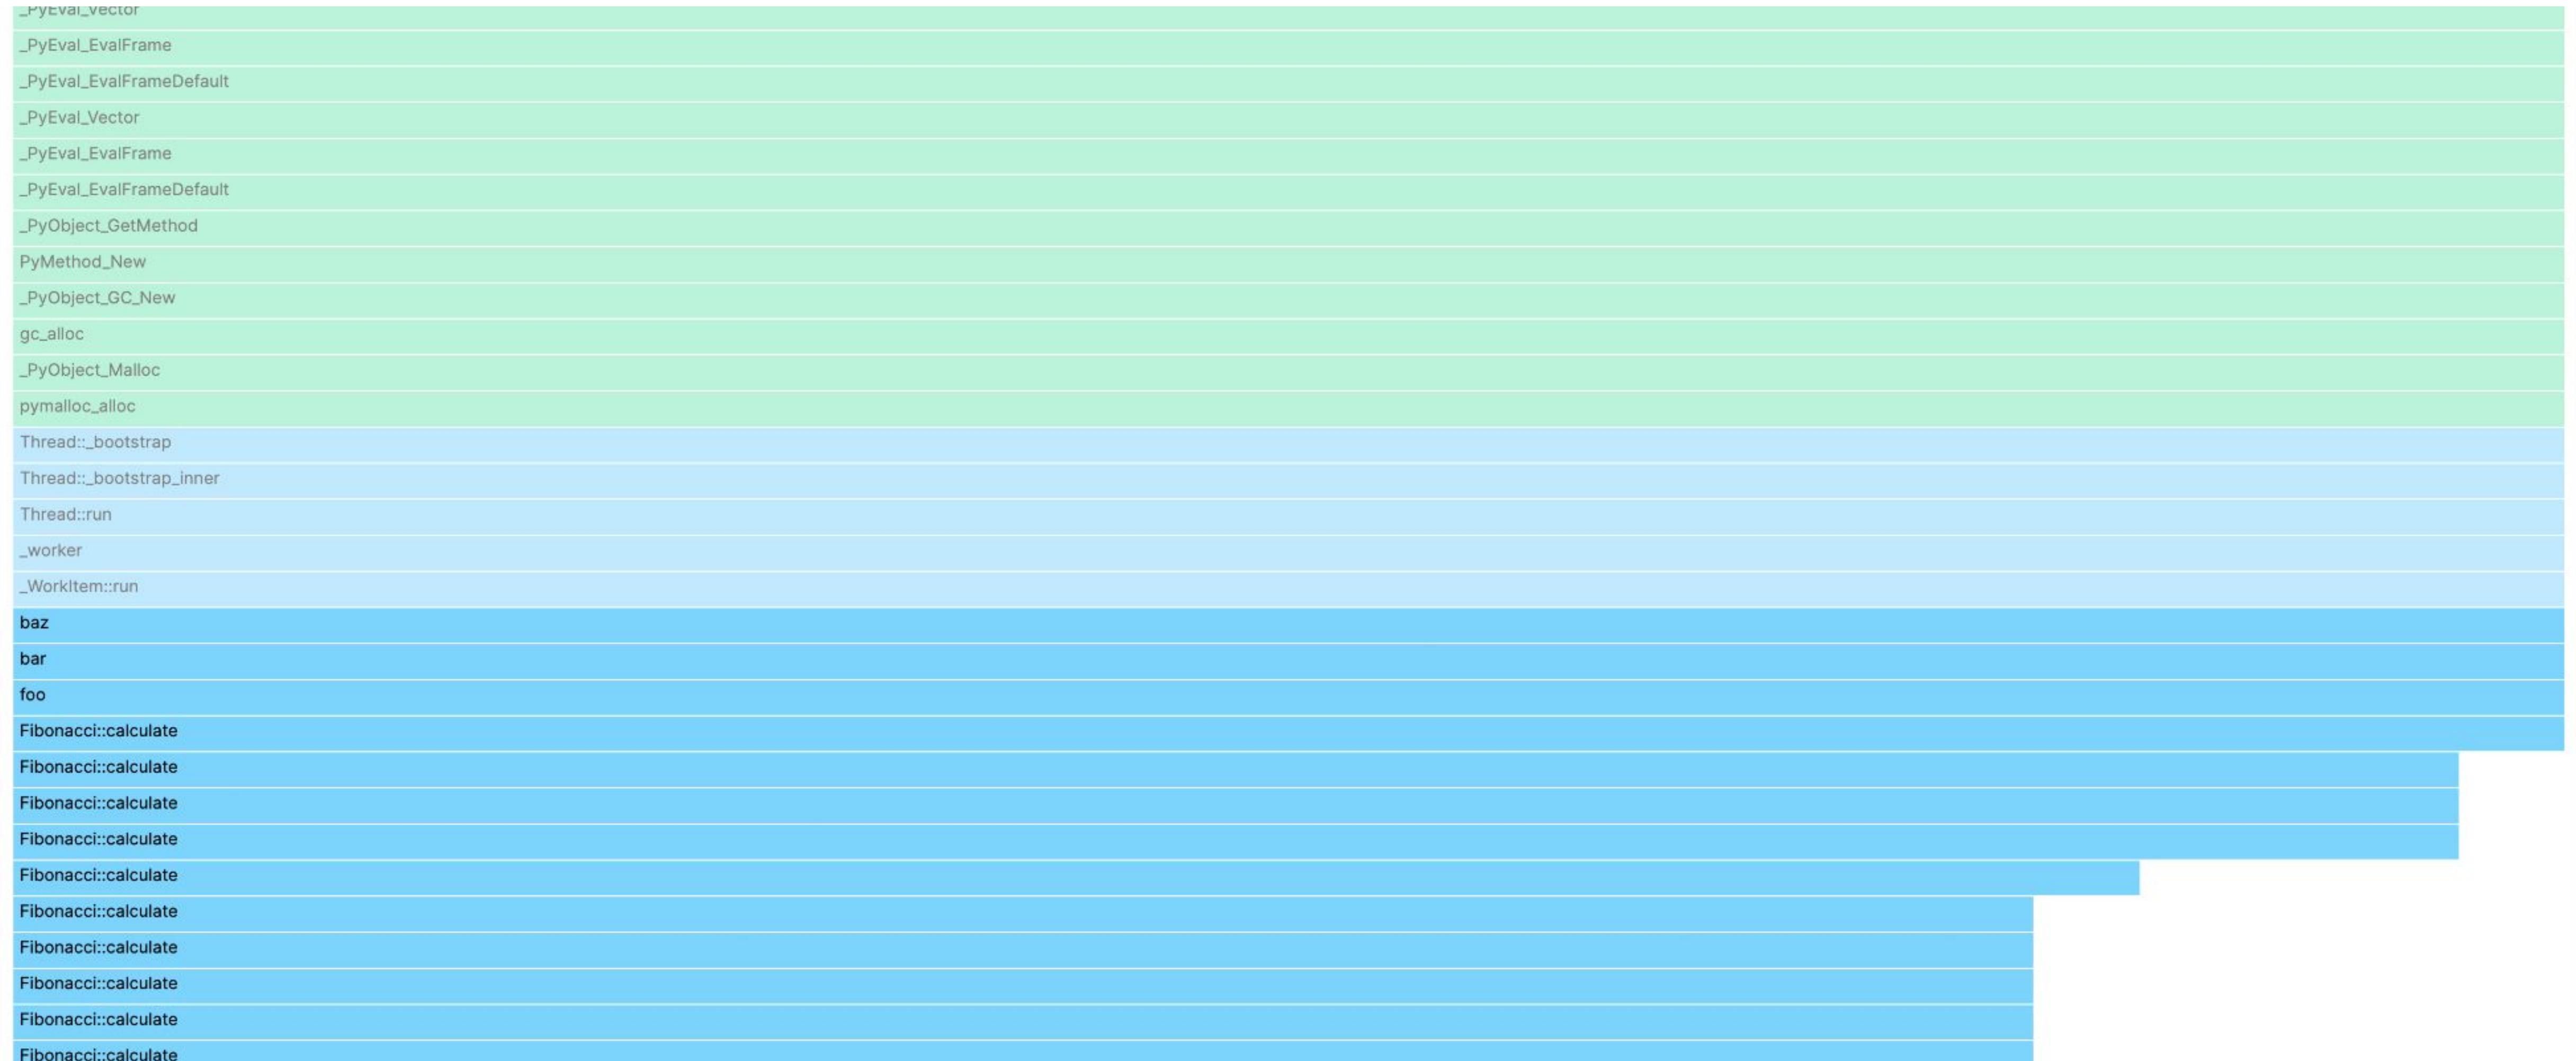
With Perf map or jitdump

- C#
- Erlang
- JVM(with async-profiler)
- Julia
- NodeJS

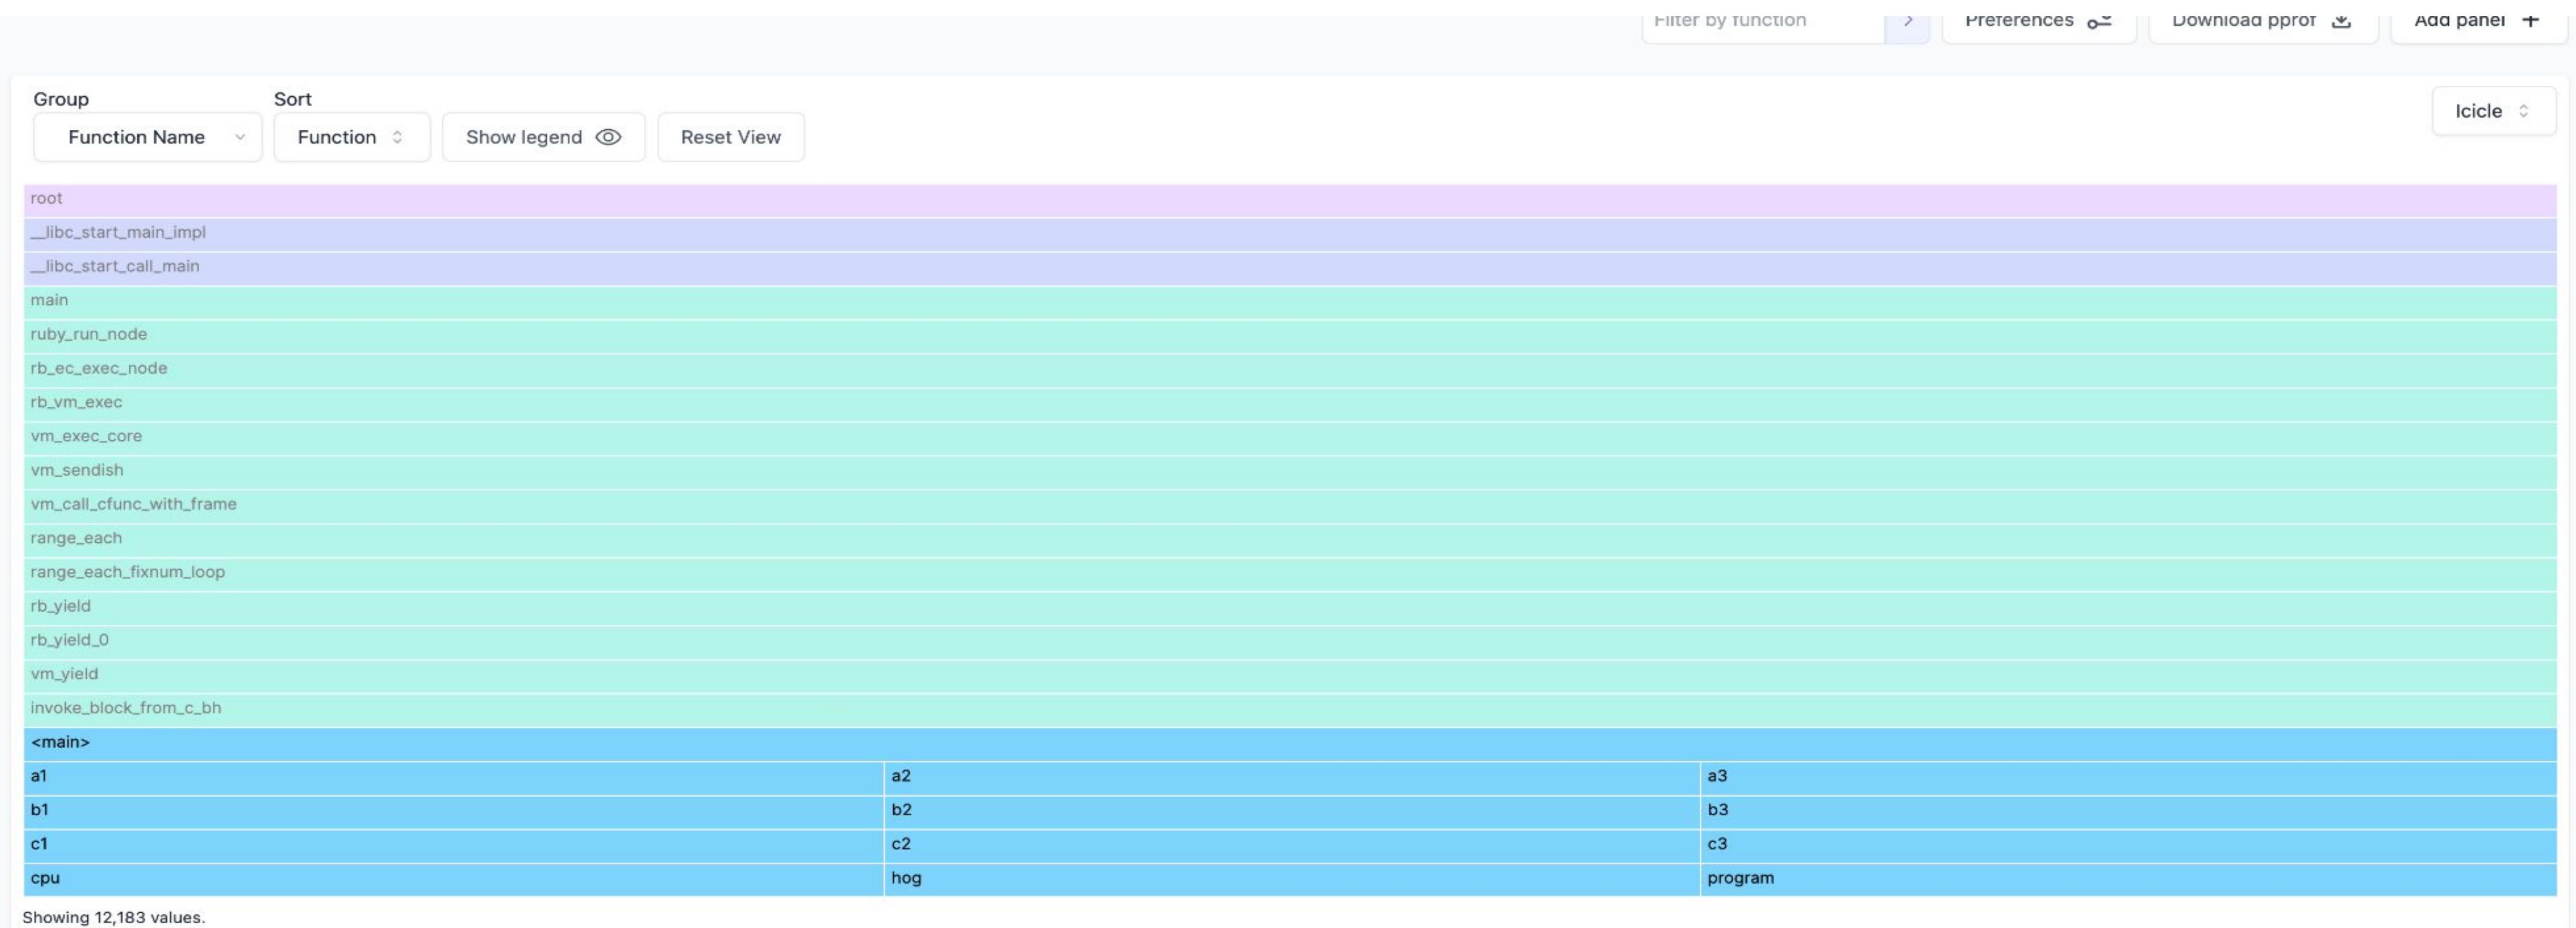
Interpreted Languages

- Python
- Ruby

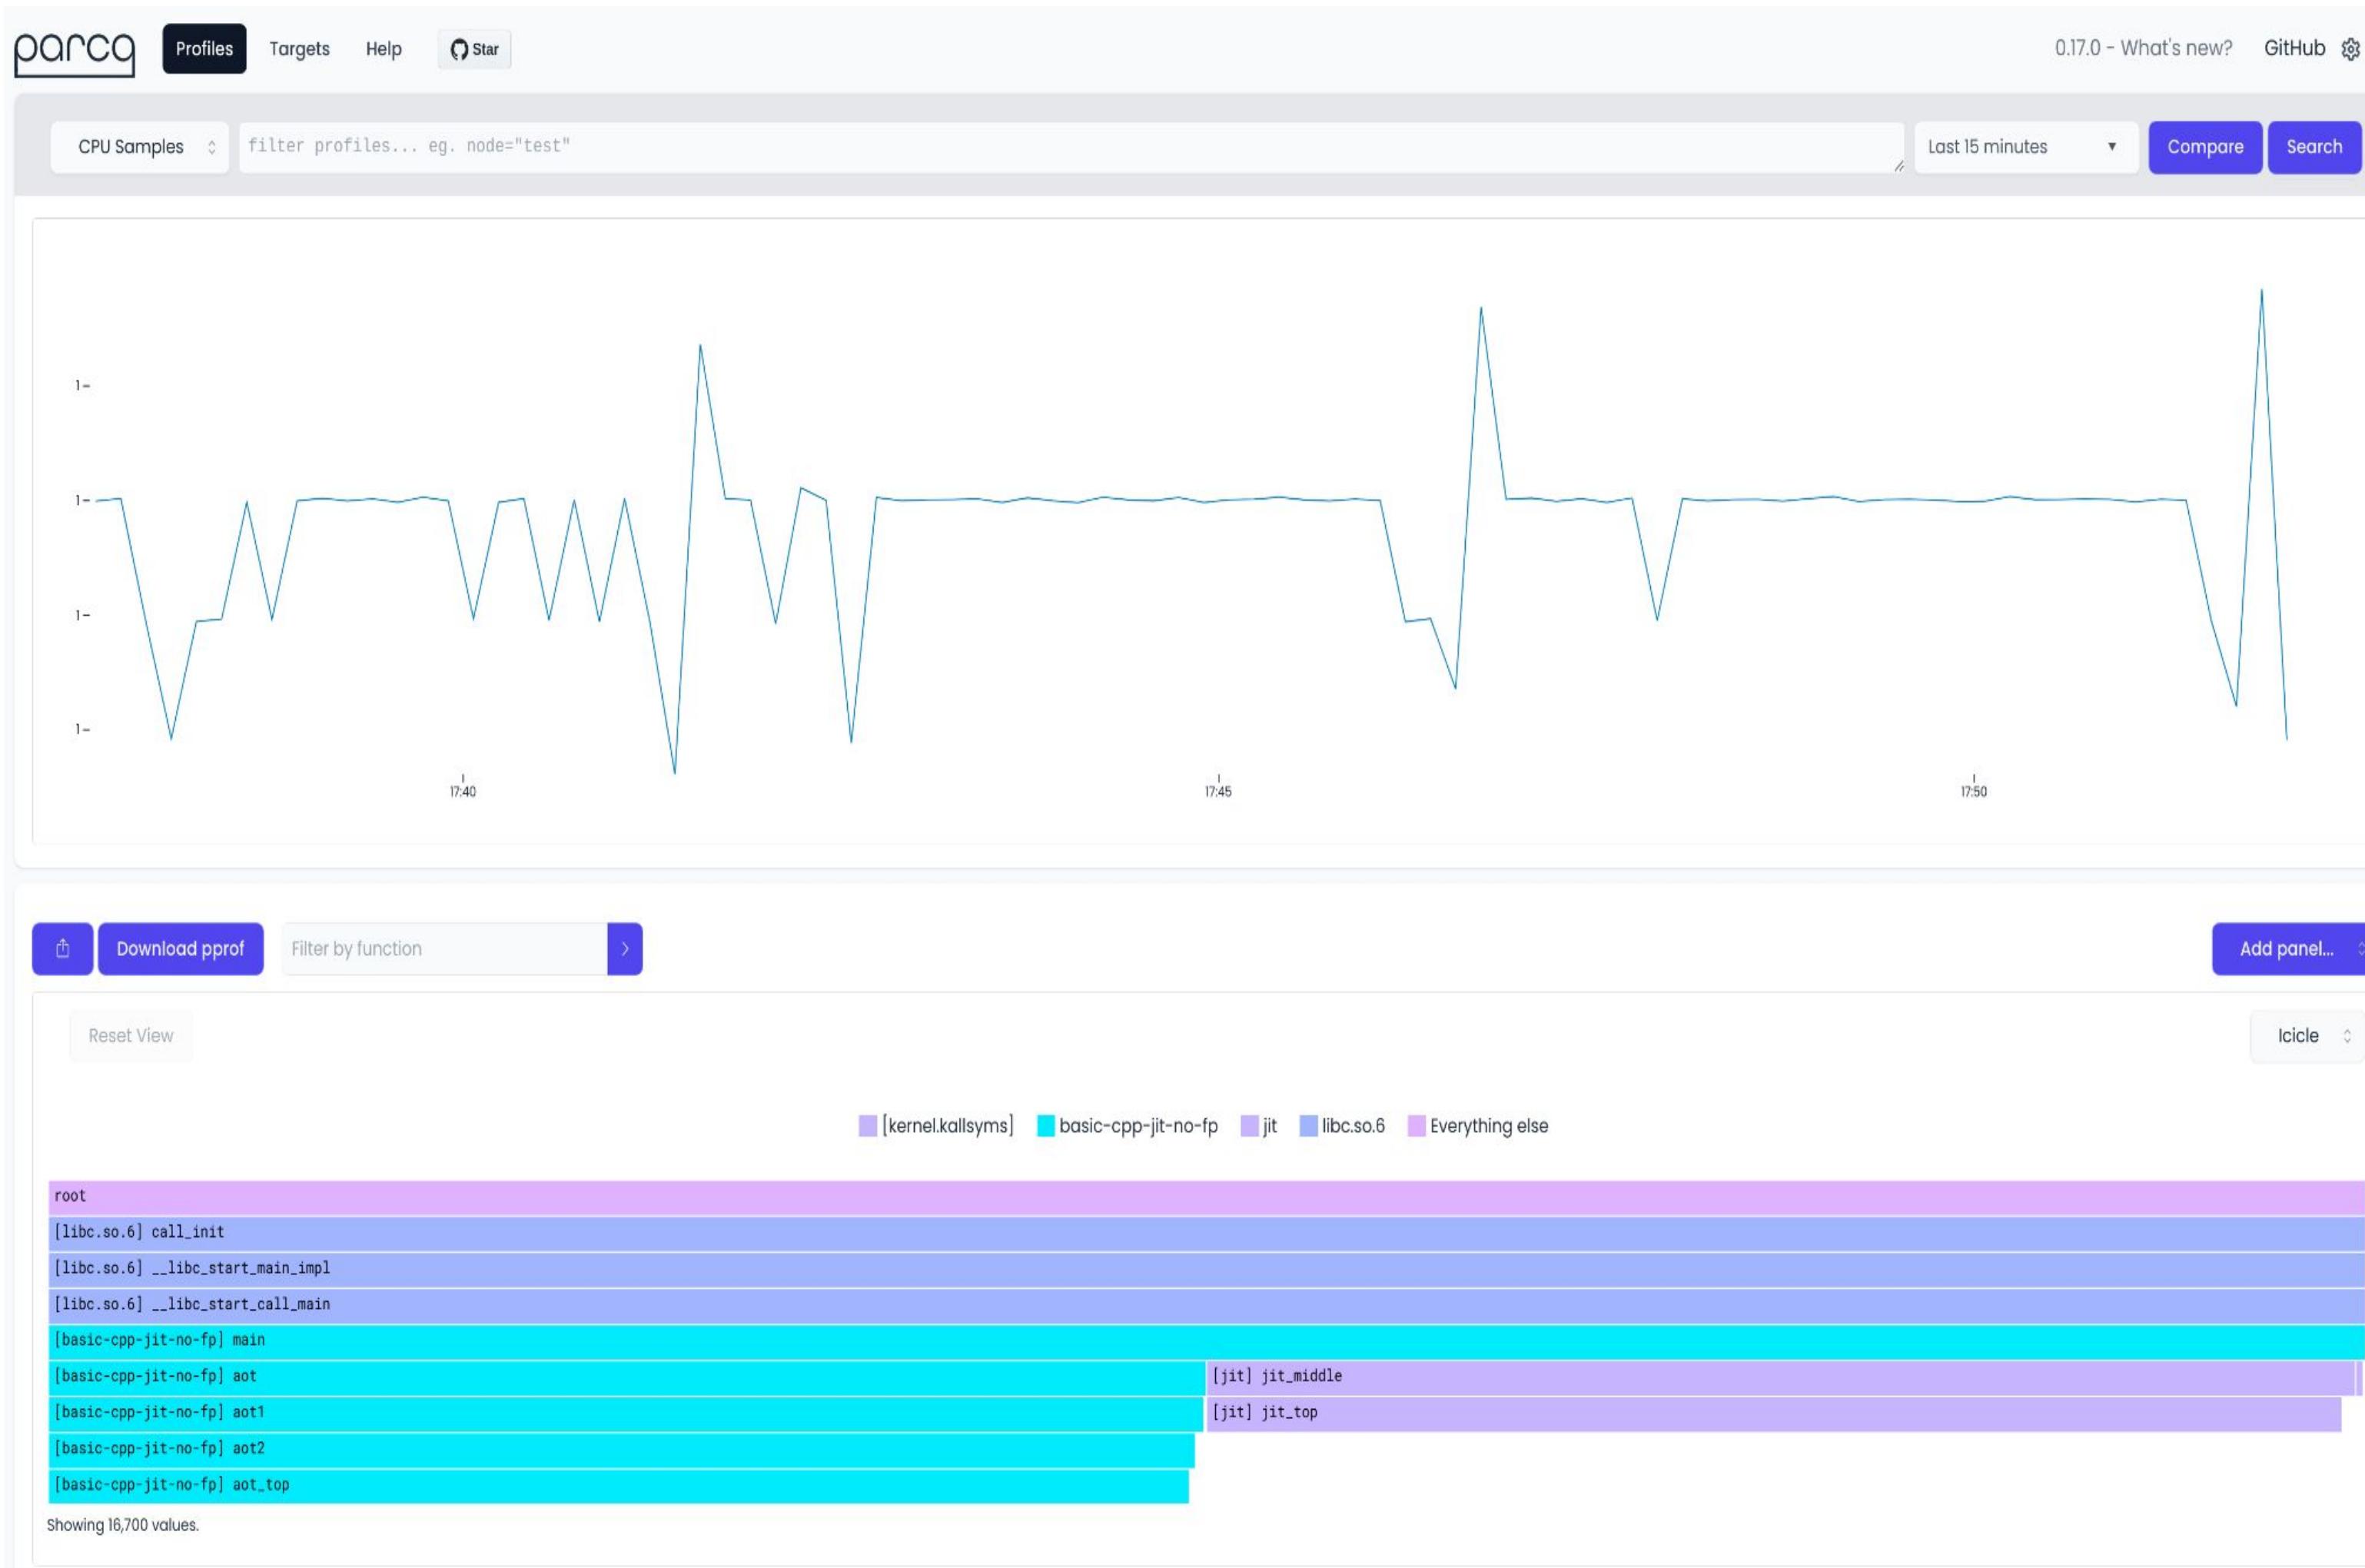
Interpreted Languages: Python



Interpreted Languages: Ruby



Parca Agent Language Support: JITed Runtimes and Stacks



```
// This implements a simple JIT for x86_64 with support for symbolization  
// with perfmap. We don't use any JIT framework or assembler for the sake  
// of simplicity.
```

```
// Some tools have heuristics to unwind the stack even  
// with frame pointers *and* unwind information omitted.  
#define ENABLE_FRAME_POINTERS_IN_JIT true  
// Amount of items to push into the stack to simulate a  
// more realistic stack usage. Otherwise stack unwinding with  
// frame pointers might work by accident.  
#define STACK_ITEMS 30
```

```
int __attribute__((noinline)) aot_top() {  
    for (int i = 0; i < 1000; i++) {  
    }  
  
    return 0;  
}  
  
// ahead of time  
int __attribute__((noinline)) aot2() { return aot_top(); }  
  
int __attribute__((noinline)) aot1() { return aot2(); }  
  
int __attribute__((noinline)) aot() { return aot1(); }  
  
void add_preamble(char **mem) {  
    if (!ENABLE_FRAME_POINTERS_IN_JIT) {  
        return;  
    }  
  
    *(mem)++ = 0x55; // push %rbp  
    *(mem)++ = 0x48; // mov %rsp,%rbp  
    *(mem)++ = 0x89;  
    *(mem)++ = 0xe5; // < difference between this and 0xec?
```

JITed Runtimes: Symbolising VSCode



- Bottom of the Stack in C++
- Correctly Profiling NodeJS
- Am I contained ?

How Compilers affect Icicle Graphs:

the cool edition

Clouds and Kernels

Kernels in the Cloud

- GKE by Google Cloud
- Graviton by AWS
- Custom kernels
- Custom cloud kernels

.... and others

**Linux Kernel v5.3+ with
BTF**

**Kernels versions
we support**

- Bad kernel bug causing a CPU lockup in Linux Kernel $\geq \text{v5.19}$ $\&\& < \text{v6.1}$
 - <https://github.com/parca-dev/parca-agent/issues/1675>
- Fix backported to ~stable for 6.1 and 6.3

Kernels versions
we prefer and
recommend

Linux Kernel $\geq \text{v6.4}$

BUT....

- Most OS Distros don't update as frequently as upstream
- Custom kernels don't backport fixes
- Cloud providers largely use v5.xx and <v6.2x kernels

Future Work

Future Roadmap

- Language Support
 - JVM
 - PHP
 - LuaJIT
 - Python
- Improve Memory Usage in the Agent
- More code coverage in Testing:
 - Kernel tests with QEMU
 - Integration tests to ensure correctness
 - Using ASAN and TSAN while developing
 - Snapshot testing of the unwind tables

- <https://www.polarsignals.com/blog/posts/2023/10/17/profiling-arm64-with-ebpf-in-parca-agent>
- <https://www.polarsignals.com/blog/posts/2023/10/04/profiling-python-and-ruby-with-ebpf>
- <https://www.polarsignals.com/blog/posts/2022/11/29/dwarf-based-stack-walking-using-ebpf>
- <https://www.polarsignals.com/blog/posts/2023/03/28/how-to-read-icicle-and-flame-graphs/>
- <https://www.polarsignals.com/blog/posts/2022/01/13/fantastic-symbols-and-where-to-find-them/>
- <https://demo.parca.dev>

Resources



Linux
Plumbers
Conference | Richmond, VA | Nov. 13-15, 2023



Thank you for listening!

[by Sumera | @sylfrena | @SumoOfShinovar](#)