BPF Programmable Netdevice

Daniel Borkmann (Isovalent)

LPC Networking & BPF Track 2023



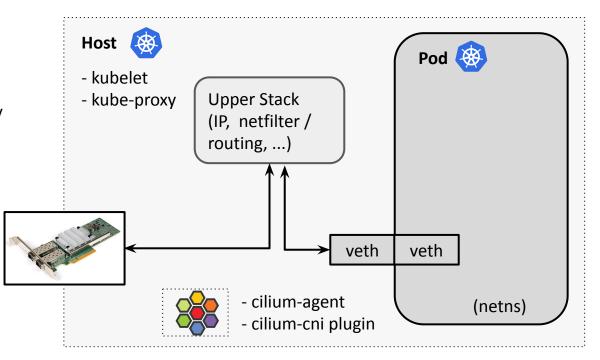
Goal for this talk: Can we achieve lowest possible networking overhead for K8s Pods?

Standard K8s Datapath Architecture:



Problems:

- kube-proxy scalability
- Routing via upper stack
- Potential reasons:
 - Cannot replace kube-proxy
 - Custom netfilter rules
 - Just "went with defaults"



Standard K8s Datapath Ard And And And And And And And Angle Angle And Angle An

Problems:

- kube-proxy scalability
- Routing via upper stack
- Potential reasons:
 - Cannot replace kube-proxy
 - Custom netfilter rules
 - Just "went with defaults"



:DOCKER-ISOLATION-STAGE-2 - [0:0] :DOCKER-USER - [0:0] :KUBE-EXTERNAL-SERVICES - [0:0] :KUBE-FIREWALL - [0:0] :KUBE-FORWARD - [0:0] :KUBE-SERVICES - [0:0]

-A FORWARD -m comment --comment "kubernetes forwarding rules" -j KUBE-FORWARD -A FORWARD -m conntrack --ctstate NEW -m comment --comment "kubernetes service portals" -j KUBE-SERVICES A FORWARD -i DOCKER-USER -A FORWARD -j DOCKER-ISOLATION-STAGE-1 -A FORWARD -o docker0 -m conntrack --ctstate RELATED, ESTABLISHED -j ACCEPT -A FORWARD -o docker0 -j DOCKER -A FORWARD -i docker0 ! -o docker0 -j ACCEP -A FORWARD -i docker0 -o docker0 -j ACCEPT -A FORWARD -o docker_gwbridge -m conntrack --ctstate RELATED,ESTABLISHED -j ACCEPT A FORWARD -o docker gwbridge -i DOCKER -A FORWARD -i docker gwbridge! -o docker gwbridge -i ACCEPT -A FORWARD -i docker_gwbridge -o docker_gwbridge -j DROP -A OUTPUT -m conntrack --ctstate NEW -m comment --comment "kubernetes service portals" -j KUBE-SERVICES -A OUTPUT - KUBE-FIREWALL, I. -A DOCKER-ISOLATION-STAGE-1 -i docker0 ! -o docker0 -j DOCKER-ISOLATION-STAGE-2 -A DOCKER-ISOLATION-STAGE-1 -i docker_gwbridge!-o docker_gwbridge-j DOCKER-ISOLATION-STAGE-2 -A DOCKER-ISOLATION-STAGE-1 - RETURN -A DOCKER-ISOLATION-STAGE-2 -o docker0 -j DROP -A DOCKER-ISOLATION-STAGE-2 -o docker_gwbridge -j DROF -A DOCKER-ISOLATION-STAGE-2 - j RETURN -A DOCKER-USER - RETURN -A KUBE-FIREWALL -m comment --comment "kubernetes firewall for dropping marked packets" -m mark --mark 0x8000/0x8000 -i DROP -A KUBE-FORWARD -m conntrack --ctstate INVALID -j DROP -A KUBE-FORWARD -m comment --comment "kubernetes forwarding rules" -m mark --mark 0x4000/0x4000 -i ACCEPT -A KUBE-FORWARD -s 10.217.0.0/16 -m comment --comment "kubernetes forwarding conntrack pod source rule" -m conntrack --ctstate RELATED,ESTABLISHEI -A KUBE-FORWARD -d 10.217.0.0/16 -m comment -comment "kubernetes forwarding conntrack pod destination rule" -m conntrack --ctstate RELATED, ESTABL -A KUBE-SERVICES -d 10.99.38.155/32 -p tcp -m comment --comment "default/nginx-59: has no endpoints" -m tcp --dport 80 -j REJECT --reject-with icmp-port--A KUBE-SERVICES -d 10.96.61.252/32 -p tcp -m comment --comment "default/nginx-64: has no endpoints" -m tcp --dport 80 -j REJECT --reject-with icmp-port--A KUBE-SERVICES -d 10.104.166.10/32 -p tcp -m comment --comment "default/nginx-67: has no endpoints" -m tcp --dport 80 -j REJECT --reject-with icmp-por -A KUBE-SERVICES -d 10.98.85.41/32 -p tcp -m comment --comment "default/nginx-9: has no endpoints" -m tcp --dport 80 -j REJECT --reject-with icmp-port-un -A KUBE-SERVICES -d 10.97.138.144/32 -p tcp -m comment --comment "default/nginx-17: has no endpoints" -m tcp --dport 80 -i REJECT --reject-with icmp-por -A KUBE-SERVICES -d 10.106.49.80/32 -p tcp -m comment --comment "default/nginx-37: has no endpoints" -m tcp --dport 80 -j REJECT --reject-with icmp-port-

A KUBE-SERVICES -d 10.106.234.213/32 -p tcp -m comment --comment "default/nginx-88: has no endpoints" -m tcp --dport 80 -i REJECT --reject-with icmp-po -A KUBE-SERVICES -d 10.109.209.136/32 -p tcp -m comment --comment "default/nginx-33: has no endpoints" -m tcp --dport 80 -j REJECT --reject-with icmp-po -A KUBE-SERVICES -d 10.106.196.105/32 -p tcp -m comment --comment "default/nginx-49: has no endpoints" -m tcp --dport 80 -j REJECT --reject-with icmp-po -A KUBE-SERVICES -d 10.111.101.6/32 -p tcp -m comment -comment "default/nginx-53: has no endpoints" -m tcp --dport 80 -j REJECT --reject-with icmp-port--A KUBE-SERVICES -d 10.110.226.230/32 -p tcp -m comment --comment "default/nginx-79: has no endpoints" -m tcp --dport 80 -j REJECT --reject-with icmp-po -A KUBE-SERVICES -d 10.98.99.136/32 -p tcp -m comment --comment "default/nginx-6: has no endpoints" -m tcp --dport 80 -j REJECT --reject-with icmp-port-u -A KUBE-SERVICES -d 10.99.75.233/32 -p tcp -m comment --comment "default/nginx-7: has no endpoints" -m tcp --dport 80 -j REJECT --reject-with icmp-port-u -A KUBE-SERVICES -d 10.108.41.202/32 -p tcp -m comment -comment "default/nginx-14: has no endpoints" -m tcp --dport 80 -j REJECT --reject-with icmp-por -A KUBE-SERVICES -d 10.97.36.249/32 -p tcp -m comment -comment "default/nginx-99: has no endpoints" -m tcp --dport 80 -j REJECT --reject-with icmp-port--A KUBE-SERVICES -d 10.98.213.37/32 -p tcp -m comment -comment "default/nginx-77: has no endpoints" -m tcp --dport 80 -j REJECT --reject-with icmp-port--A KUBE-SERVICES -d 10.107.229.31/32 -p tcp -m comment --comment "default/ngipx-92; has no endpoints" -m tcp --dport 80 -i REJECT --reject-with icmp-por

Standard K8s Datapath Ard And And And And And And And Angle Angle And Angle An

Problems:

- kube-proxy scalability
- Routing via upper stack
- Potential reasons:
 - Cannot replace kube-proxy
 - Custom netfilter rules
 - Just "went with defaults"

Host

- kubelet
- kube-proxy

KubeCon '23:

K8s sig-networking community is now starting to move this to nft since RHEL10 will not ship iptables anymore.



:DOCKER-ISOLATION-STAGE-2 - [0:0] :DOCKER-USER - [0:0] :KUBE-EXTERNAL-SERVICES - [0:0] :KUBE-FIREWALL - [0:0] :KUBE-FORWARD - [0:0] :KUBE-SERVICES - [0:0]

-A FORWARD -m comment --comment "kubernetes forwarding rules" -j KUBE-FORWARD

-A FORWARD -m conntrack --ctstate NEW -m comment --comment "kubernetes service portals" -i KUBE-SERVICES

A FORWARD -i DOCKER-USER

-A FORWARD -j DOCKER-ISOLATION-STAGE-1

-A FORWARD -o docker0 -m conntrack --ctstate RELATED, ESTABLISHED -j ACCEPT

-A FORWARD -o docker0 -j DOCKER

-A FORWARD -i docker0 ! -o docker0 -j ACCEP

-A FORWARD -i docker0 -o docker0 -j ACCEPT

-A FORWARD -o docker_gwbridge -m conntrack --ctstate RELATED,ESTABLISHED -j ACCEPT

 A FORWARD -o docker gwbridge -i DOCKER -A FORWARD -i docker_gwbridge ! -o docker_gwbridge -j ACCEPT

-A FORWARD -i docker_gwbridge -o docker_gwbridge -j DROP

-A OUTPUT -m conntrack --ctstate NEW -m comment --comment "kubernetes service portals" -j KUBE-SERVICES -A OUTPUT - KUBE-FIREWALL, I.

-A DOCKER-ISOLATION-STAGE-1 -i docker0 ! -o docker0 -j DOCKER-ISOLATION-STAGE-2

-A DOCKER-ISOLATION-STAGE-1 -i docker_gwbridge!-o docker_gwbridge-j DOCKER-ISOLATION-STAGE-2

-A DOCKER-ISOLATION-STAGE-1 - RETURN

-A DOCKER-ISOLATION-STAGE-2 -o docker0 -j DROP

A DOCKER-ISOLATION-STAGE-2 -o docker gwbridge -i DROF

-A DOCKER-ISOLATION-STAGE-2 - j RETURN -A DOCKER-USER - RETURN

-A KUBE-FIREWALL -m comment --comment "kubernetes firewall for dropping marked packets" -m mark --mark 0x8000/0x8000 -i DROP

-A KUBE-FORWARD -m conntrack --ctstate INVALID -j DROP

-A KUBE-FORWARD -m comment --comment "kubernetes forwarding rules" -m mark --mark 0x4000/0x4000 -i ACCEPT

-A KUBE-FORWARD -s 10.217.0.0/16 -m comment --comment "kubernetes forwarding conntrack pod source rule" -m conntrack --ctstate RELATED,ESTABLISHEI

-A KUBE-FORWARD -d 10.217.0.0/16 -m comment -comment "kubernetes forwarding conntrack pod destination rule" -m conntrack --ctstate RELATED, ESTABL

-A KUBE-SERVICES -d 10.99.38.155/32 -p tcp -m comment --comment "default/nginx-59: has no endpoints" -m tcp --dport 80 -j REJECT --reject-with icmp-port-

-A KUBE-SERVICES -d 10.96.61.252/32 -p tcp -m comment --comment "default/nginx-64: has no endpoints" -m tcp --dport 80 -j REJECT --reject-with icmp-port-

A KUBE-SERVICES -d 10.104.166.10/32 -p tcp -m comment -comment "default/nginx-67: has no endpoints" -m tcp --dport 80 -j REJECT --reject-with icmp-por

A KUBE-SERVICES -d 10.98.85.41/32 -p tcp -m comment -comment "default/nginx-9: has no endpoints" -m tcp --dport 80 -j REJECT --reject-with icmp-port-un -A KUBE-SERVICES -d 10.97.138.144/32 -p tcp -m comment --comment "default/nginx-17: has no endpoints" -m tcp --dport 80 -j REJECT --reject-with icmp-por

-A KUBE-SERVICES -d 10.106.49.80/32 -p tcp -m comment --comment "default/nginx-37: has no endpoints" -m tcp --dport 80 -j REJECT --reject-with icmp-port-

-A KUBE-SERVICES -d 10.104.164.205/32 -p tcp -m comment --comment "default/nginx-5: has no endpoints" -m tcp --dport 80 -j REJECT --reject-with icmp-por -A KUBE-SERVICES -d 10.104.25.150/32 -p tcp -m comment --comment "default/nginx-19: has no endpoints" -m tcp --dport 80 -j REJECT --reject-with icmp-por

4 KUBE-SERVICES -d 10.106.234.213/32 -p tcp -m comment --comment "default/nginx-88: has no endpoints" -m tcp --dport 80 -i REJECT --reject-with icmp-po

-A KUBE-SERVICES -d 10.109.209.136/32 -p tcp -m comment --comment "default/nginx-33: has no endpoints" -m tcp --dport 80 -j REJECT --reject-with icmp-po

-A KUBE-SERVICES -d 10.106.196.105/32 -p tcp -m comment --comment "default/nginx-49: has no endpoints" -m tcp --dport 80 -j REJECT --reject-with icmp-po -A KUBE-SERVICES -d 10.111.101.6/32 -p tcp -m comment -comment "default/nginx-53: has no endpoints" -m tcp --dport 80 -j REJECT --reject-with icmp-port-

-A KUBE-SERVICES -d 10.110.226.230/32 -p tcp -m comment --comment "default/nginx-79: has no endpoints" -m tcp --dport 80 -j REJECT --reject-with icmp-po

-A KUBE-SERVICES -d 10.98.99.136/32 -p tcp -m comment --comment "default/nginx-6: has no endpoints" -m tcp --dport 80 -j REJECT --reject-with icmp-port-u -A KUBE-SERVICES -d 10.99.75.233/32 -p tcp -m comment --comment "default/nginx-7: has no endpoints" -m tcp --dport 80 -j REJECT --reject-with icmp-port-u

-A KUBE-SERVICES -d 10.108.41.202/32 -p tcp -m comment --comment "default/nginx-14: has no endpoints" -m tcp --dport 80 -i REJECT --reject-with icmp-por

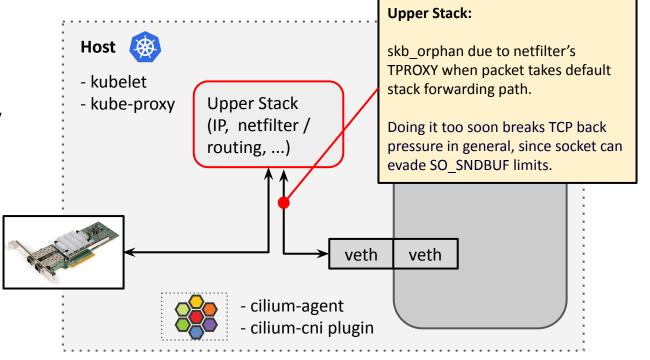
-A KUBE-SERVICES -d 10.97.36.249/32 -p tcp -m comment -comment "default/nginx-99: has no endpoints" -m tcp --dport 80 -j REJECT --reject-with icmp-port--A KUBE-SERVICES -d 10.98.213.37/32 -p tcp -m comment -comment "default/nginx-77: has no endpoints" -m tcp --dport 80 -j REJECT --reject-with icmp-port--A KUBE-SERVICES -d 10.107.229.31/32 -p tcp -m comment --comment "default/ngipx-92; has no endpoints" -m tcp --dport 80 -i REJECT --reject-with icmp-por

Standard K8s Datapath Architecture:



Problems:

- kube-proxy scalability
- Routing via upper stack
- Potential reasons:
 - Cannot replace kube-proxy
 - Custom netfilter rules
 - Just "went with defaults"



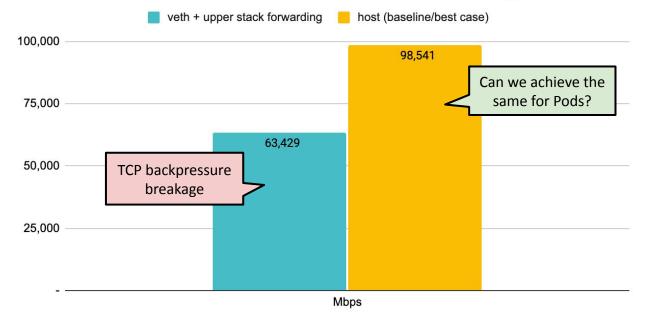
Standard K8s Datapath Architecture:



Problems:

- kube-proxy scalability
- Routing via upper stack
- Potential reasons:
 - Cannot replace kube-proxy
 - Custom netfilter rules
 - Just "went with defaults"

TCP stream single flow Pod to Pod over wire, 8k MTU (higher is better)

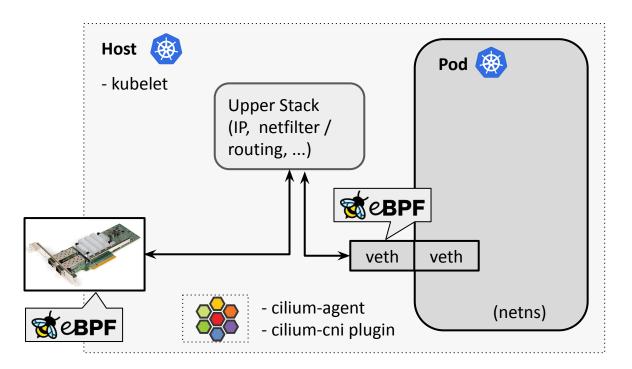


* 8264 MTU for data page alignment in GRO





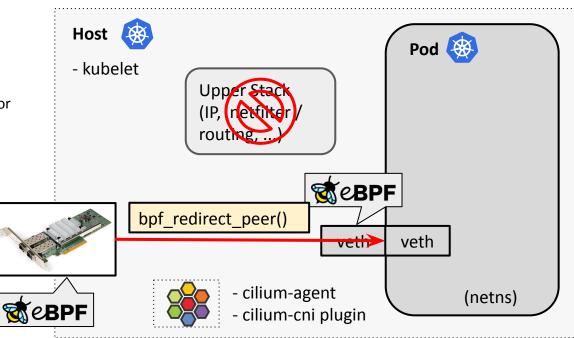
- BPF Host Routing





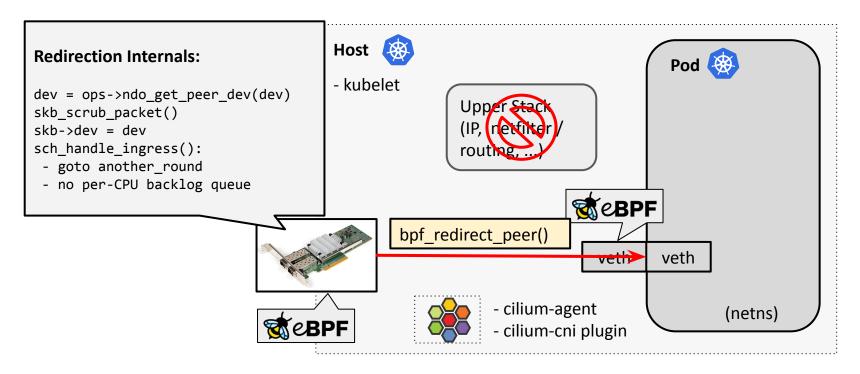


- BPF Host Routing
 - → Routing only via tc BPF layer
 - → Fast netns switch on ingress
 - → Helper for fib + dynamic neighbor resolution on egress





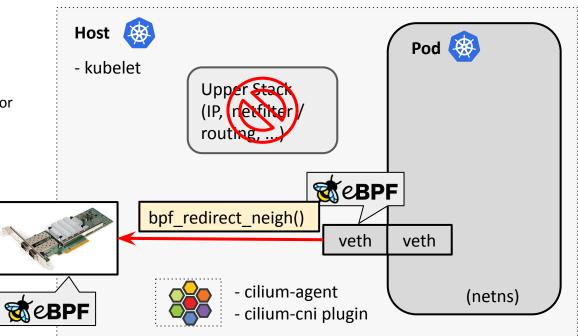






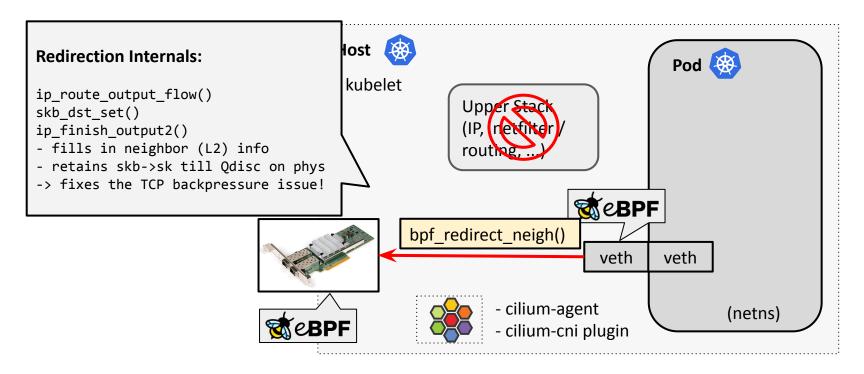


- BPF Host Routing
 - → Routing only via tc BPF layer
 - → Fast netns switch on ingress
 - → Helper for fib + dynamic neighbor resolution on egress





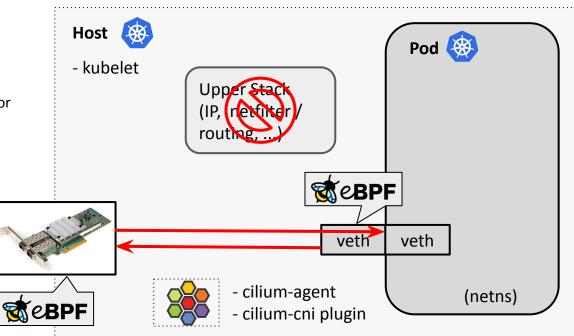








- BPF Host Routing
 - → Routing only via tc BPF layer
 - → Fast netns switch on ingress
 - → Helper for fib + dynamic neighbor resolution on egress

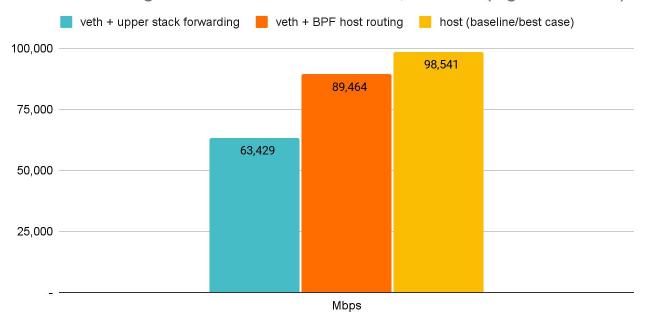






- BPF Host Routing

TCP stream single flow Pod to Pod over wire, 8k MTU (higher is better)



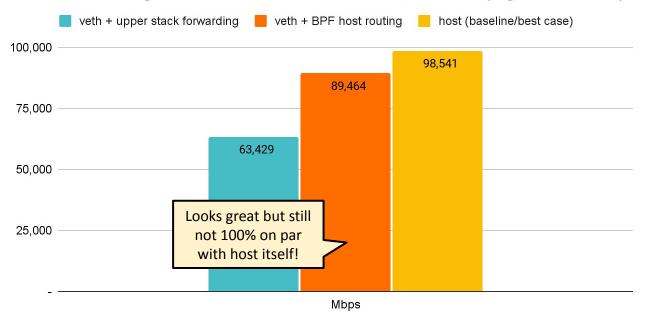
* 8264 MTU for data page alignment in GRO





- BPF Host Routing

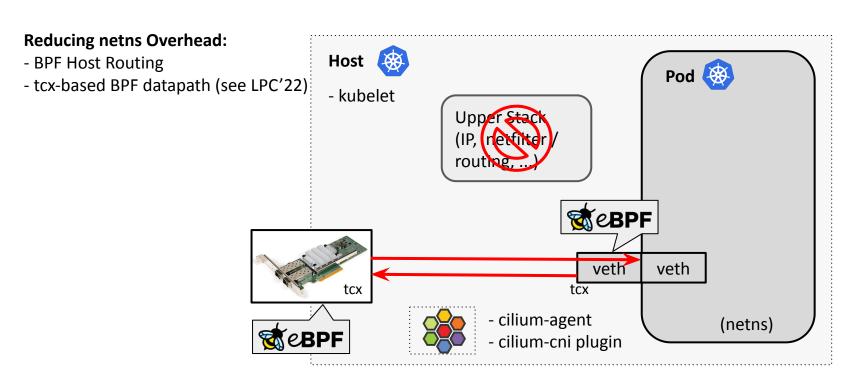
TCP stream single flow Pod to Pod over wire, 8k MTU (higher is better)



^{* 8264} MTU for data page alignment in GRO



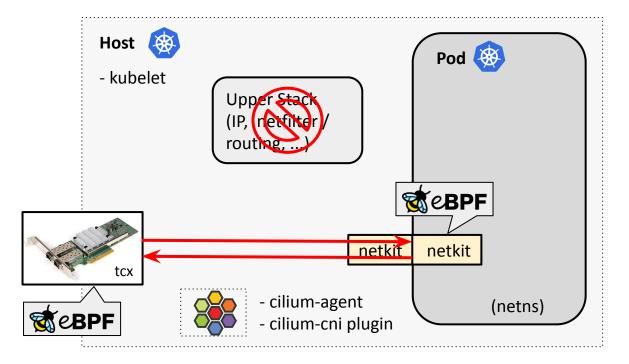


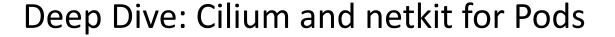






- BPF Host Routing
- tcx-based BPF datapath
- netkit devices for Pods

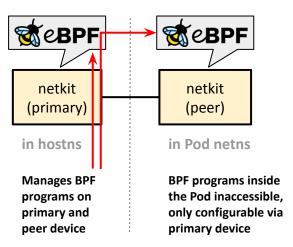






netkit programmable virtual devices for BPF:

- Going forward, Cilium's CNI code will set up netkit devices for Pods instead of veth
- BPF program via bpf mprog is part of the driver's xmit routine, allowing fast egress netns switch
- Driver implements ndo_get_peer_dev, allowing fast ingress netns switch
- Configurable as L3 device (default) or L2 device
- Configurable default drop-all if no BPF is attached



```
static netdev tx t netkit xmit(struct sk buff *skb, struct net device *dev)
       struct netkit *nk = netkit priv(dev);
       enum netkit action ret = READ ONCE(nk->policy);
       netdev tx t ret dev = NET XMIT SUCCESS;
       const struct bpf mprog entry *entry;
       struct net device *peer;
                                                                                default policy if nothing attached
       rcu read lock();
       peer = rcu dereference(nk->peer);
       if (unlikely(!peer | !(peer->flags & IFF UP) ||
                                                                                skb scrubbing
                    !pskb may pull(skb, ETH HLEN)
                    skb orphan frags(skb, GFP ATOMIC)))
               goto drop;
       netkit prep forward(skb, !net eq(dev net(dev), dev net(peer)));
                                                                                 netns switch, provides hostns context for
       skb->dev = peer;
       entry = rcu dereference(nk->active);
                                                                                bpf fib lookup(), bpf redirect*(), etc.
       if (entry)
               ret = netkit run(entry, skb, ret);
       switch (ret) {
                                                                                executes active bpf mprog array
       case NETKIT NEXT:
       case NETKIT PASS:
               skb->protocol = eth type trans(skb, skb->dev);
               skb postpull rcsum(skb, eth hdr(skb), ETH HLEN);
                                                                                for redirect into local hostns
               netif rx(skb);
               break;
       case NETKIT REDIRECT:
               skb do redirect(skb);
                                                                                performs direct redirect to phys device
               break:
       case NETKIT DROP:
       default:
drop:
               kfree skb(skb);
               dev core stats tx dropped inc(dev);
               ret dev = NET XMIT DROP;
               break;
       rcu read unlock();
       return ret dev;
```

19

Deep Dive: netkit Internals /2



BPF entry point

```
struct bpf mprog fp {
        struct bpf prog *prog;
};
struct bpf mprog cp {
        struct bpf link *link;
};
struct bpf mprog entry {
        struct bpf mprog fp fp items[BPF MPROG MAX];
        struct bpf mprog bundle *parent;
};
struct bpf mprog bundle {
        struct bpf mprog entry a;
        struct bpf mprog entry b;
        struct bpf mprog cp cp items[BPF MPROG MAX];
        struct bpf prog *ref;
        atomic64 t revision;
        u32 count;
};
```

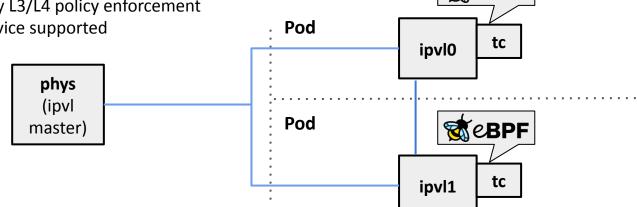
Cilium and ipvlan



Cilium used to have limited ipvlan support in the past

- Cilium CNI sets up ipvlan slave device for the target netns with a simple BPF program doing a tailcall
 - BPF tailcall map resides in hostns so that Cilium agent has access to it (it cannot enter Pod's netns)
- Cilium's bpf_lxc program added to tailcall map
- Limitations:
- L3 mode comes with netfilter asymmetry -> L3S needed
- BPF programs @ ipvlan slave could be unloaded
- No L7 proxy (host), only L3/L4 policy enforcement





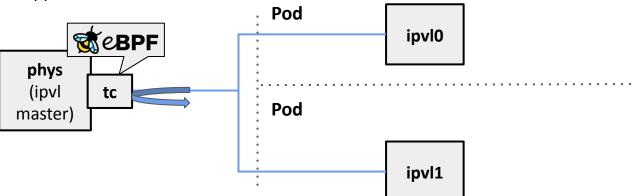
TEBPF

Cilium and ipvlan



Cilium used to have limited ipvlan support in the past

- What could have been done better? Cilium could have used L3(S) private mode (forces all traffic out to hostns)
- Cilium's bpf_lxc program added to tailcall map which is accessed at tc egress of physical device
- Pod to Pod communication goes to phys egress, bpf_redirect to ingress to loop back to Pod
- Limitations (remainder):
- Traffic classification for tailcall map entry needed at tc egress, might have allowed Pod to spoof their source
- L3 mode comes with netfilter asymmetry -> L3S needed
- Only single physical device supported

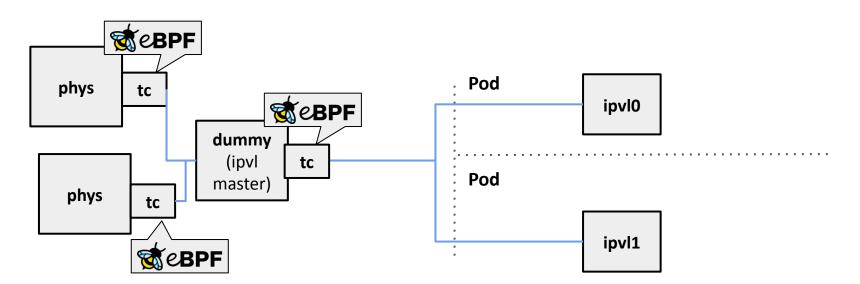






Cilium used to have limited ipvlan support in the past

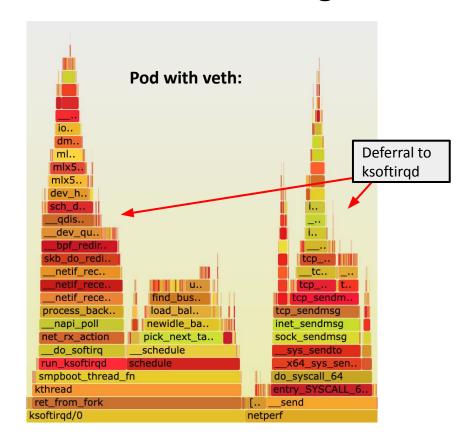
- How might multiple physical devices look in this architecture picture?

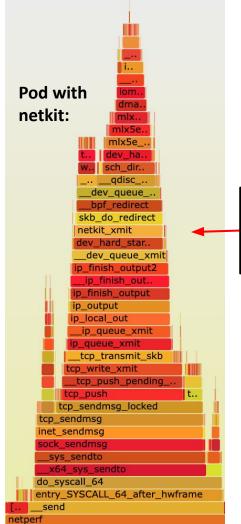


Comparison veth, ipvlan, netkit

	veth	ipvlan	netkit
Operation mode:	L2	L3 (or L2)	L3 (or L2)
Device "legs":	pair (e.g. 1 host, 1 Pod) veth0 veth1	1 "master" device (e.g. physical device), n "slave" devices ipvl0 ipvl1	pair (e.g. 1 host, 1 Pod) with "primary" and "peer" device
BPF programming:	tc(x) BPF on host device*	In host with tc(x) via "master" device (only entity in host) *	In Pod, BPF is native part of "peer" device internals
Routing:	L2 gateway (+ kernel FIB)	ipvlan internal FIB + kernel FIB	kernel FIB e.g. bpf_fib_lookup()
Problems:	Needs L2 neigh resolution, Higher overhead due to per-CPU backlog queue, native XDP support but very slow and hard to use.	Inflexible for multiple physical devices & troubleshooting, cumbersome to program BPF on "master". ipvlan needs to be operated in L3/private mode for Pod policy enforcement.	Still one device per Pod inside host, for some use-cases the host device can be removed fully (wip).

veth vs netkit: Backlog Queue





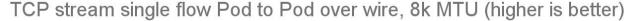


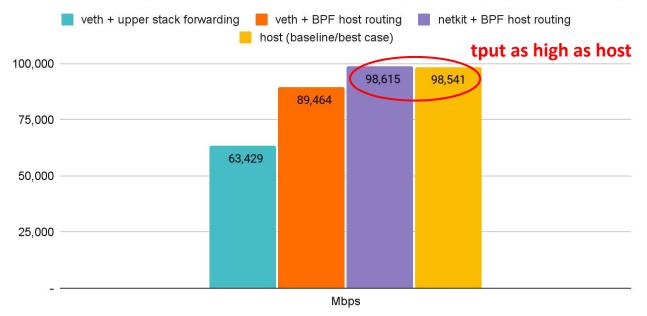
Remains in process context all the way, allows for better scheduler accounting





- BPF Host Routing
- tcx-based BPF datapath layer
- netkit devices for Pods



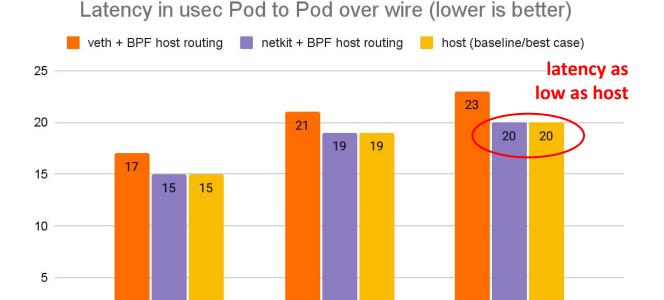


* 8264 MTU for data page alignment in GRO





- BPF Host Routing
- tcx-based BPF datapath layer
- netkit devices for Pods



P90

MIN

P99





iproute2 and <u>vishvananda/netlink</u> library support

- Goal: Basic device setup and introspection support

```
# ip link add type netkit
# ip -d a
[...]
7: nk0@nk1: <BROADCAST,MULTICAST,NOARP,M-DOWN> mtu 1500 qdisc noop state DOWN group default qlen 1000
    link/ether 00:00:00:00:00:00 brd ff:ff:ff:ff:ff:ff promiscuity 0 allmulti 0 minmtu 68 maxmtu 65535
    netkit mode 13 type peer policy forward numtxqueues 1 numrxqueues 1 [...]
8: nkl@nk0: <BROADCAST,MULTICAST,NOARP,M-DOWN> mtu 1500 qdisc noop state DOWN group default qlen 1000
    link/ether 00:00:00:00:00:00 brd ff:ff:ff:ff:ff:ff promiscuity 0 allmulti 0 minmtu 68 maxmtu 65535
    netkit mode 13 type primary policy forward numtxqueues 1 numrxqueues 1 [...]
```

- Support base setup and delegate BPF program management to applications (via <u>libbpf</u>, ebpf-go)
- vishvananda/netlink needed for Cilium CNI plugin integration (Go-based)





Fixing networking stats for netkit in general and for peer-redirection

- Goal: Proper network stats accounting for netkit and veth

From: Peilin Ye <peilin.ye@bytedance.com>

Traffic redirected by bpf_redirect_peer() (used by recent CNIs like Cilium) is not accounted for in the RX stats of supported devices (that is, veth and netkit), confusing user space metrics collectors such as cAdvisor [0], as reported by Youlun.

- Fix is calling dev_sw_netstats_rx_add() in skb_do_redirect() and move netkit & veth to dev->tstats
- Guard if drivers implementing ndo_get_peer_dev and do not use dev->tstats
- Suggestion from Jakub Kicinski to move {I,t,d}stats allocation into net core



Adding peer pointer into struct net device

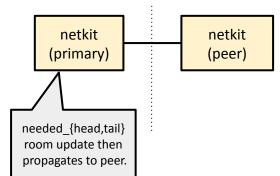
- Goal: Get rid of ndo get peer deventirely and add peer pointer to net device (suggestion from Jakub Kicinski)
- The latter is only implemented by veth and netkit
- Helps performance for ingress direction due to the current indirect call in skb do redirect()
- Short-term: INDIRECT CALL 1() macro can be utilized on the ndo if driver is built-in
- Mid-term: Rework veth and netkit and remove private peer pointers in favor of dev->peer

```
if (flags & BPF F PEER) {
        const struct net device ops *ops = dev->netdev ops;
        if (unlikely(!ops->ndo get peer dev |
                     !skb at tc ingress(skb)))
                goto out drop;
        dev = ops->ndo get peer dev(dev);
        if (unlikely(!dev | |
                     !(dev->flags & IFF UP)
                     net eq(net, dev net(dev))))
                goto out drop:
        skb->dev = dev:
        return -EAGAIN;
```



Experimenting with head/tailroom customization

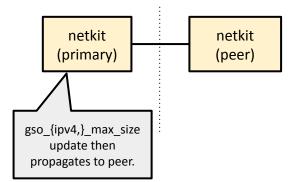
- Goal: Being able to control dev->needed_headroom and dev->needed_tailroom
- Could benefit datapath performance under tunneling (vxlan, geneve) or encryption (wireguard)
 - Potentially avoids pskb_expand_head() reallocation costs
- Idea: Have actual IFLA_HEADROOM and IFLA_TAILROOM attributes to dump and set on a device
- needed_{head,tail}room is by default 0, vxlan/geneve adjusts needed_headroom, wireguard also needed_tailroom
- Performance benefit: to be measured, references from <u>old patches mention</u> costs around 5% on realloc





Adding new ndo for setting dev->gso_{ipv4,}_max_size

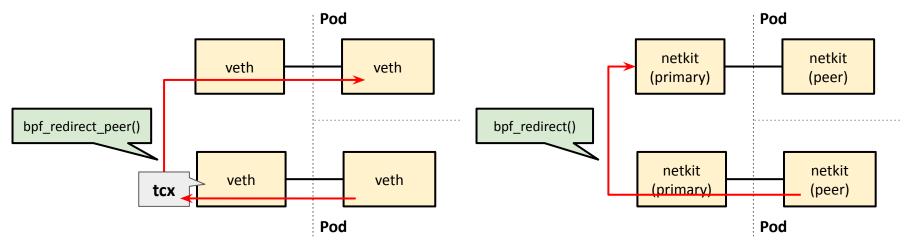
- Goal: Enabling BIG TCP for Pods without having to restart Pods
- Cilium agent is not able to exec into the Pod's netns at runtime and mounting host procfs into Cilium container is not desired (security reasons). Only the Cilium CNI plugin has access when setting up devices.
- Downside: Enabling BIG TCP on an existing cluster requires restart of application Pod
- New ndo for updating dev->gso_{ipv4,}_max_size in similar style as dev->needed_{head,tail}room would be desirable.. e.g. picks max of primary/peer and applies it to both





veth vs netkit difference in terms of bpf_redirect_peer

- Goal: Usability improvements to avoid users running into pitfalls
- Performance-wise local node's Pod-Pod redirection optimizations are on-par for veth and netkit
- Just that for veth bpf_redirect_peer (ingress->ingress) is needed whereas netkit bpf_redirect (egress->egress)
- Perhaps rather documentation/awareness issue?





Semi-related to netkit: Removing indirect calls for IPv6 in bpf_fib_lookup()

- Goal: Further improve bpf_fib_lookup() IPv6 performance by removing indirect calls from fast-path.

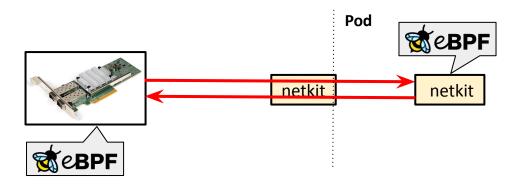
 Useful when netkit is utilising the latter out of its BPF programs.
- Worst case 4 indirect calls from a single BPF helper call (!) via bpf_ipv6_fib_lookup():
- ipv6_stub->fib6_{table_,}lookup()
- ipv6_stub->fib6_select_path()
- ipv6_stub->ip6_mtu_from_fib6()
- ipv6_bpf_stub->ipv6_dev_get_saddr()
- Two options to overcome:
- Remove IPv6 tristate from Kbuild, move to bool and then successively get rid of the stub helpers (my preference)
- INDIRECT_CALL_1() wrappers for the case when IPv6 is built-in





BPF kfunc for dev_queue_xmit_nit to allow for tcpdump at specific points in BPF program

- Goal: Improve troubleshooting for Pod-traffic
- bpf_redirect_peer() skips host device on ingressing traffic (netkit, veth), and bpf_redirect() from the BPF program inside the netkit device skips host device on egressing traffic
- One either needs to nsenter into Pod for troubleshooting or capture all traffic on the phys device
- BPF kfunc for dev_queue_xmit_nit-like functionality helps to dynamically insert tcpdump-tracing at custom points





Semi-related to netkit: Reorganize ndos in struct net_device_ops

- Goal: Reduce cacheline access for dev->netdev_ops to improve performance
- Picking up on the work Coco Li kicked off in order to look at fast-path RX ndo ops and fast-path TX ndo ops
- Wrapping them into __cacheline_group_{begin,end} markers with member and size assertions
- struct net_device had nice gains:

Reorganize fast path variables on tx-txrx-rx order Fastpath variables end after npinfo.

Below data generated with pahole on x86 architecture.

Fast path variables span cache lines before change: 12 Fast path variables span cache lines after change: 4

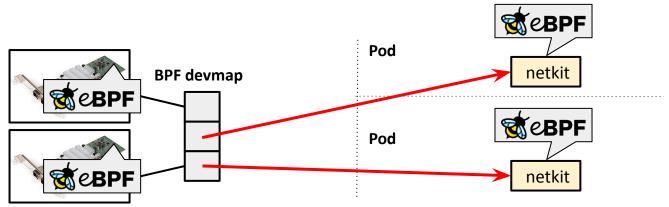
Signed-off-by: Coco Li Lixiaoyan@google.com>
Suggested-by: Eric Dumazet <edumazet@google.com>
Reviewed-by: David Ahern <dsahern@kernel.org>

Patches: optimize cacheline access (Coco Li)



netkit single device mode

- Goal: Could we even go further and remove the host-facing device?
- Single netkit device mode could act as sink within Pod
- BPF devmap in hostns keeps the device pointers of the peer device as map values, addressable by id key. Modified bpf_redirect_peer() would look up BPF devmap instead.
- BPF devmap can be shared among multiple phys devices inside hostns





Getting rid of indirect call overhead for bpf_mprog

- Goal: Get rid of indirect calls when invoking BPF programs (tcx, netkit)
- Similar to XDP use dynamically generated branch funnel
- More involved due to the array as opposed to just single program invocation, but might be needed anyway for XDP bpf_mprog support
- Needs benchmarking as branch funnel for netkit will be much larger than XDP

Cilium Datapath Architecture (journey 2019 - today):



All Building Blocks:

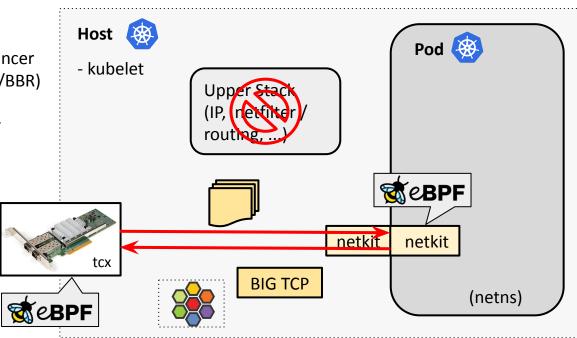
- BPF kube-proxy replacement
- XDP-based Service Load-Balancer
- Bandwidth Manager (fq/EDT/BBR)
- BPF Host Routing
- tcx-based BPF datapath layer
- netkit devices for Pods

Pushing even further:

- BIG TCP (IPv4/IPv6)

Future integration:

- TCP usec resolution ($\sqrt{6.7}$)
- BBRv3 (once upstream)





Thank you! Questions?

github.com/cilium/cilium

cilium.io

ebpf.io

BPF Host Routing

tcx BPF datapath

netkit devices

BIG TCP for IPv4/IPv6