

Unlocking the softirq lock for PREEMPT_RT

Linux Plumbers Conference

Sebastian A. Siewior

Linutronix GmbH

November 15, 2023

Softirq on PREEMPT_RT vs vanilla

- ❏ Preemptible. Context switch is possible.
- ❏ Runs after the threaded handler.
- ❏ No piggyback after hardirq.
- ❏ Everything from hardirq goes to ksoftirqd.
- ❏ Due to preemption *local_bh_disable()* is a per-CPU lock.

Results of the lock

- 🗄 Ressources depending on BH locking are protected. 👍
- 🗄 Long-running forced threaded interrupts block other forced threaded interrupts. 👎
- 🗄 Timer, Tasklets, ...block forced threaded interrupts. 👎
- 🗄 Best we can do is a PI-boost. The need-resched condition is never observed. 👎

Trace force-threaded interrupts preempted

```

irq/40-eno0-2034 D...2 681 softirq_raise: vec=3 [action=NET_RX]
irq/40-eno0-2034 ..s.2 681 softirq_entry: vec=3 [action=NET_RX]
irq/40-eno0-2034 d.H.3 690 irq_handler_entry: irq=35
irq/40-eno0-2034 dNH33 692 sched_wakeup: irq/35-ahci prio=44
irq/40-eno0-2034 d.s23 694 sched_switch: prio=49 R+>irq/35-ahci prio=44

irq/35-ahci-837 d..31 696 sched_pi_setprio: irq/40-eno0 prio 49 -> 44
irq/35-ahci-837 d..21 699 sched_switch: prio=44 D->irq/40-eno0 prio=44

irq/40-eno0-2034 d.s34 715 sched_wakeup: iperf3 prio=120
irq/40-eno0-2034 d..21 736 sched_switch: prio=49 R+>irq/35-ahci prio=44

irq/35-ahci-837 D..13 740 softirq_raise: vec=4 [action=BLOCK]
irq/35-ahci-837 ..s.2 740 softirq_entry: vec=4 [action=BLOCK]

```

But why exactly the lock

- ❏ A FPGA interrupt handler doing only wake up of userland.
- ❏ A CAN driver needs to inject packet.
- ❏ AHCI driver needs to do IO.
- ❏ `local_bh_disable()` and `enable()` plus the lock:
 - Preserve the raise softirq and invoke semantic.
 - Protect per-CPU resources. Mostly.

How to get rid of the lock?

- ❏ **Identify the per-CPU resources. Add a lock.**
- ❏ **The working PoC**
 - `local_lock_nested_bh()` (followed `preempt_disable_nested()`).
 - lockdep check for BH. CPU migration must be off (due to disabled BH).
 - A per-CPU lock only on RT.
 - The raise and invoke semantic of `softirqs` is the same.
 - The macros `in_serving_softirq()`, `softirq_count()` work unchanged.
 - guard notation.

Example struct napi_alloc_cache

```
@@ -264,6 +264,7 @@ static void *page_frag_alloc_1k(struct
    struct napi_alloc_cache {
+        local_lock_t bh_lock;
        struct page_frag_cache page;
        struct page_frag_1k page_small;
        unsigned int skb_count;
@@ -295,6 +298,7 @@ void *__napi_alloc_frag_align(unsigned int
    struct napi_alloc_cache *nc = this_cpu_ptr(&napi_alloc_cache);

    fragsz = SKB_DATA_ALIGN(fragsz);
+    guard(local_lock_nested_bh)(&napi_alloc_cache.bh_lock);

    return page_frag_alloc_align(&nc->page, fragsz, GFP_ATOMIC,
                                align_mask);
}
```

Example struct napi_alloc_cache

```
@@ -725,9 +730,11 @@ struct sk_buff *__netdev_alloc_skb(struct net_device
    pfmemalloc = nc->pfmemalloc;
} else {
    local_bh_disable();
-   nc = this_cpu_ptr(&napi_alloc_cache.page);
-   data = page_frag_alloc(nc, len, gfp_mask);
-   pfmemalloc = nc->pfmemalloc;
+   scoped_guard(local_lock_nested_bh, &napi_alloc_cache.bh_lock) {
+       nc = this_cpu_ptr(&napi_alloc_cache.page);
+       data = page_frag_alloc(nc, len, gfp_mask);
+       pfmemalloc = nc->pfmemalloc;
+   }
    local_bh_enable();
}
```


Example softnet_data: xmit.recursion_lock

```
@@ -3208,6 +3208,10 @@ struct softnet_data {
    #endif
        /* written and read only by owning cpu: */
        struct {
+       #ifdef CONFIG_PREEMPT_RT
+           struct task_struct *recursion_owner;
+           local_lock_t recursion_lock;
+       #endif
            u16 recursion;
            u8  more;
+       #ifdef CONFIG_NET_EGRESS
```

Example softnet_data: xmit.recursion_lock

```
@@ -3272,6 +3276,27 @@ static inline bool dev_xmit_recursion(void)
                                XMIT_RECURSION_LIMIT);
    }

#ifdef CONFIG_PREEMPT_RT
+
+static inline void dev_xmit_recursion_inc(void)
+{
+    if (__this_cpu_read(softnet_data.xmit.recursion_owner) != current) {
+        local_lock_nested_bh(&softnet_data.xmit.recursion_lock);
+        __this_cpu_write(softnet_data.xmit.recursion_owner, current);
+    }
+    __this_cpu_inc(softnet_data.xmit.recursion);
+}
```

Example BPF, net/core/filter.c

```
@@ -85,6 +85,11 @@
+DEFINE_PER_CPU(struct bpf_run_lock, bpf_run_lock) = {
+    .redirect_lock = INIT_LOCAL_LOCK(redirect_lock),
+};
+EXPORT_PER_CPU_SYMBOL_GPL(bpf_run_lock);
@@ -4205,6 +4210,7 @@ static const struct bpf_func_proto
    bpf_xdp_adjust_meta_proto = {
        void xdp_do_flush(void)
        {
+            guard(local_lock_nested_bh)(&bpf_run_lock.redirect_lock);
            __dev_flush();
            __cpu_map_flush();
            __xsk_map_flush();
```

Example BPF supports redirect

```
--- a/drivers/net/ethernet/amazon/ena/ena_netdev.c
+++ b/drivers/net/ethernet/amazon/ena/ena_netdev.c
@@ -385,6 +385,7 @@ static int ena_xdp_execute(struct ena_ring *rx_ring,
    struct xdp_buff *xdp)
    {
        if (!xdp_prog)
            goto out;

+       guard(local_lock_nested_bh)(&bpf_run_lock.redirect_lock);
        verdict = bpf_prog_run_xdp(xdp_prog, xdp);

        switch (verdict) {
```

Example BPF driver, no redirect

```
--- a/drivers/net/ethernet/cavium/thunder/nicvf_main.c
+++ b/drivers/net/ethernet/cavium/thunder/nicvf_main.c
@@ -554,7 +554,8 @@ static inline bool nicvf_xdp_rx(struct nicvf *nic,
    xdp_prepare_buff(&xdp, hard_start, data - hard_start, len, false);
    orig_data = xdp.data;

-   action = bpf_prog_run_xdp(prog, &xdp);
+   scoped_guard(local_lock_nested_bh, &bpf_run_lock.redirect_lock)
+       action = bpf_prog_run_xdp(prog, &xdp);

    len = xdp.data_end - xdp.data;
    /* Check if XDP program has changed headers */
```

Example BPF driver, move REDIRECT bits, cpsw_priv.c

```
    return CPSW_XDP_PASS;

- act = bpf_prog_run_xdp(prog, xdp);
- /* XDP prog might have changed packet data and boundaries */
- *len = xdp->data_end - xdp->data;
+ scoped_guard(local_lock_nested_bh, &bpf_run_lock.redirect_lock) {
+     act = bpf_prog_run_xdp(prog, xdp);
+     *len = xdp->data_end - xdp->data;
+     if (act == XDP_REDIRECT) {
+         if (xdp_do_redirect(ndev, xdp, prog))
+             goto drop;
+     }
+ }

    switch (act) {
case XDP_PASS:
```

Trace force-threaded interrupts preempted, patched

```

irq/38-eno0-2006 D...1 032 softirq_raise: vec=3 [action=NET_RX]
irq/38-eno0-2006 ..s.1 032 softirq_entry: vec=3 [action=NET_RX]
irq/38-eno0-2006 d.H.1 033 irq_handler_entry: irq=35 name=ahci
irq/38-eno0-2006 dNH31 034 sched_wakeup: irq/35-ahci prio=44
irq/38-eno0-2006 d.s21 035 sched_switch: prio=49 R+>irq/35-ahci prio=44

irq/35-ahci-842 D..12 038 softirq_raise: vec=4 [action=BLOCK]
irq/35-ahci-842 ..s.1 039 softirq_entry: vec=4 [action=BLOCK]
irq/35-ahci-842 d.s32 041 sched_wakeup: grep prio=120
irq/35-ahci-842 ..s.1 042 softirq_exit: vec=4 [action=BLOCK]
irq/35-ahci-842 d..2. 043 sched_switch: prio=44 S->irq/38-eno0 prio=49

irq/38-eno0-2006 ..s.1 044 softirq_exit: vec=3 [action=NET_RX]
irq/38-eno0-2006 d..2. 051 sched_switch: prio=49 S->swapper/2 prio=120

```

Need to look at tw_timer_handler, allocation

```
struct inet_timewait_sock *inet_twsk_alloc(const struct sock *sk,
{
    struct inet_timewait_sock *tw;
    ...

    tw = kmem_cache_alloc(sk->sk_prot_creator->twsk_prot->twsk_slab,
                          GFP_ATOMIC);
    if (tw) {
    ...

        timer_setup(&tw->tw_timer, tw_timer_handler, TIMER_PINNED);
    }
```



```
void tcp_time_wait(struct sock *sk, int state, int timeo)
{
...
    tw = inet_twsk_alloc(sk, &net->ipv4.tcp_death_row, state);
    if (tw) {
...
        /* tw_timer is pinned, so we need to make sure BH are disabled
         * in following section, otherwise timer handler could run before
         * we complete the initialization.
         */
        local_bh_disable();
        inet_twsk_schedule(tw, timeo);
        /* Linkage updates.
         * Note that access to tw after this point is illegal.
         */
        inet_twsk_hashdance(tw, sk, net->ipv4.tcp_death_row.hashinfo);
        local_bh_enable();
    }
}
```

The page_pool is probably safe

- ❏ Acquires memory in softirq.
- ❏ Returns memory in softirq.
- ❏ If preempted by another NAPI instance then it ends up in slowpath.

Thank you for your attention

**Special thanks to the Linux Foundation
for supporting our efforts to
bring PREEMPT_RT mainline.**

`<bigeasy@linutronix.de>`